

# CS2102: Database Systems (AY2022-2023 – Sem 2)

## Final Exam

### Instructions

1. Please read **ALL** instructions carefully.
2. This assessment contains **SIXTEEN (16)** questions including subquestions:
  - (a) There are 11 Multiple Response Questions (MRQ):  
1.2 - 1.4, 2.2, 3.1b, 3.2b, 4.1(a-d) - 4.2
  - (b) There is 1 Fill in the Blank Questions (FITB): 1.1
  - (c) There are 2 Multiple Choice Question (MCQ): 2.1, 4.3,
  - (d) There are 2 True/False Question: 3.1a, 3.2a
3. There is a total of **EIGHTY (80)** points for this assessment.
4. Answer **ALL** questions.
5. All the assessment is be done using Exemplify:
  - (a) This is a secure assessment (no internet ; no other software)
  - (b) This is a closed-book exam ; you are allowed 1 A4 cheat sheet double-sided
  - (c) The choices on Exemplify may be in a different order
6. No additional time will be given to submit.
7. Use the question number shown on Exemplify when asking question.
  - If the answer is clear from the question pdf/Exemplify, we will reply with “No Comment”.
8. Failure to follow each of the instructions above may result in deduction of your marks.

**Good Luck!**

# 1 SQL

We want to create a database for room booking. Our database should contain:

- **Employees**
  - Identified by employee id (INT).
  - Must store the name of the employee (TEXT).
- **Rooms**
  - Identified by floor number (INT) and room number (INT).
  - Must store the name of the room (TEXT).

A room can be booked by employees following the constraints below.

- When a room is booked, it is booked for the whole day.
- A room can be booked by at most one employee on any single day. In other words, no two different employees can book the same room on the same day. However, it is possible that the room is not booked by any employee.
- An employee can book at most one room on any single day. In other words, no two different rooms can be booked by the same employee on the same day. However, it is possible that the employee does not book any room.

## 1.1 (14 points) **CREATE TABLE.** *This questions is split into two on Exemplify*

Your answer for each blank should only be exactly **one** simple constraint (e.g., *UNIQUE*, etc) which cannot be a **CHECK** constraint or an attribute. You should not add unnecessary constraints. Unnecessary constraints includes constraints not specified above as well as constraints already enforced in other parts (e.g., *NOT NULL* or *UNIQUE* on an attribute with *PRIMARY KEY* constraint). If you feel that there should be nothing inside the box (e.g., *because it will only be unnecessary constraints*), simply put in “-” (*without the quote*). An empty box is treated as *unanswered*. Lastly, in some cases, the order may not matter (e.g, *some primary key or foreign key attributes ordering*).

```
CREATE TABLE Employees (  
    eid      INT      ,  
    ename    TEXT       
);
```

```
CREATE TABLE Rooms (  
    rfloor INT      ,  
    rnum   INT      ,  
    rname  TEXT     ,  
    PRIMARY KEY ( ,  )  
);
```

```

CREATE TABLE Bookings (
    bfloor INT  ,
    bnum INT  ,
    beid INT  ,
    bday DATE  ,
    PRIMARY KEY (  ,  ,  ),
    UNIQUE (  ,  ),
    FOREIGN KEY (  ,  ) REFERENCES Rooms(rnum, rfloor),
    FOREIGN KEY (  ) REFERENCES Employees
);

```

### NOTES:

Primary key of **Rooms** has to be in that specific order because of foreign key of **Bookings** referencing that specific order. Primary key of **Bookings** as well as its unique constraints can be in any order.

**beid** of **Bookings** (*and only this from **Bookings***) has to have NOT NULL because we require candidate keys to be enforced additionally with NOT NULL. Note that **beid** is in UNIQUE constraints because this is a candidate key with 2 attributes. **bday** of **Bookings** should not have NOT NULL because it is already inside a PRIMARY KEY constraint.

**1.2** (6 points) **INSERT INTO**. Assume that the following insert statements have been performed successfully.

```

INSERT INTO Employees VALUES (1, 'A');
INSERT INTO Employees VALUES (2, 'B');
INSERT INTO Employees VALUES (3, 'C');
INSERT INTO Rooms VALUES (1, 1, 'R1');
INSERT INTO Rooms VALUES (1, 2, 'R2');
INSERT INTO Rooms VALUES (2, 1, 'R3');

```

Select ALL insert statement that will **NOT** be rejected. For choices with two insert statements, select only if none are rejected. Consider each choices independently and assumes all the constraints are enforced correctly. Note that dates can be inserted as a string in the format 'YYYY-MM-DD', where YYYY is the year, MM is the month, and DD is the date.

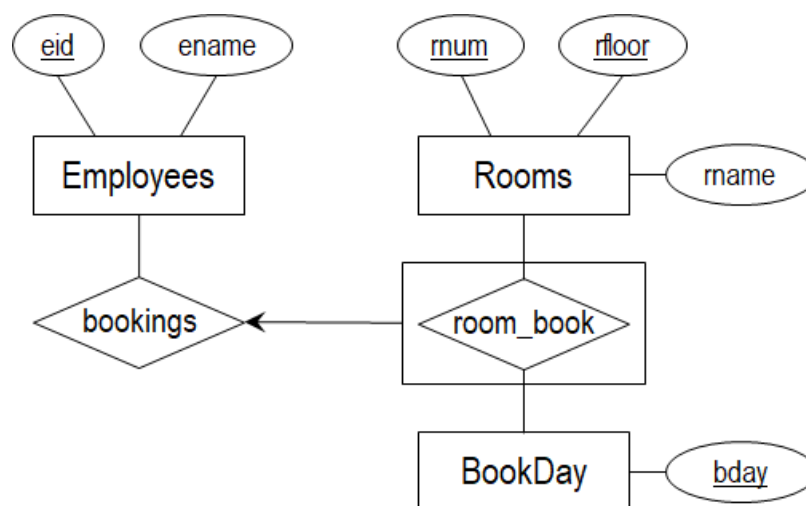
- ☒ (A) INSERT INTO Bookings VALUES (1, 2, 3, '2023-04-29');
- ☐ (B) INSERT INTO Bookings VALUES (2, 2, 2, '2023-04-29');
- ☐ (C) INSERT INTO Bookings VALUES (3, 2, 1, '2023-04-29');
- ☒ (D) INSERT INTO Bookings VALUES (1, 1, 1, '2023-04-29');
- ☐ (E) INSERT INTO Bookings VALUES (1, 2, 1, '2023-04-29');  
INSERT INTO Bookings VALUES (2, 1, 1, '2023-04-29');

- (✓) INSERT INTO Bookings VALUES (1, 2, 1, '2023-04-29');  
INSERT INTO Bookings VALUES (2, 1, 2, '2023-04-29');
- (G) INSERT INTO Bookings VALUES (2, 2, 2, '2023-04-29');  
INSERT INTO Bookings VALUES (2, 2, 1, '2023-04-29');
- (H) INSERT INTO Bookings VALUES (1, 2, 3, '2023-04-29');  
INSERT INTO Bookings VALUES (3, 2, 1, '2023-04-29');
- (I) *None of the above*

#### NOTES:

This does not depend on the order of **PRIMARY KEY** but only on the order of attributes in the **CREATE TABLE**. So, for **Bookings** specifically, we have the following order of attributes (**bfloor**, **bnun**, **beid**, **bday**).

**1.3 (5 points) ER Diagram.** Now consider the following ER diagram representation.



Select ALL constraints that are **NOT** enforced by the ER diagram.

- (A) Employees are identified by their employee id.
- (B) Rooms are identified by floor number and room number.
- (C) No two different employees can book the same room on the same day.
- (D) It is possible that a room is not booked by any employee.
- (✓) No two different rooms can be booked by the same employee on the same day.
- (F) It is possible that an employee does not book any room.
- (G) *None of the above*

**NOTES:**

For some, it is probably easier to see if you look at the functional dependencies.

(A)  $\{eid\} \rightarrow \{ename\}$

(B)  $\{rnum, rfloor\} \rightarrow \{rname\}$

(C)  $\{rnum, rfloor, bday\} \rightarrow \{eid\}$   
*(so now, you cannot have two different employees)*

(D) This is because there is no total participation constraints of room\_book with respect to bookings (*as well as Rooms with respect to room\_book*).

(✓)  $\{eid, bday\} \not\rightarrow \{rnum, rfloor\}$   
*(so now, you can have two different rooms!)*

(F) This is because there is no total participation constraints of Employees with respect to bookings.

**1.4 (5 points) SQL Query.** Consider the following instance of table Bookings.

Bookings			
bfloor	bnum	beid	bday
1	1	1	2023-04-29
1	3	3	2023-04-29
1	4	2	2023-04-29
1	2	1	2023-04-29
2	2	2	2023-04-29
1	1	3	2023-04-28
2	1	3	2023-04-28
2	2	3	2023-04-28
2	3	2	2023-04-28
1	2	1	2023-04-27
2	1	2	2023-04-27
1	1	3	2023-04-27

Select ALL the rows that will appear in the result if we run the following query.

```
SELECT bfloor, COUNT(DISTINCT beid)
FROM   Bookings
WHERE  bday > '2023-04-27'
GROUP BY bday, bfloor;
```

(✓) 

1	1
---	---

(✓) 

1	3
---	---

(✓) 

2	1
---	---

(G) 

2	3
---	---

(B) 

1	2
---	---

(D) 

1	4
---	---

(✓) 

2	2
---	---

(H) 

2	4
---	---

(I) *None of the above*

### NOTES:

Firstly, while PostgreSQL can compare DATE in a logical way, even comparing this as string is sufficient given the way our dates are written. So exclude the last three rows.

Then, group this by **bday** and **bfloor**, you get the following grouping with **bday** moved next to **bfloor**

Bookings			
bfloor	bday	bnum	beid
1	2023-04-29	1	1
1	2023-04-29	3	3
1	2023-04-29	4	2
1	2023-04-29	2	1
2	2023-04-29	2	2
1	2023-04-28	1	3
2	2023-04-28	1	3
2	2023-04-28	2	3
2	2023-04-28	3	2

Lastly, simply compute the number of ***distinct*** beid for each grouping. This gives us:

Bookings	
bfloor	count
1	3
2	1
1	1
2	2

## 2 Stored Procedures

Consider the following Scores table. The data types of `sid` and `name` are `TEXT` while the data type of `score` is `INT`.

sid	name	score
s1	Alice	71
s2	Bob	78
s3	Cathy	84
s4	David	89
s5	Eric	93

**2.1** (6 points) Consider the `test_func` function below.

```
CREATE OR REPLACE FUNCTION test_func()
RETURNS INT AS $func$
DECLARE
    curs CURSOR for (SELECT * FROM Scores ORDER BY score desc);
    r RECORD;
    score1 INT;
    score2 INT;
    gap1 INT;
    gap2 INT;
BEGIN
    score1 := -1;
    gap2 := -1;

    OPEN curs;
    LOOP
        FETCH curs INTO r;
        EXIT WHEN NOT FOUND;

        IF (score1 = -1) THEN
            score1 := r.score;
        ELSE
            gap1 := score1 - r.score;
            IF (gap1 > gap2) THEN
                gap2 := gap1;
                score2 := score1;
            END IF;
            score1 := r.score;
        END IF;
    END LOOP;
    CLOSE curs;
    // continue on the next page
```

```

IF (gap2 > -1) THEN
    RETURN score2;
ELSE
    RETURN -1;
END IF;
END;
$func$ LANGUAGE plpgsql;

```

Suppose that we execute the following query. What will be the query result?

```
SELECT * FROM test_func();
```

- (A) 4                      (C) 6                      (E) 8                      (✓) 78                      (I) 89  
 (B) 5                      (D) 7                      (F) 71                      (H) 84                      (J) 93  
 (K) *None of the above*

#### NOTES:

The function is finding the largest gap when sorted in descending order.

Scores			
sid	name	score	gap
s5	Eric	93	4
s4	David	89	5
s3	Cathy	84	6
s2	Bob	78	7
s1	Alice	71	-

But it does not return the gap, it returns the larger score from the gap. Notice that the gap is between 78 and 71. So it returns 78.

3 marks when answering either 71 or 84.

**2.2** (6 points) Consider the `scores_check_func` function below.

```

CREATE OR REPLACE FUNCTION scores_check_func()
RETURNS TRIGGER AS $func$
DECLARE
    fail_cnt INT;
BEGIN
    SELECT COUNT(*) INTO fail_cnt
    FROM Scores
    where score < 60;

    IF (fail_cnt > 3) THEN

```



```

        RAISE EXCEPTION 'You failed too many students!';
    END IF;

    RETURN NEW;
END;
$func$ LANGUAGE plpgsql;

```

Suppose that we create a trigger `scores_check_trigger` based on the above function, and then execute the following transaction.

```

BEGIN
    INSERT INTO Scores values ('s6', 'Fred', 59);
    INSERT INTO Scores values ('s7', 'Gigi', 58);
    INSERT INTO Scores values ('s8', 'Helen', 57);
    INSERT INTO Scores values ('s9', 'Ivan', 56);
COMMIT;

```

Assumes that `Scores` has **nine (9)** tuples after the transaction is executed. In that case, which of the following definitions of `scores_check_trigger` are possible?

- (✓) CREATE TRIGGER `scores_check_trigger`  
 BEFORE INSERT ON `Scores`  
 FOR EACH ROW  
 EXECUTE FUNCTION `scores_check_func()`;
- (B) CREATE TRIGGER `scores_check_trigger`  
 AFTER INSERT ON `Scores`  
 FOR EACH ROW  
 EXECUTE FUNCTION `scores_check_func()`;
- (C) CREATE CONSTRAINT TRIGGER `scores_check_trigger`  
 BEFORE INSERT ON `Scores`  
 DEFERRABLE INITIALLY IMMEDIATE  
 FOR EACH ROW  
 EXECUTE FUNCTION `scores_check_func()`;
- (D) CREATE CONSTRAINT TRIGGER `scores_check_trigger`  
 AFTER INSERT ON `Scores`  
 DEFERRABLE INITIALLY IMMEDIATE  
 FOR EACH ROW  
 EXECUTE FUNCTION `scores_check_func()`;
- (E) CREATE CONSTRAINT TRIGGER `scores_check_trigger`  
 BEFORE INSERT ON `Scores`  
 DEFERRABLE INITIALLY DEFERRED  
 FOR EACH ROW  
 EXECUTE FUNCTION `scores_check_func()`;
- (F) CREATE CONSTRAINT TRIGGER `scores_check_trigger`  
 AFTER INSERT ON `Scores`  
 DEFERRABLE INITIALLY DEFERRED

```
FOR EACH ROW  
EXECUTE FUNCTION scores_check_func();
```

(G) *None of the above*

**NOTES:**

Deferrable constraints can only be used on **AFTER** trigger. Since deferrable **BEFORE** trigger is not even a possible definition, it cannot be part of the answer. However, we do need to have a **BEFORE** trigger and, in fact, the only **BEFORE** trigger that we have is the only answer.

### 3 Functional Dependencies

**3.1** (6 points) Consider the following relation  $R(A, B, C, D, E, F)$  with the following set of functional dependencies

$$\Sigma = \{\{F\} \rightarrow \{B\}, \{D\} \rightarrow \{A\}, \{C\} \rightarrow \{F\}, \{E, F\} \rightarrow \{C\}, \\ \{A, F\} \rightarrow \{E\}, \{A, B\} \rightarrow \{F\}\}$$

Suppose that we decompose  $R$  into  $R1(A, B, C, D)$  and  $R2(C, D, E, F)$ .

(a) (2 points) Is this a lossless-join decomposition?

☒ True

☐ (B) False

(b) (4 points) Is this a dependency-preserving decomposition? If not, please select all FDs that are **NOT** preserved.

☒  $\{F\} \rightarrow \{B\}$

☒  $\{A, F\} \rightarrow \{E\}$

☐ (B)  $\{D\} \rightarrow \{A\}$

☐ (E)  $\{A, F\} \rightarrow \{B\}$

☐ (C)  $\{E, F\} \rightarrow \{C\}$

☐ (F)  $\{A, B\} \rightarrow \{F\}$

☐ (G) *None of the above*

**3.2** (7 points) Consider the following relation  $R(A, B, C, D, E, F)$  with the following set of functional dependencies

$$\Sigma = \{\{D\} \rightarrow \{E\}, \{C\} \rightarrow \{A\}, \{C, F\} \rightarrow \{E\}, \{A, F\} \rightarrow \{C\}, \\ \{C, E\} \rightarrow \{D\}, \{B, E\} \rightarrow \{A\}, \{A, D\} \rightarrow \{B\}, \{B, C\} \rightarrow \{F\}\}$$

Suppose that we decompose  $R$  into  $R1(A, B, C, D)$  and  $R2(C, D, E, F)$ .

(a) (2 points) Is this a lossless-join decomposition?

☒ True

☐ (B) False

(b) (5 points) Is this a dependency-preserving decomposition? If not, please select all FDs that are **NOT** preserved.

☐ (A)  $\{D\} \rightarrow \{E\}$

☐ (E)  $\{C, E\} \rightarrow \{D\}$

☐ (B)  $\{C\} \rightarrow \{A\}$

☒  $\{B, E\} \rightarrow \{A\}$

☐ (C)  $\{C, F\} \rightarrow \{E\}$

☐ (G)  $\{A, D\} \rightarrow \{B\}$

☒  $\{A, F\} \rightarrow \{C\}$

☐ (H)  $\{B, C\} \rightarrow \{F\}$

☐ (I) *None of the above*

## 4 Normal Forms

4.1 (15 points) Consider the following relation  $R(A, B, C, D, E)$  with the following set of functional dependencies

$$\Sigma = \{\{C\} \rightarrow \{E\}, \{B\} \rightarrow \{E\}, \{B\} \rightarrow \{A\}, \{C, E\} \rightarrow \{D\}, \\ \{B, E\} \rightarrow \{C\}, \{A, D\} \rightarrow \{B\}, \{A, D\} \rightarrow \{C\}\}$$

(a) (4 points) Write ALL the keys of  $R$  with respect to  $\Sigma$ . (*choices in Exemplify*)

**NOTES:**

The keys are

- $\{B\}$
- $\{A, C\}$
- $\{A, D\}$

(b) (3 points) Is  $R$  in BCNF with respect to  $\Sigma$ ? If not, write down all FDs in  $\Sigma$  that violates the BCNF requirements. (*choices in Exemplify*)

**NOTES:**

The violations are

- $\{C\} \rightarrow \{E\}$
- $\{C, E\} \rightarrow \{D\}$

- (c) (5 points) Is  $R$  in BCNF with respect to  $\Sigma$ ? If not, apply the BCNF decomposition algorithm (*introduced in CS2102 lectures*) on  $R$ , and select the relations in the final result of your BCNF decomposition. (*choices in Exemplify*)

**NOTES:**

The decomposed schema are

- $R1(A, B, C)$
- $R2(C, D, E)$

$R3(C, E)$  and  $R3(D, E)$  are also possible decomposition when using the algorithm because we did not require subsumed relation to be removed from the result. No other relations are reachable using only the algorithm.

- (d) (3 points) Is  $R$  in 3NF with respect to  $\Sigma$ ? If not, write down all FDs in  $\Sigma$  that violates the 3NF requirements. (*choices in Exemplify*)

**NOTES:**

The violations are

- $\{C\} \rightarrow \{E\}$

**4.2** (5 points) Consider the following relation  $R(A, B, C, D, E, F)$  with the following set of functional dependencies

$$\Sigma = \{\{F\} \rightarrow \{D\}, \{B, F\} \rightarrow \{C\}, \{B, C, D, F\} \rightarrow \{E\}, \{D\} \rightarrow \{B\}, \\ \{C, D\} \rightarrow \{A\}, \{B, D\} \rightarrow \{C\}, \{A, D\} \rightarrow \{C\}, \{B\} \rightarrow \{A\}\}$$

Derive a minimal basis of  $\Sigma$ . (*choices in Exemplify*)

**NOTES:**

The minimal basis is the following set

- $\{B\} \rightarrow \{A\}$
- $\{D\} \rightarrow \{B\}$
- $\{D\} \rightarrow \{C\}$
- $\{F\} \rightarrow \{D\}$
- $\{F\} \rightarrow \{E\}$

These properties need to be satisfied:

- Must be a cover (*i.e., must be equivalent to the original set of FD*).
- The right-hand side must be *singular*.
- No redundant attributes.
- No redundant functional dependencies.

**4.3** (5 points) Consider an arbitrary relation  $R(A, B, C, D, E, F)$ . Suppose that  $R$  is not in 3NF with respect to some unknown set of functional dependencies  $\Sigma$ . In that case, what is the maximum number of keys that  $R$  can have? (*choices in Exemplify*)

**NOTES:**

To derive this, note that one of the attributes cannot be part of a key. In that case, we can work with 5 attributes.

Note that if we include a set of attributes as a key, then all superset of this cannot be a key. As such, we want to keep the number of attributes to be the same in all keys. This gives us 5 choices for the size. To find the number of set with this size, we use binomial coefficient:  ${}^nC_k$ . We have 5 choices:

- ${}^5C_1$ : 5
- ${}^5C_2$ : 10
- ${}^5C_3$ : 10
- ${}^5C_4$ : 5
- ${}^5C_5$ : 1

So the maximum is 10.