

CS2102: Database Systems

Lecture 2 — SQL (Part 1)

Course Logistics

- Project registration
 - Canvas Groups with self sign-up
 - Canvas → CS2102 → People → Project (tab) → Project 1-125
 - Group size: 4 (members do not have to be in the same tutorial)
 - Use Canvas Discussion to look for members or team (random assignment to groups after Friday, 17:00)
- Tutorials
 - Appeals regarding allocation to be done on CourseReg
 - Mandatory attendance (you can skip up to 2 tutorials without penalty)
 - We expect students to come prepared (Check out the questions before coming to the tutorial!)

Quick Recap: Relational DBMS (RDBMS)

- RDBMS = DBMS + Relational Model
 - Unified representation of all data as relations (tables)
 - Integrity constraints to specify restrictions on what constitutes correct/valid data
 - Transactions with ACID properties to guarantee integrity of the data
 - Levels of abstraction for data independence

A Relational Model of Data for Large Shared Data Banks

E. F. CODD

IBM Research Laboratory, San Jose, California

Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation). A prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed. Changes in data representation will often be needed as a result of changes in query, update, and report traffic and natural growth in the types of stored information.

Table "Employees"

| id | name | dob | salary |
|-----------|-------------|------------|---------------|
| 1 | Alice | 10-08-1988 | 7,500 |
| 2 | Bob | 06-11-2001 | 4,800 |
| 3 | Carol | 25-02-1995 | 5,500 |

Quick Clarifications

- Terminology: "key" vs. "candidate key"
 - Same concept; terms used interchangeably
("candidate" highlights that there might be more than one key)
 - Additionally: (candidate) keys cannot be *null*
(otherwise they could not serve as chosen primary key)

Overview

- **SQL — overview**
 - History and usages
 - SQL language groups
- **Creating a database with SQL**
 - Basic DDL & DML commands
 - Defining integrity constraints
 - Advanced: deferrable constraints
- **Modifying a database with SQL**
 - Basic DDL commands

SQL — Structured Query Language

- De-facto standard language to "talk" to a RDBMS: **SQL**
 - Developed Donald D. Chamberlin and Raymond F. Boyce (IBM Research, 1974)
 - Originally called SEQUEL (Structured English Query Language)
 - SQL is not a general-purpose language (such as Python, Java, C++, etc.) but a **domain-specific language**
 - SQL is a **declarative language**: focus on *what* to compute, not on *how* to compute
- **SQL Standard**
 - First standard: SQL-86; most recent standard: SQL-2019 (new standard every ~3-5 years)
 - New standards introduce new language concepts (e.g., support new features of RDBMS)
 - Many RDBMS add their own "flavor" to SQL

Using SQL

- Interactive SQL: directly writing SQL statements to an interface
 - Command line interface
e.g., PostgreSQL's **psql** [1]
 - Graphical user interface
e.g., PostgreSQL's **pgAdmin** [2]

```
List of relations
Schema | Name   | Type  | Owner
-----+-----+-----+-----
public | aircrafts | table | postgres
public | airports  | table | postgres
public | countries | table | postgres
public | flightcodes | table | postgres
public | flights   | table | postgres
(5 rows)

flightsdb=# SELECT code, name FROM aircrafts LIMIT 3;
code | name
----+---
75D | Boeing 757-200
737 | Boeing 737
E75 | Embraer 175 (short wing)
(3 rows)

flightsdb=#
```

The screenshot shows the pgAdmin 4 interface. On the left, the 'Browser' panel displays a tree view of database objects under 'Servers (1) > vdw-com2 > Databases (4) > flightsdb'. The 'flightsdb' node is expanded, showing 'Tables (1)', 'Views (0)', 'Functions (0)', 'Procedures (0)', 'Triggers (0)', 'Extensions (0)', 'Foreign Data Wrappers (0)', 'Languages (0)', 'Publications (0)', 'Schemas (1)', and 'Subscriptions (0)'. The 'Tables (1)' item is highlighted with a blue selection bar. On the right, the 'Query Editor' tab is active, showing the following SQL query:

```
1 SELECT code, name
2 FROM aircrafts
3 LIMIT 3;
```

The 'Data Output' tab is selected, displaying the results of the query:

| code | name |
|------|--------------------------|
| 75D | Boeing 757-200 |
| 737 | Boeing 737 |
| E75 | Embraer 175 (short wing) |

[1] <https://www.postgresql.org/docs/current/static/app-psql.html>

[2] <https://www.pgadmin.org/>

Using SQL

- Non-interactive
 - SQL statements are included in an application written in a host language
 - Two basic approaches to include SQL in host languages: SLI & CLI
- Statement Level Interface (SLI)
 - Application is a mixture of host language statements and SQL statements
 - Examples: Embedded SQL, Dynamic SQL
- Call Level Interface (CLI)
 - Application is completely written in host language
 - SQL statements are strings passed as arguments to host language procedures or libraries
 - Examples: ODBC (Open DataBase Connectivity), JDBC (Java DataBase Connectivity)

Statement Level Interface (SLI) — Example

```
int main()
{
    EXEC SQL WHENEVER NOT FOUND DO BREAK;
    EXEC SQL BEGIN DECLARE SECTION;
    char v_code[32], v_name[32];
    EXEC SQL END DECLARE SECTION;

    // Connect to database
    EXEC SQL BEGIN DECLARE SECTION;
    const char *target = "flightsdb@localhost";
    const char *user = "postgres";
    const char *passwd = "_____";
    EXEC SQL END DECLARE SECTION;
    EXEC SQL CONNECT TO :target USER :user USING :passwd;

    // Declare cursor
    EXEC SQL DECLARE c CURSOR FOR
    SELECT code, name FROM aircrafts LIMIT 3;

    // Open cursor
    EXEC SQL OPEN c;

    // Loop through cursor and display results
    for(;;) {
        EXEC SQL FETCH NEXT FROM c INTO :v_code, :v_name;
        printf(">>> code: %s, name: %s\n", v_code, v_name);
    }

    // Cleanup (close cursor, commit, disconnect)
    EXEC SQL CLOSE c;
    EXEC SQL COMMIT;
    EXEC SQL DISCONNECT;

    return 0;
}
```

```
#!/bin/bash

# Run ecpg preprocessor to convert C program with embedded SQL statements
# to normal C code; replaces the SQL invocations with special function calls.
ecpg flightsdb.pgc

# Compile generated C code; requires to include all header files the compiler
# needs to understand the special function calls (files come with PostgreSQL).
gcc -g -I/usr/include/postgresql -c flightsdb.c

# Build output to executable file; also needs access to the header files.
gcc -o flightsdb flightsdb.o -L/usr/include/postgresql -lecpq
```

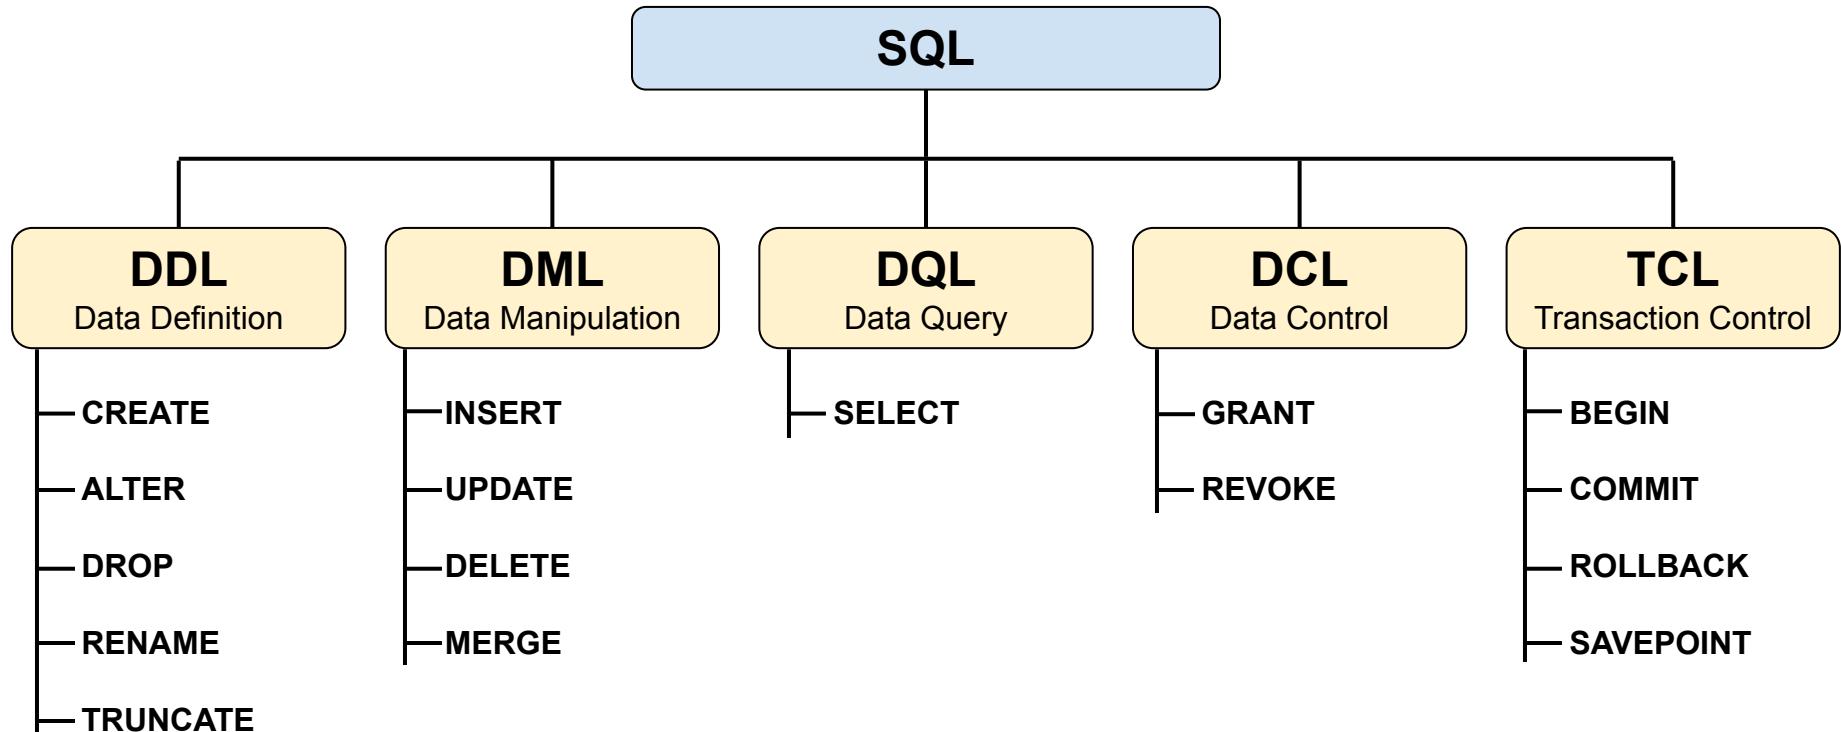
```
>>> code: 75D, name: Boeing 757-200
>>> code: 737, name: Boeing 737
>>> code: E75, name: Embraer 175 (short wing)
```

Call Level Interface (CLI) — Example

```
1 import psycopg2 # Host language library (here psycopg2 for Python)
2
3 # Connect to database
4 db = psycopg2.connect(host="localhost", database="flightsdb", user="postgres", password="████████")
5
6 # Create cursor
7 cursor = db.cursor()
8
9 # Open cursor by executing query (string parameter passed to execute() method)
10 cursor.execute("SELECT code, name FROM aircrafts LIMIT 3")
11
12 # Loop over all results until no next tuple is returned
13 while True:
14     row = cursor.fetchone()
15     if row is None:
16         break
17     print(row)
18
19 # Cleanup
20 cursor.close()
21 db.close()
```

```
('75D', 'Boeing 757-200')
('737', 'Boeing 737')
('E75', 'Embraer 175 (short wing)')
```

SQL — Types of Commands/Statements



Overview

- SQL — overview
 - History and usages
 - SQL language groups
- Creating a database with SQL
 - Basic DDL & DML commands
 - Defining integrity constraints
 - Advanced: deferrable constraints
- Modifying a database with SQL
 - Basic DDL commands

DDL — Creating Tables

- Basic syntax: definition of table name and attributes (with data types)

Employees (id: **integer**, name: **text**, age: **integer**, role: **text**)



```
CREATE TABLE Employees (
    id      INTEGER,
    name   VARCHAR(50),
    age    INTEGER,
    role   VARCHAR(50)
);
```

- Extended syntax: definition of additional data integrity constraints

Data Types (PostgreSQL)

- Basic data types

(supported by most RDBMS)

| | |
|------------------------|--|
| boolean | logical Boolean (true/false) |
| integer | signed four-byte integer |
| float8 | double precision floating-point number (8 bytes) |
| numeric [(p,s)] | exact numeric of selectable precision |
| char(n) | fixed-length character string |
| varchar(n) | variable-length character string |
| text | variable-length character string |
| date | calendar date (year, month, day) |
| timestamp | date and time |

- Many extended data types

- Document types: XML, JSON
- Spatial types: point, line, polygon, circle, box, path
- Special types: money/currency, MAC/IP address

- Definition user-defined types (UDTs)

DML — Inserting Data (Basic Examples)

```
CREATE TABLE Employees (
    id    INTEGER,
    name  VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50)
);
```

- Example: Inserting 3 employees

- Specifying all attribute values

```
INSERT INTO Employees VALUES (101, 'Sarah', 25, 'dev');
```

- Specifying selected attribute values

```
INSERT INTO Employees (id, name) VALUES (102, 'Judy'), (103, 'Max');
```



Employees

| id | name | age | role |
|-----|-------|------|------|
| 101 | Sarah | 25 | dev |
| 102 | Judy | null | null |
| 103 | Max | null | null |

DML — Inserting Data (Basic Examples)

- Example: Inserting 3 employees

- Specifying all attribute values

```
INSERT INTO Employees VALUES (101, 'Sarah', 25, 'dev');
```

- Specifying selected attribute values

```
INSERT INTO Employees (id, name) VALUES (102, 'Judy'), (103, 'Max');
```



Employees

| id | name | age | role |
|-----|-------|------|--------------|
| 101 | Sarah | 25 | dev |
| 102 | Judy | null | sales |
| 103 | Max | null | sales |

```
CREATE TABLE Employees (
    id      INTEGER,
    name   VARCHAR(50),
    age    INTEGER,
    role   VARCHAR(50) DEFAULT 'sales'
);
```

DML — Deleting Data (Basic Examples)

-- Delete all tuples

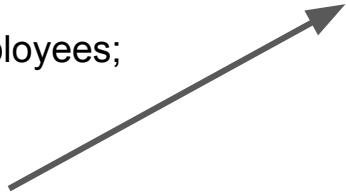
DELETE FROM Employees;

Employees

| id | name | age | role |
|----|------|-----|------|
|----|------|-----|------|

Employees

| id | name | age | role |
|-----|-------|------|-------|
| 101 | Sarah | 25 | dev |
| 102 | Judy | null | sales |
| 103 | Max | null | sales |

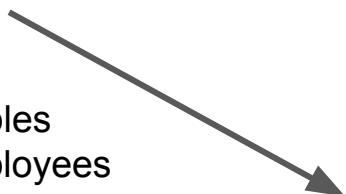


-- Delete selected tuples

DELETE FROM Employees
WHERE role = 'dev';

Employees

| id | name | age | role |
|-----|------|------|-------|
| 102 | Judy | null | sales |
| 103 | Max | null | sales |



DML — Updating Data (Basic Examples)

Employees

| id | name | age | role |
|-----|-------|------|-------|
| 101 | Sarah | 25 | dev |
| 102 | Judy | null | sales |
| 103 | Max | null | sales |

-- Sarah's birthday
UPDATE Employees
SET age = age + 1
WHERE name = 'Sarah';

Employees

| id | name | age | role |
|-----|-------|-----------|-------|
| 101 | Sarah | 26 | dev |
| 102 | Judy | null | sales |
| 103 | Max | null | sales |



-- New privacy law
UPDATE Employees
SET age = 0;

Employees

| id | name | age | role |
|-----|-------|----------|-------|
| 101 | Sarah | 0 | dev |
| 102 | Judy | 0 | sales |
| 103 | Max | 0 | sales |

Employees

| id | name | age | role |
|-----|--------------|------|--------------|
| 101 | SARAH | 25 | DEV |
| 102 | JUDY | null | SALES |
| 103 | MAX | null | SALES |

Overview

- SQL — overview
 - History and usages
 - SQL language groups
- Creating a database with SQL
 - Basic DDL & DML commands
 - **Defining integrity constraints**
 - Advanced: deferrable constraints
- Modifying a database with SQL
 - Basic DDL commands

Prerequisite — Handling *null* Values

- Recall: rules of handling *null* values

- The result of a comparison operation with *null* is *unknown*
- The result of an arithmetic operation with *null* is *null*

Assume that value of x is *null*

$x < 2020$ → *unknown*

$x = \text{null}$ → *unknown*

$x <> \text{null}$ → *unknown*

$x + 5$ → *null*

→ Three-valued logic: **true**, **false**, **unknown**

- Questions

- How to check if a value is equal to null?
- How to treat null values as ordinary values for comparison?

}

Important for writing SQL queries & checking integrity constraints!

IS (NOT) NULL Comparison Predicate

- Check if a value is equal to null (since "`=`" would return unknown)
 - If x is a *null* value \rightarrow " $x \text{ IS NULL}$ " evaluates to **true**
 - If x is a non-*null* value \rightarrow " $x \text{ IS NULL}$ " evaluates to **false**
- Equivalence
 - " $x \text{ IS NOT NULL}$ " is equivalent to "**NOT** ($x \text{ IS NULL}$)"

} vice versa for " $x \text{ IS NOT NULL}$ "

| x | y |
|-------------|-------------|
| 1 | 1 |
| 1 | 2 |
| <i>null</i> | 1 |
| <i>null</i> | <i>null</i> |



| x | y | x IS NULL | y IS NULL |
|-------------|-------------|--------------|--------------|
| 1 | 1 | <i>false</i> | <i>false</i> |
| 1 | 2 | <i>false</i> | <i>false</i> |
| <i>null</i> | 1 | <i>true</i> | <i>false</i> |
| <i>null</i> | <i>null</i> | <i>true</i> | <i>true</i> |

IS (NOT) NOT DISTINCT Comparison Predicate

- "x **IS DISTINCT FROM** y"
 - equivalent to "x <> y" if x and y are non-null values
 - if x and y both null → evaluates to **false**
 - if only one value is null → evaluates to **true**
- Equivalence
 - "x **IS NOT DISTINCT FROM** y" is equivalent to "**NOT** (x **IS DISTINCT FROM** y)"

| x | y | x <> y | x IS DISTINCT FROM y |
|------|------|--------|----------------------|
| 1 | 1 | FALSE | FALSE |
| 1 | 2 | TRUE | TRUE |
| null | 1 | null* | TRUE |
| null | null | null* | FALSE |

* PostgreSQL represents "unknown" using null

DDL — Data Integrity Constraints: Overview

- **Types of Constraints** ("named" or "unnamed")

- Not-null constraints
- Unique constraints
- Primary key constraints
- Foreign key constraints
- General constraints



A constraint is violated
if it evaluates to **false**

- **Constraint specifications** (difference "where" a constraint is specified)

- Column constraint: applies to single column, specified at column definition
- Table constraint: applies to one or more columns, specified after all column definitions
- Assertion: stand-alone command (**create assertion** ...)

Not-Null Constraints

- Example: the id or name of an employee cannot be *null*

unnamed constraint (name assigned by DBMS)

```
CREATE TABLE Employees (
    id    INTEGER NOT NULL,
    name  VARCHAR(50) NOT NULL,
    age   INTEGER,
    role  VARCHAR(50),
);
```

named constraint (easier bookkeeping)

```
CREATE TABLE Employees (
    id    VARCHAR(50) CONSTRAINT nn_id NOT NULL,
    name  VARCHAR(50) CONSTRAINT nn_name NOT NULL,
    age   INTEGER,
    role  VARCHAR(50),
);
```

- Not-null constraint violation:
 - There exists a tuple $t \in \text{Employees}$ where " $t.id \text{ IS NOT NULL}$ " evaluates to **false**
 - There exists a tuple $t \in \text{Employees}$ where " $t.name \text{ IS NOT NULL}$ " evaluates to **false**

Unique Constraints

- Example: the id of an employee must be unique

unnamed column constraint

```
CREATE TABLE Employees (
    id   INTEGER UNIQUE,
    name VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50)
);
```

named column constraint

```
CREATE TABLE Employees (
    id   INTEGER CONSTRAINT u_id UNIQUE,
    name VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50)
);
```

unnamed table constraint

```
CREATE TABLE Employees (
    id   INTEGER,
    name VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50),
    UNIQUE (id)
);
```

named table constraint

```
CREATE TABLE Employees (
    id   INTEGER,
    name VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50),
    CONSTRAINT u_id UNIQUE (id)
);
```

Unique Constraints

- Unique constraint for more than one attribute / column
 - Can only be specified using table constraints
 - Example: Each pair of employee name and project name must be unique

Teams (eid: **integer**, pname: **text**, hours: **integer**)

unnamed table constraint

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    UNIQUE (eid, pname)
);
```

named table constraint

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    CONSTRAINT u_allocation UNIQUE (eid, pname)
);
```

Unique Constraints

Quick Quiz: Is the unique constraint of table "Teams" violated in the example below?

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    UNIQUE (eid, pname)
);
```

Teams

| eid | pname | hours |
|-----|----------|-------|
| 101 | BigAI | 10 |
| 105 | BigAI | 5 |
| 102 | GlobalDB | 20 |
| 101 | null | null |
| 101 | null | null |
| 103 | CoreOS | 40 |
| 109 | CoreOS | null |

- Unique constraint violation

- For any two tuples $t_i, t_k \in \text{Teams}$:
- " $(t_i.\text{eid} <> t_k.\text{eid}) \text{ or } (t_i.\text{pname} <> t_k.\text{pname})$ " evaluates to **false**

Primary Key Constraints

Quick Quiz: What is the difference between using "primary key" and "unique not null"?

- Quick recap: primary key
 - Selected key uniquely identifying tuples in a table
 - Prime attributes (i.e. attributes of primary key) cannot be null

Employees (id: integer, name: **text**, age: **integer**, role: **text**)

```
CREATE TABLE Employees (
    id    INTEGER PRIMARY KEY,
    name  VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50)
);
```

```
CREATE TABLE Employees (
    id    INTEGER UNIQUE NOT NULL,
    name  VARCHAR(50),
    age   INTEGER,
    role  VARCHAR(50)
);
```

same effect

Primary Key Constraints

- Primary key constraint for more than one attribute / column

Teams (eid: integer, pname: text, hours: integer)

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    PRIMARY KEY (eid, pname)
);
```

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    CONSTRAINT pk_allocation PRIMARY KEY (eid, pname)
);
```

Sidenote

- Specification of constraints — basic rules
 - All constraints can be specified "named" or "unnamed"
(unnamed constraints still get named by the DBMS in a meaningful way; names can be looked up)
 - All column constraints can be specified as table constraints
(exception: "not null" only possible as column constraint)
 - Table constraints referring to a single column can be specified as column constraint
 - Column and table constraints can be combined (even w.r.t. to the same column)

```
CREATE TABLE Employees (
    id      INTEGER NOT NULL,
    name   VARCHAR(50),
    age    INTEGER,
    role   VARCHAR(50),
    UNIQUE (id)
);
```

Foreign Key Constraints

- Quick recap: foreign key constraint
 - Subset of attributes of relation A if it refers to the primary key in a relation B

Employees (id: integer, name: text, age: integer, role: text, hired: date)

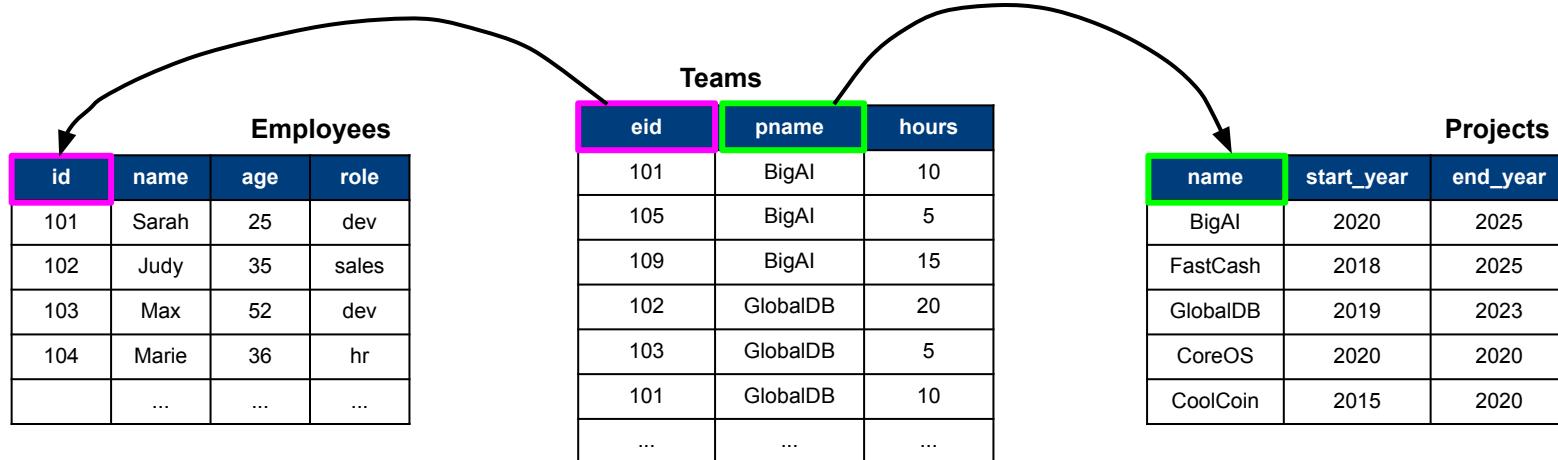
Teams (eid: integer, pname: text, hours: integer)

Projects (name: text, start_year: integer, end_year: integer)

Foreign key constraints

Teams.eid → Employees.id

Teams.pname → Projects.name



Foreign Key Constraints

```
CREATE TABLE Employees (
    id      INTEGER PRIMARY KEY,
    name   VARCHAR(50),
    age    INTEGER,
    role   VARCHAR(50)
);
```

```
CREATE TABLE Projects (
    name    VARCHAR(50) PRIMARY KEY,
    start_year  INTEGER,
    end_year   INTEGER
);
```

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    PRIMARY KEY (ename, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id),
    FOREIGN KEY (pname) REFERENCES Projects (name)
);
```

Foreign Key Constraints — Violations

- Quick recap: each foreign key in referencing relation must
 - appear as primary key in referenced relation OR
 - be a *null* value

The diagram illustrates a foreign key relationship between the 'Teams' table and the 'Projects' table. An arrow points from the 'pname' column in the 'Teams' table to the 'name' column in the 'Projects' table, indicating that 'pname' is a foreign key referencing 'name'.

| Teams | | |
|-------|----------|-------|
| eid | pname | hours |
| 101 | BigAI | 10 |
| 105 | BigAI | 5 |
| 109 | BigAI | 15 |
| 102 | GlobalDB | 20 |
| 103 | GlobalDB | 5 |
| 101 | GlobalDB | 10 |
| ... | ... | ... |

| Projects | | |
|----------|------------|----------|
| name | start_year | end_year |
| BigAI | 2020 | 2025 |
| FastCash | 2018 | 2025 |
| GlobalDB | 2019 | 2023 |
| CoreOS | 2020 | 2020 |
| CoolCoin | 2015 | 2020 |

Questions:

- What happens if the first tuple in "Project" should be deleted?
- What if the project "BigAI" should be renamed to "SmartAI"?

Note: Trying to insert or update a tuple in "Teams" with a new project name that is not in "Project" will always violate the foreign constraint.

Foreign Key Constraints — Violations

- Extend syntax to specify behavior when data in referenced table changes
 - Specify action in case of violation of a foreign key constraint
 - ON DELETE/UPDATE <action>** to distinguish action w.r.t. to a delete or update in referenced table
 - Both specifications are optional

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    PRIMARY KEY (ename, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE <action> ON UPDATE <action>,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON DELETE <action> ON UPDATE <action>
);
```

Foreign Key Constraints — Violations

- Possible actions for **on delete** and **on update**

NO ACTION rejects delete/update if it violates constraint (default value)

RESTRICT similar to "no action" except that check of constraint cannot be deferred
(deferrable constraints are discussed in a bit)

CASCADE propagates delete/update to referencing tuples

SET DEFAULT updates foreign keys of referencing tuples to some default value
(important: default value must be a primary key in the referenced table!)

SET NULL updates foreign keys of referencing tuples to *null*
(important: corresponding column must be allowed to contain *null* values!)

Foreign Key Constraints

Quick Quiz: The SQL command below is correct but what will cause problems. Why?

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE NO ACTION ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON DELETE SET NULL ON UPDATE CASCADE
);
```

optional since it is the default action



- Effects on handling violations of foreign key constraints
 - Updates of "Employees.id" and "Projects.name" are propagated to affected tuples in "Teams"
 - Deleting a project will set "Teams.pname" to *null* for employees working on that project
 - Deleting an employee will raise an error if that employee is still assigned to a team

Foreign Key Constraints — Example

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON UPDATE CASCADE
);
```

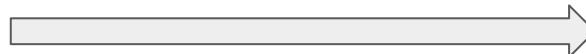
Projects

| name | start_year | end_year |
|----------|------------|----------|
| BigAI | 2020 | 2025 |
| FastCash | 2018 | 2025 |
| ... | ... | ... |

Teams

| eid | pname | hours |
|-----|----------|-------|
| 101 | BigAI | 10 |
| 105 | BigAI | 5 |
| 109 | BigAI | 15 |
| 102 | GlobalDB | 20 |
| ... | ... | ... |

UPDATE Projects
SET name = 'SmartAI'
WHERE name = 'BigAI';



Projects

| name | start_year | end_year |
|----------|------------|----------|
| SmartAI | 2020 | 2025 |
| FastCash | 2018 | 2025 |
| ... | ... | ... |

Teams

| eid | pname | hours |
|-----|----------|-------|
| 101 | SmartAI | 10 |
| 105 | SmartAI | 5 |
| 109 | SmartAI | 15 |
| 102 | GlobalDB | 20 |
| ... | ... | ... |

Foreign Key Constraints — Example

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname    VARCHAR(100) DEFAULT 'FastCash', -- default value must be primary key in "Projects"!
    hours    INTEGER,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id) ON UPDATE CASCADE,
    FOREIGN KEY (pname) REFERENCES Projects (name) ON UPDATE CASCADE ON DELETE SET DEFAULT
);
```

Projects

| name | start_year | end_year |
|----------|------------|----------|
| BigAI | 2020 | 2025 |
| FastCash | 2018 | 2025 |
| ... | ... | ... |

Teams

| eid | pname | hours |
|-----|----------|-------|
| 101 | BigAI | 10 |
| 105 | BigAI | 5 |
| 109 | BigAI | 15 |
| 102 | GlobalDB | 20 |
| ... | ... | ... |

**DELETE FROM Projects
WHERE name = 'BigAI';**



Projects

| name | start_year | end_year |
|----------|------------|----------|
| FastCash | 2018 | 2025 |
| ... | ... | ... |

Teams

| eid | pname | hours |
|-----|----------|-------|
| 101 | FastCash | 10 |
| 105 | FastCash | 5 |
| 109 | FastCash | 15 |
| 102 | GlobalDB | 20 |
| ... | ... | ... |

Foreign Key Constraints

- Practical considerations
 - Specified constraints might not behave as expected (e.g., **SET NULL** issue with prime attributes)
 - Particularly **ON DELETE CASCADE** can have very bad consequences
 - **CASCADE** may significantly affect overall performance
- Careful design and specification of foreign key constraints is crucial!

Check Constraints

- **CHECK** constraint
 - Most basic general constraint (i.e., not a structural integrity constraint)
 - Allows to specify that column values must satisfy a Boolean expression
 - Scope: one table, single row
- Example: The hours an employee is allocated to a project must be > 0

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER CHECK (hours > 0),
    -- hours  INTEGER CONSTRAINT positive_hours CHECK (hours > 0),
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id),
    FOREIGN KEY (pname) REFERENCES Projects (name)
);
```

Check Constraints

- **CHECK** constraints can refer to multiple columns
 - Example: The start year of a project cannot be larger value than the end year

```
CREATE TABLE Projects (
    name      VARCHAR(50) PRIMARY KEY,
    start_year INTEGER,
    end_year   INTEGER,
    -- CHECK (start_year <= end_year),
    CONSTRAINT valid_lifetime CHECK (start_year <= end_year)
);
```

Check Constraints

- **CHECK** constraints can be arbitrarily complex Boolean expressions
 - Example: minimum hour requirements for different projects

```
CREATE TABLE Teams (
    eid      INTEGER,
    pname   VARCHAR(100),
    hours   INTEGER,
    PRIMARY KEY (eid, pname),
    FOREIGN KEY (eid) REFERENCES Employees (id),
    FOREIGN KEY (pname) REFERENCES Projects (name),
    CHECK (
        (pname = 'CoreOS' AND hours >= 30)
        OR
        (pname <> 'CoreOS' AND hours > 0)
    )
);
```

Assertions

- **CREATE ASSERTION** statement (since SQL-92)
 - Formulation of (almost) arbitrary constraints
 - Scope: multiple tables, multiple rows
 - Example: "Each project must have at least one team member being 30 or older"
 - Assertion in practice: various potential side effects and limitations, e.g.:
 - Assertions cannot modify the data
 - No proper error handling
 - Not linked to a specific table
(e.g., dropping a table does not affect assertion)
- Most RDBMS do not support assertions but **triggers** (more powerful alternative)

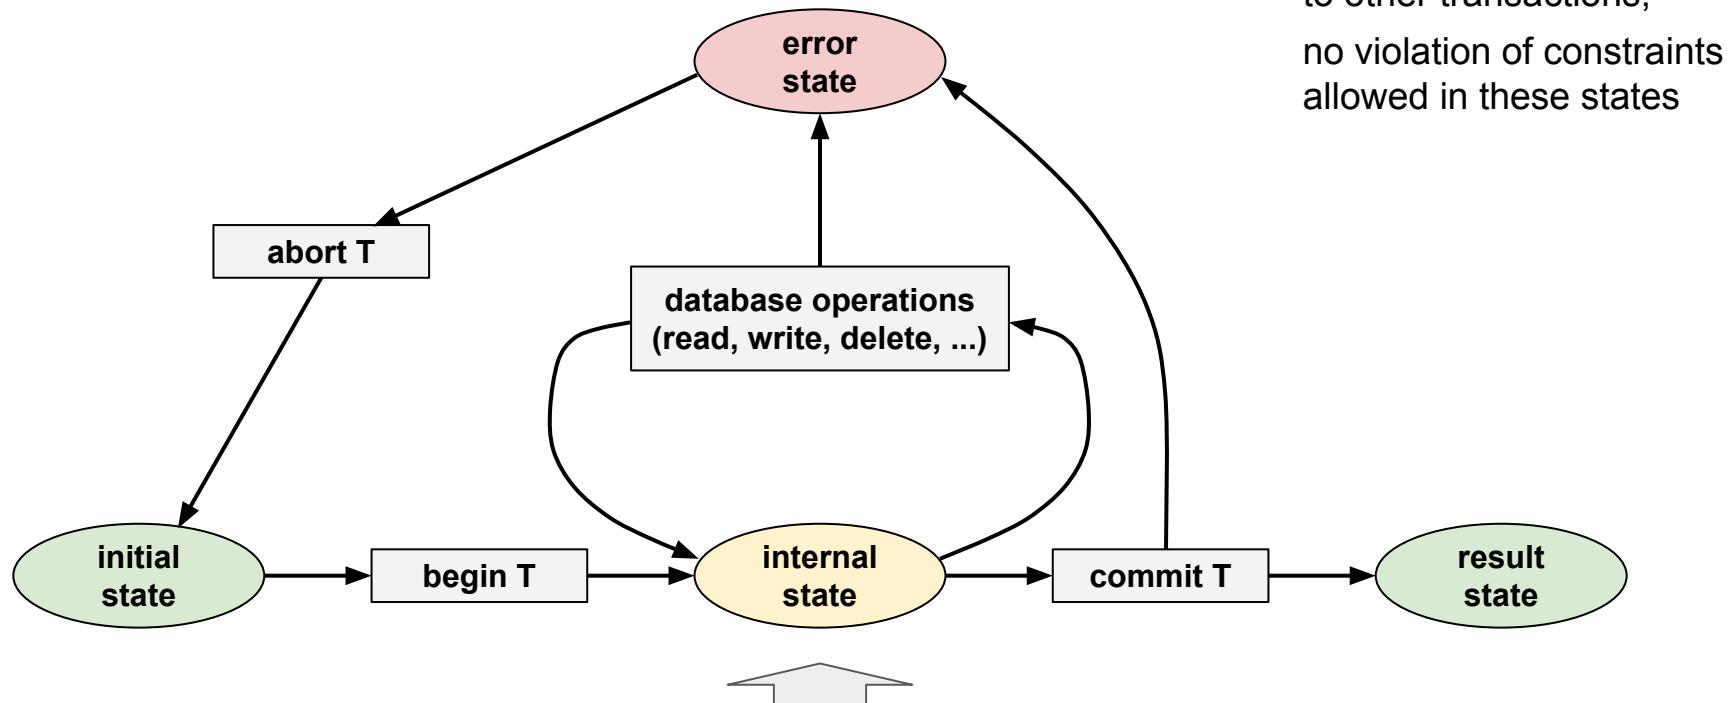
Overview

- SQL — overview
 - History and usages
 - SQL language groups
- Creating a database with SQL
 - Basic DDL & DML commands
 - Defining integrity constraints
 - Advanced: deferrable constraints
- Modifying a database with SQL
 - Basic DDL commands

Deferrable Constraints — Motivation

- Default behavior for constraints
 - Constraints are checked immediately at the end of SQL statement execution (even within a transaction containing multiple SQL statements)
 - A violation will cause the statement to be rolled back
- Relaxed constraint checks: **Deferrable Constraints**
 - Check can be deferred for some constraints to the end of a transaction
 - Available for: UNIQUE, PRIMARY KEY, FOREIGN KEY

Deferrable Constraints — Motivation



Deferrable constraints may (temporarily) be violated within the scope of a transaction

Deferrable Constraints — Example

- Motivating example without deferrable constraints

| Employees | | |
|-----------|-------|---------|
| id | name | manager |
| 101 | Sarah | null |
| 102 | Judy | 101 |
| 103 | Max | 102 |

```
CREATE TABLE Employees (
    id      INTEGER PRIMARY KEY,
    name    VARCHAR(50),
    manager INTEGER,
    CONSTRAINT manager_fkey FOREIGN KEY (manager) REFERENCES Employees (id)
        NOT DEFERRABLE -- default value (optional), check if constraint is immediate and cannot be changed
);
INSERT INTO Employees VALUES (101, 'Sarah', null), (102, 'Judy', 101), (103, 'Max', 102);

BEGIN;
DELETE FROM Employees WHERE id = 102;           -- Judy got fired → constraint violated → ABORT
UPDATE Employees SET manager = 101 WHERE id = 103; -- Max gets a new manager
COMMIT;
```

Deferrable Constraints — Example

Employees

| id | name | manager |
|-----|-------|---------|
| 101 | Sarah | null |
| 102 | Judy | 101 |
| 103 | Max | 102 |

```
CREATE TABLE Employees (
    id      INTEGER PRIMARY KEY,
    name    VARCHAR(50),
    manager INTEGER,
    CONSTRAINT manager_fkey FOREIGN KEY (manager) REFERENCES Employees (id)
        DEFERRABLE INITIALLY DEFERRED -- check of constraint deferred by default
);
INSERT INTO Employees VALUES (101, 'Sarah', null), (102, 'Judy', 101), (103, 'Max', 102);

BEGIN;
DELETE FROM Employees WHERE id = 102;           -- Judy got fired → constraint violated but not checked
UPDATE Employees SET manager = 101 WHERE id = 103; -- Max gets a new manager → constraint re-established
COMMIT;
```

Deferrable Constraints — Example

Employees

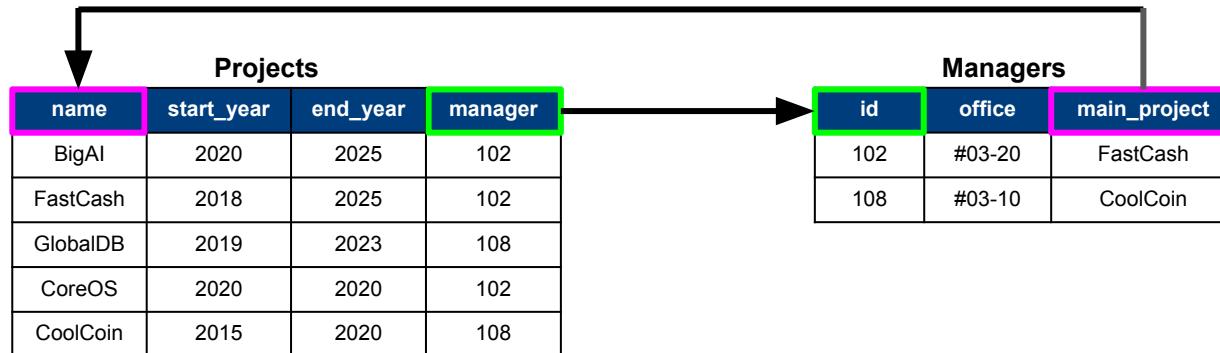
| id | name | manager |
|-----|-------|---------|
| 101 | Sarah | null |
| 102 | Judy | 101 |
| 103 | Max | 102 |

```
CREATE TABLE Employees (
    id      INTEGER PRIMARY KEY,
    name    VARCHAR(50),
    manager INTEGER,
    CONSTRAINT manager_fkey FOREIGN KEY (manager) REFERENCES Employees (id)
        DEFERRABLE INITIALLY IMMEDIATE -- check of constraint immediate by default, but can be changed
);
INSERT INTO Employees VALUES (101, 'Sarah', null), (102, 'Judy', 101), (103, 'Max', 102);

BEGIN;
SET CONSTRAINT manager_fkey DEFERRED;           -- Set check of constraint from "immediate" to "deferred"
DELETE FROM Employees WHERE id = 102;          -- Judy got fired → constraint violated buy not checked
UPDATE Employees SET manager = 101 WHERE id = 103; -- Max gets a new manager → constraint re-established
COMMIT;
```

Deferrable Constraints — Benefits

- No need to care about order of SQL statements within a transaction
- Allows for cyclic foreign key constraints



- Performance boost when constraint checks are bottleneck
 - Example: batch insert of large number of tuples

Deferrable Constraints — (Potential) Downsides

- Troubleshooting can be more difficult
- Data definitions no longer unambiguous
- Performance penalty when performing queries

Overview

- **SQL — overview**
 - History and usages
 - SQL language groups
- **Creating a database with SQL**
 - Basic DDL & DML commands
 - Defining integrity constraints
 - Advanced: deferrable constraints
- **Modifying a database with SQL**
 - Basic DDL commands

DDL — Modifying a Schema

- **ALTER TABLE** statements to modify an existing data definition
 - CREATE TABLE statements do not have to be final data definition
 - Common: adding/dropping column, adding dropping constraints, changing data types
- Examples: Change specification of a single column

```
ALTER TABLE Projects ALTER COLUMN name TYPE VARCHAR(200); -- change data type to VARCHAR(200)
```

```
ALTER TABLE Projects ALTER COLUMN start_year SET DEFAULT 2021; -- set default value of column "start_year"
```

```
ALTER TABLE Projects ALTER COLUMN start_year DROP DEFAULT; -- drop default value of column "start_year"
```

DDL — Modifying a Schema

- Examples: Adding and dropping columns

```
ALTER TABLE Projects ADD COLUMN budget NUMERIC DEFAULT 0.0; -- add new column with a default value
```

```
ALTER TABLE Projects DROP COLUMN budget; -- drop column from table
```

- Examples: Adding and dropping constraints

```
ALTER TABLE Teams ADD CONSTRAINT eid_fkey FOREIGN KEY (eid) REFERENCES Employees (id);  
-- add foreign key constraint
```

```
ALTER TABLE Teams DROP CONSTRAINT eid_fkey;  
-- drop foreign key constraint (name of constraint might be retrieved from metadata)
```

DDL — Drop Tables

- **DROP TABLE** to delete tables from database

- Without dependent objects (incl. foreign key constraints, views, etc.)

```
DROP TABLE Projects;
```

```
DROP TABLE IF EXISTS Projects; -- check first if table exists; avoids throwing an error
```

- With dependent objects (assume foreign key constraint Teams.pname→Projects.name)

```
DROP TABLE Projects; -- will throw an error because of foreign key constraint
```

```
DROP TABLE Projects CASCADE; -- will delete table "Projects" and foreign key constraint  
-- (will not delete table "Teams"!)
```

Summary

- SQL — *the standard language for RDBMS*
 - Different language groups: DDL, DML, DQL, DCL, TCL
- Focus in this lecture: DDL and DML
 - DDL: **CREATE TABLE, ALTER TABLE, DROP TABLE**
 - DML: **INSERT, UPDATE, DELETE**
- Key challenge: specification of integrity constraints
 - **NOT NULL, UNIQUE, PRIMARY KEY, FOREIGN KEY, CHECK**
 - Specification actions in case of foreign key constraint violations (**ON UPDATE/DELETE**)
 - Relaxed checks of violations with deferrable constraints

Quick Quiz Solutions

Quick Quiz (Slide 26)



Quick Quiz (Slide 27)



Quick Quiz (Slide 35)

