

**CS2106 Introduction to Operating Systems**  
**Lab 2 – Shell and Process Operations in Unix**  
**Release: Week of 8 September 2025**  
**Demo: Week of 15 September 2025**  
**Final Submission: Sunday 21 September 2025, 2359 hours**

## 1 Instruction

Some points to note about this lab:

- a. This lab should be run using the slurm cluster, accessible via xlog.comp.nus.edu.sg. Be sure to connect using the SoC VPN and NOT the NUS VPN.
- b. You may complete this lab on your own, or with ONE partner from any lab group. You do not need to partner with the same person with whom you completed Lab 1.
- c. There are two deliverables for this lab; a zip file that you can submit with your partner or individually IF you are doing the lab alone, and a practical demo that must be done individually.
- d. Only ONE copy of the submission zip file is to be submitted. Therefore, if you do the report with a partner, decide who should submit the report. DO NOT submit two copies.
- e. Use the enclosed AxxxxxY.docx answer book for your report, renaming it to the student ID of the submitter.
- f. Please indicate the student number, name, and group number of each person in the lab report.
- g. Create a ZIP file called AxxxxxY.zip with the following files (rename AxxxxxY to the student ID of the submitter):
  - The AxxxxxY.docx, appropriately renamed.
  - The myshell.c file
- h. Upload your ZIP file to Canvas by **SUNDAY, 21 September 2025 2359 hours**.
- i. Some extra time is given for submission, but once the folder closes, **NO SUBMISSIONS WILL BE ENTERTAINED AND YOU WILL RECEIVE 0 MARKS FOR THE LAB.**

## 2 Introduction

As programmers, command-line interpreters, also known as **shells**, are an important and ubiquitous part of our lives. A command line interpreter (or command prompt, or shell) allows the execution of multiple user commands with various number of arguments. The user inputs the commands one after another, and the commands are executed by the command line interpreter.

A shell is actually just another program implemented using the process-related system calls discussed in the lectures. In this lab, you are going to implement a simple shell with the following functionalities:

- Running commands in foreground and background
- Chaining commands
- Redirecting input, output, and error streams
- Terminating commands
- Managing the processes launched by this shell

The main purpose of this lab is to familiarize you with:

- advanced aspects of C programming,
- system calls,
- process operations in Unix-based operating system

This lab has a total of **three** exercises.

For exercise 1, only some simple functionalities are required. Take the opportunity to design your code in a modular and extensible way. You may want to look through the rest of the exercises before starting to code.

### Shell Implementation

The driver of the shell has been implemented for you in **driver.c**. The driver will read and tokenize the user commands for you, where tokens will be separated by spaces (whitespace or tabs).

We provide a structure **PCBTable** in **myshell.h** as follows:

```
struct PCBTable {
    pid_t pid;
    int status; // 4: Stopped, 3: Terminating, 2: Running, 1: exited
    int exitCode; // -1 not exit, else exit code status
};
```

where

- **pid** is the process id
- **status** indicates the status of the process ; (e.g., 2: Running, 1: Exited)
- **exitCode** indicate the exit status of process, -1 if still running

Please use the **PCBTable** to maintain the details of all the processes you fork. Please define a structure variable of type **PCBTable**, which can be an array, a linked list or other containers. You can assume there won't be more than 50 **PCBTables**.

You should implement the **three** functions in the file **myshell.c**:

- **my\_init()** is called when your shell starts up, before it starts accepting user commands. It should initialize your shell appropriately.
- **my\_process\_command()** is called in the main loop. It should handle every user command except the **quit** command. The function accepts two arguments:
  - **tokens** is an array of tokens with the last element of the array being **NULL**.
  - **size** is the size of that array, including the **NULL**.

You may want to print out the **tokens** array or look at the **driver.c** file for your own understanding.

- `my_quit()` is called when user inputs the `quit` command.

We have also provided some executable programs with various runtime behaviours in the `programs` folder to aid your testing later. You can also create your own programs and run them with your shell.

- `programs/goingtosleep`: Prints “Good night!” and then prints “Going to sleep” 10 times with a 2-second sleep between each print.
- `programs/lazy`: Wake up, echo “Good morning...” and loops for 2 minutes. It takes a while (at least 5 seconds) to respond to the SIGTERM signal.
- `programs/result`: Takes in a single number `x` and exits with a return status of `x`.
- `programs/showCmdArg`: Shows the command line arguments passed in by the user.
- `programs/Makefile`: To compile the above programs.

### Preventing Fork Bombs

During implementation, you might mistakenly call `fork` infinitely and ignite a `fork bomb`. Thus, we have provided a fork monitor (`in monitor.c` and `fork-wrapper.c`) to prevent the shell from creating too many subprocesses. If your shell creates too many processes, it will be killed with a “YOU ARE HITTING THE FORK LIMIT” message. Please check your code if you see that message.

Although we have such precautions, you should still take care of the problem. Add in the “`fork()`” call only after thoroughly testing the basic code. You may want to test it separately without putting it in a loop first. For any child process, make sure you have a “`return ...`” or “`exit()`” as the last line of code (even if there is an `exec` before as the `exec` can fail).

If you accidentally ignite a fork bomb on one of the SoC Compute Cluster node, your account may be frozen. Please contact your lab tutor who will then contact SoC Technical Service to unfreeze your account and kill off all your processes.

### Build and Run

Use the commands below to build and launch your shell.

```
$ make          # Build
$ ./monitor myshell    # Launch the shell with fork monitor
```

You may add `-Werror` into the `CFLAGS` in the `Makefile` to make all warnings errors. You will not be penalized for warnings, but you are strongly encouraged to resolve all warnings. Warnings are indications of potential `undefined behavior` which may cause your program to behave in inconsistent and unexpected ways that may not always manifest (i.e., it may work when you test, but it may fail when grading). You are also advised to use `valgrind` to ensure your program is free of memory errors.

On launching your shell, you will then see an `myshell>` prompt where you can input commands in a similar manner to a Bash shell. The commands that this shell should accept are elaborated in the following sections.

### File Structure

When you unzip `Lab2.zip`, you should find the following files:

<code>myshell.c</code>	Your implementation.
<code>myshell.h</code>	<b>Do not modify.</b> Modifications will be ignored when grading.
<code>driver.c</code>	<b>Do not modify.</b> The file will be replaced when grading.
<code>monitor.c</code>	<b>Do not modify.</b> Fork Monitor
<code>fork-wrapper.c</code>	<b>Do not modify.</b> Fork Monitor
<code>check_zip.sh</code>	<b>Do not modify.</b> For you to check your zip file before submission.
<code>Makefile</code>	<b>Do not modify.</b> The file will be replaced when grading.
<code>programs/*</code>	<b>Do not modify.</b> The file will be replaced when grading. Tiny programs for testing.

## Terminologies

The **Table 1** contains information about some terminologies used in this lab document.

Table 1: Terminology Description

Terminology	Description
{program}	A single string with no whitespaces. Will be a <b>valid file path</b> that only contains the following characters: a-z, A-Z, 0-9, forward slash (/), dash (-), and period (.). Refer to the definition of a valid file path at the end of additional details.
(args...)	<b>Optional</b> arguments to the program. Each argument will be a single string with no whitespaces, and will only contain the following characters: a-z, A-Z, 0-9, forward slash (/), dash (-), and period (.). If there is more than 1 argument, the arguments will be separated by whitespaces.
option	It is numeric numbers (both +ve and -ve)
{PID}	The process id of the child process. Your process ids may be different from the sample output, this is expected. Will only contain the following characters: 0-9.
{file}	A single string with no whitespaces. Will be a <b>valid file path</b> that only contains the following characters: a-z, A-Z, 0-9, forward slash (/), dash (-), and period (.). Refer to the definition of a valid file path at the end of additional details.
Any other token, e.g., quit, info, &	They refer to the token itself.

### 3 Exercise 1: Basic Shell

Let us implement a simple shell with limited features in this first exercise. The shell should accept any of the following commands, in a loop:

`{program} (args...)`

Example commands:

```
/bin/ls  
/bin/cat -n file.txt
```

If `{program}` is a readable and executable file:

- Execute the file in a child process with the supplied arguments.
- **Wait** till the child process is done.

Else:

- Print “`{program}` not found” to stderr.

`{program} (args...) &`

**Note the & symbol at the end of the user command.**

If `{program}` is a readable and executable file:

- Run `{program}` in a child process with the supplied arguments.
- Print “**Child [PID] in background**”, where `PID` is the process id of the child process.
- **Continue to accept user commands.**

Else:

- Print “`{program}` not found” to stderr.

`info option`

If `option` is **0**

- Print details of all processes in the order in which they were run. You will need to print their process IDs, their current status (Exited or Running)
- For Exited processes, print their exit codes.
- Please check the sample output for the printing format below.

If `option` is **1**

- Print the number of exited processes.

If `option` is **2**

- Print the number of running processes.

For all other cases print “`Wrong command`” to stderr.

**quit**

- Terminate all RUNNING processes using **SIGTERM** and print “Killing [pid]” for each terminated process.
- Do not wait for the child processes. Print “Goodbye !” and exit the shell.

**Notes:**

- Read **Section 6** to find out additional details about these requirements.
- Note that **info** can only display a return result after a process has exited.

**Question 1.1 (1 marks)**

Explain how do you keep track of all the states of processes?

**Question 1.2 (1 marks)**

Explain how would you find out if a child process in the background has exited when calling info 1?

**Question 1.3 (1 marks)**

How do you check if a program exists and is an executable?

**Example output**

```

1 myshell> info 0
2 myshell> info
3 Wrong command
4 myshell> info -1
5 Wrong command
6 myshell> /bin/echo hello
hello
7 myshell> info 0
8 [3281554] Exited 0
9 myshell>
10 myshell> /bin/notaprogram
11 /bin/notaprogram not found
12 myshell>
13 myshell> info 0
14 [3281554] Exited 0
15 myshell> /bin/sleep 10 &
16 child [3281571] in background
17 myshell> info 0
18 [3281554] Exited 0
19 [3281571] Running
20 myshell> info 2
21 Total running process: 1
22 myshell> info 0
23 [3281554] Exited 0
24 [3281571] Running
25 myshell> info 0
26 [3281554] Exited 0
27 [3281571] Exited 0
28

```

```
29 myshell> info 1
30 Total exited process: 2
31 myshell>
32 myshell> ./programs/result 7
33 myshell> info 0
34 [3281554] Exited 0
35 [3281571] Exited 0
36 [3281579] Exited 7
37 myshell> info 2
38 Total running process: 0
39 myshell> info 1
40 Total exited process: 3
41 myshell> ./programs/result 256
42 myshell> info 0
43 [3281554] Exited 0
44 [3281571] Exited 0
45 [3281579] Exited 7
46 [3281585] Exited 0
47 myshell> ./programs/showCmdArg 5 23 1 &
48 Child [3281597] in background
49 myshell> [Arg 0]: 5
50 [Arg 1]: 23
51 [Arg 2]: 1
52
53 myshell> /bin/sleep 15 &
54 Child [3281599] in background
55 myshell> quit
56 Killing [3281599]
57
58 Goodbye
```

---

**DEMO (4 marks)**

The TA might ask you to type any command or all commands in this exercise. Failing which might result in marks deduction.

## 4 Exercise 2: Advanced Shell

Implement the following commands, in addition to the ones in exercise 1.

---

```
wait {PID}
```

Example command:

```
wait 226
```

{PID} is a process id created using “{program} (args...) &” syntax and has not yet been waited before.

If the process indicated by the process id is RUNNING, **wait** for it.

Else, continue accepting user commands.

No output should be produced.

---

```
terminate {PID}
```

Example command:

```
terminate 226
```

If the process indicated by the process ID {PID} is RUNNING:

- Terminate it by sending it the **SIGTERM** signal.
- You should **not** wait for {PID}.
- The state of {PID} should be “Terminating” until {PID} exits.

Continue accepting user commands. No output should be produced.

---

```
{program1} (args1...) ; {program2} (args2...) ; ...
```

Example command:

```
/bin/ls ; /bin/sleep 5 ; /bin/pwd ; /bin/ls
```

; is an operator that allows multiple “{program} (args...)” to be chained together and executed **sequentially**, and in the foreground.

1. If {program1} exists and is readable and executable:
  - Run and **wait** for {program1}.
  - There might be error output from the program if it fails, which is fine.
2. Else:
  - Print “{program} not found” to stderr.
3. Go back to step 1 with the next {program2}.

Note: There will always be spaces around “;”. It would be helpful to start with 2 chained commands, before extending your implementation to any number of chained commands.

**Extend** the following command from exercise 1:

## info option

Should now have an additional status “Terminating”, in addition to the original “Running” and “Exited”.  
If option is **0**

- Print details of all processes in the order in which they were run. You will need to print their process IDs, their current status (Exited or Running)
- For Exited processes:
  - If the child process ended normally, print the exit code of the child process.
  - If the child process ended abnormally, print which signal (the signal number) caused the child process to exit.
- Please check the sample output for the printing format below.

If option is **3**

- Print the **number** of terminating processes

**Notes:**

- Read **Section 6** to find out additional details about these requirements.

### Question 2.1 (1 marks)

Explain how do you send wait signal to the child process?

### Question 2.2 (1 marks)

Explain how do you send terminate signal to the child process?

### Question 2.3 (1 marks)

What is stdout and stderr? What is the difference between stdout and stderr?

## wait Command

```

1 myshell> /bin/sleep 15 &
2 Child [3285066] in background
3 myshell> info 0
4 [3285066] Running
5 myshell> wait 3285066
6 myshell> info 0
7 [3285066] Exited 0
8 myshell> quit
9
10 Goodbye

```

## terminate Command

```

1 myshell> /bin/sleep 30 &
2 Child [1004267] in background
3 myshell> info 0
4 [1004267] Running
5 myshell> terminate 1004267
6 myshell> info 0
7 [1004267] Exited 15
8 myshell> info 2
9 Total running process: 0

```

---

```

10 myshell>
11 myshell> ./programs/lazy &
12 Child [1005093] in background
13 myshell> Good morning...
14 terminate 1005093
15 myshell> Give me 5 more seconds

16
17 myshell> info 0
18 [1004267] Exited 15
19 [1005093] Terminating
20 myshell> info 3
21 Total terminating process: 1
22 myshell> info 0
23 [1004267] Exited 15
24 [1005093] Exited 0 This is still 0 since lazy program exited properly
25 myshell> info 3 (check source code)
26 Total terminating process: 0
27 myshell> quit

28
29 Goodbye

```

---

## Chained Commands

---

```

1 myshell> /bin/sleep 5 ; /bin/echo Hello ; /bin/echo Bye
2 Hello
3 Bye
4 myshell> /bin/sleep notnumber ; /bin/echo Hello ; /bin/echo Bye
5 /bin/sleep: invalid time interval 'notnumber'
6 Try '/bin/sleep --help' for more information.
7 Hello
8 Bye
9 myshell> info 0
10 [3287915] Exited 0
11 [3287916] Exited 0
12 [3287917] Exited 0
13 [3287945] Exited 1
14 [3287946] Exited 0
15 [3287947] Exited 0
16 myshell> info 1
17 Total exited process: 6
18 myshell> info 2
19 Total running process: 0
20 myshell>
21 myshell> quit

22
23 Goodbye

```

---

**DEMO (4 marks)**

The TA might ask you to type any command or all commands in this exercise. Failing which might result in marks deduction.

## 5 Exercise 3: Redirection

In this exercise, we will implement the redirection operators for our shell by extending the `{program}` (`args...`), the `{program}` (`args...`) & commands, and the `;` operator from exercises 1 and 2.

Implement the following user commands:

```
{program} (args...) (< {file}) (> {file}) (2> {file})
```

(`< {file}`), (`> {file}`), and (`2> {file}`) are optional and may or may not be present. If there are more than 1 present, **all {file}s will be different**, i.e., no reading and writing to the same file.

### Example commands

```
/bin/cat a.txt > b.txt
/bin/sort < test.txt > sorted.txt
/bin/sleep 2 2> error.log
```

- If (`< {file}`) is present:
  - If there exists a file at `{file}`: `{program}` reads the contents of `{file}` as input.
  - Else, print “`{file}` does not exist” to `stderr` and exit the child process with exit code 1.
- If (`> {file}`) is present:
  - If there does not exist a file at `{file}`, create it, then redirect the **standard output** of `{program}` into `{file}`. The `{file}` should be opened in **write** mode, i.e., the file’s existing content will be overwritten.
- If (`2> {file}`) is present:
  - If there does not exist a file at `{file}`, create it, then redirect the **standard error** of `{program}` into `{file}`. The `{file}` should be opened in **write** mode, i.e., the file’s existing content will be overwritten.

Note:

1. A child process should always be created as long as the `{program}` is an executable file. You should check the validity of the files used for redirection within the child process before calling `exec`-family syscalls.
2. There will always be spaces around “`<`”, “`>`” and “`2>`”.

```
{program} (args...) (< {file}) (> {file}) (2> {file}) &
```

**Note the `&` symbol at the end of the user command.** Same behavior as above, except that the command will be run in the background instead, i.e., **your shell should continue accepting user commands**.

```
{program} (args...) (< {file}) (> {file}) (2> {file}) ; {program}
↪ (args...) (< {file}) (> {file}) (2> {file}) ; ...
```

**Example command**

```
/bin/printf hello\nworld\n > test.txt ; /bin/sort < test.txt >
↪ sorted.txt ; /bin/cat sorted.txt
```

The rest of the ; operator's function remains the same as in exercise 2. Note that the same file can be read/written across different programs, as shown above with `test.txt`.

If a file is not found for (< {file}), just print “{file} does not exist” to stderr. Again, continue executing the rest of the programs even if one of the program fails.

**Notes:**

- For simplicity, you may assume that (< {file}) always appears before (> {file}), and both will always appear before (2> {file}). Also, the redirection operators will always appear after {program} (args...). This differs from the actual Bash shell where even something like

```
> a.out < a.in cat
```

is still a valid command.

- Notice that with these new functionalities, you can simulate piping by doing

```
program1 > temp ; program2 < temp
```

Have you wondered how piping actually works? Find out more [here!](#)

- **Read Section 6** to find out additional details about these requirements.

- A flow chart about how the chained commands with **input redirection** (only) should be handled is shown in [Figure 1](#).

**Question 3.1 (1 marks)**

How do you check if a file exists and can be opened for read and write?

**Question 3.2 (1 marks)**

What system call would you use to redirect standard output to a file? What does this system call do?

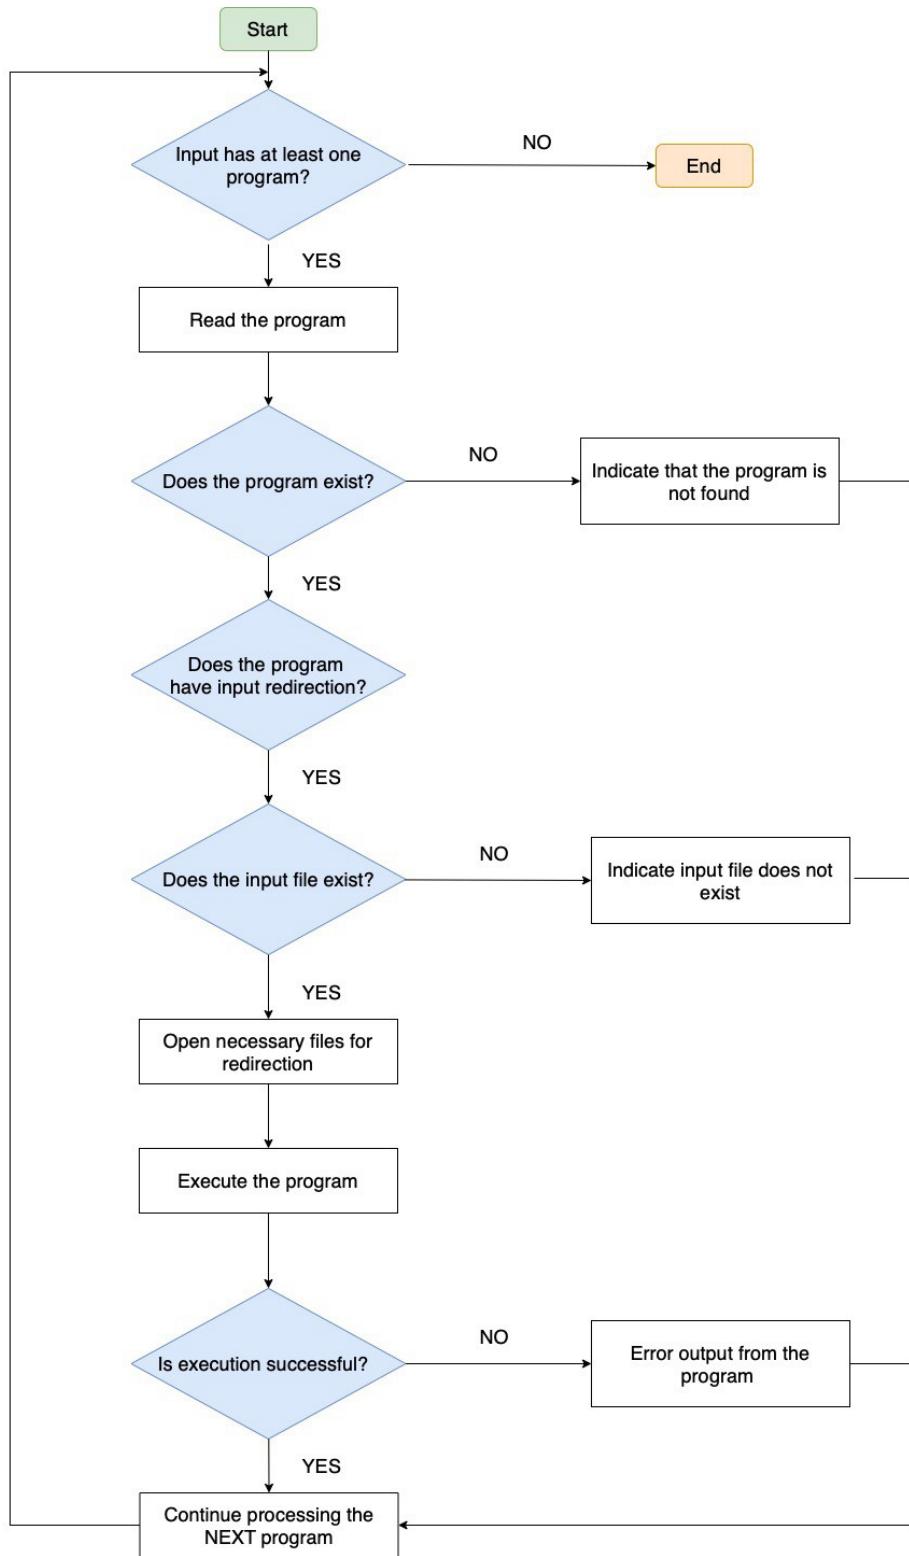


Figure 1: Chained Commands with Input Direction

**Sample Execution**

```
1 myshell> /bin/cat ./programs/result > ./a.txt
2 myshell> /bin/cat ./a.txt
3 #!/bin/bash
4 result=$1
5 exit $result
6 myshell> info 0
7 [3289566] Exited 0
8 [3289596] Exited 0
9 myshell>
10 myshell> /bin/sort < ./a.txt > ./b.txt &
11 Child [3289645] in background
12 myshell> info 0
13 [3289566] Exited 0
14 [3289596] Exited 0
15 [3289645] Exited 0
16 myshell> /bin/cat ./b.txt
17 #!/bin/bash
18 exit $result
19 result=$1
20 myshell>
21 myshell> /bin/sort < ./doesnotexit.txt
22 ./doesnotexit.txt does not exist
23 myshell> info 0
24 [3289566] Exited 0
25 [3289596] Exited 0
26 [3289645] Exited 0
27 [3289661] Exited 0
28 [3289776] Exited 1
29 myshell>
30 myshell> /bin/printf hello\nworld\n > ./a.txt ; /bin/sort < ./a.txt >
31     ↳ ./b.txt ; /bin/cat ./b.txt
32 hello
33 world
34 myshell> info 0
35 [3289566] Exited 0
36 [3289596] Exited 0
37 [3289645] Exited 0
38 [3289661] Exited 0
39 [3289776] Exited 1
40 [3289795] Exited 0
```

```
40 [3289796] Exited 0
41 [3289797] Exited 0
42 myshell>
43 myshell> /bin/echo hello ; /bin/sort < ./doesnotexit.txt
44 hello
45 ./doesnotexit.txt does not exist
46 myshell> info 0
47 [3289566] Exited 0
48 [3289596] Exited 0
49 [3289645] Exited 0
50 [3289661] Exited 0
51 [3289776] Exited 1
52 [3289795] Exited 0
53 [3289796] Exited 0
54 [3289797] Exited 0
55 [3289801] Exited 0
56 [3289802] Exited 1
57 myshell>
58 myshell> /bin/sleep notanumber 2> err.log ; /bin/echo hello
59 hello
60 myshell> /bin/cat err.log
61 /bin/sleep: invalid time interval 'notanumber'
62 Try '/bin/sleep --help' for more information.
63 myshell>
64 myshell> quit
65
66 Goodbye
```

---

**DEMO (4 marks)**

The TA might ask you to type any command or all commands in this exercise. Failing which might result in marks deduction.

## 6 Additional details for all exercises

This section will be long and verbose, but it is essential as we need to limit the scope of the lab by defining some constraints. Please read through it, perhaps after you have read through the exercises. The details below apply to all exercises in this lab, unless otherwise stated.

- Any command that does not satisfy the correct syntax details will not be tested on your shell.
- Any command that does not satisfy the formats specified in the various tables shown in sections 3-5 will not be tested on your shell. To give an example, “`/bin/ls & -a`” does not satisfy our command formats, because firstly, the ampersand must appear at the end of the command, and secondly, if we interpret the ampersand as an argument, it will be syntactically invalid by our definitions.
- It is good practice to check the return values of all syscalls (and indeed, all functions) and handle any errors that occur.
- In total, during the entire lifetime of the program, **no more than 50 (<= 50) processes will be run** (counting both currently running processes and processes that have exited).
- Programs with the same names as the user commands will not be run, i.e., there will be no name collisions.
- Programs that require interactive input will not be tested on your shell. For example, programs such as the Python shell, or your shell itself.
- Notice that **background processes** may sometimes mess up the display on your shell, which may cause your output to be slightly different from the sample session. This is fine. Other than these cases, your output, excluding the PID numbers, **should match the sample sessions exactly**.
- At any point of time, if a file is being read from or executed, we will not be running a command that writes to that same file.
- You may assume that there will be no permission issues for files that already exist (files that are not created by your shell). Your shell will have read, write, and execute permissions, where necessary, for these files that already exist. Similarly, there will be no permission issues for directories as well.
- To illustrate what a **valid file path** is for the purposes of this lab, we split a path up into two parts. The first part is the path to the containing directory (we will call it `<dir>`), such that running the command `cd <dir>` in a bash shell always succeeds. The implications of this are that every directory along the path to our file exists and is accessible. This first part is optional, and if omitted defaults to the current working directory. The second part (non-optional) is the file name, which can only contain the following characters: a-z, A-Z, 0-9, dash (-), and period (.). In `<dir>`, there might or might not be a file name that matches our second part, but either way it is still considered a **valid file path**.
- You will have to check if the process can execute the **programs** or not.
- For simplicity, only regular files/directories will be used in our testing. No symbolic links will be used.