
CS2106

Introduction to OS

Lecture 1 Introduction

Overview

■ Operating Systems basic concepts:

- What is OS?
- Brief History
 - Motivation for OS
- Overview of modern OSes

■ Operating System Structures

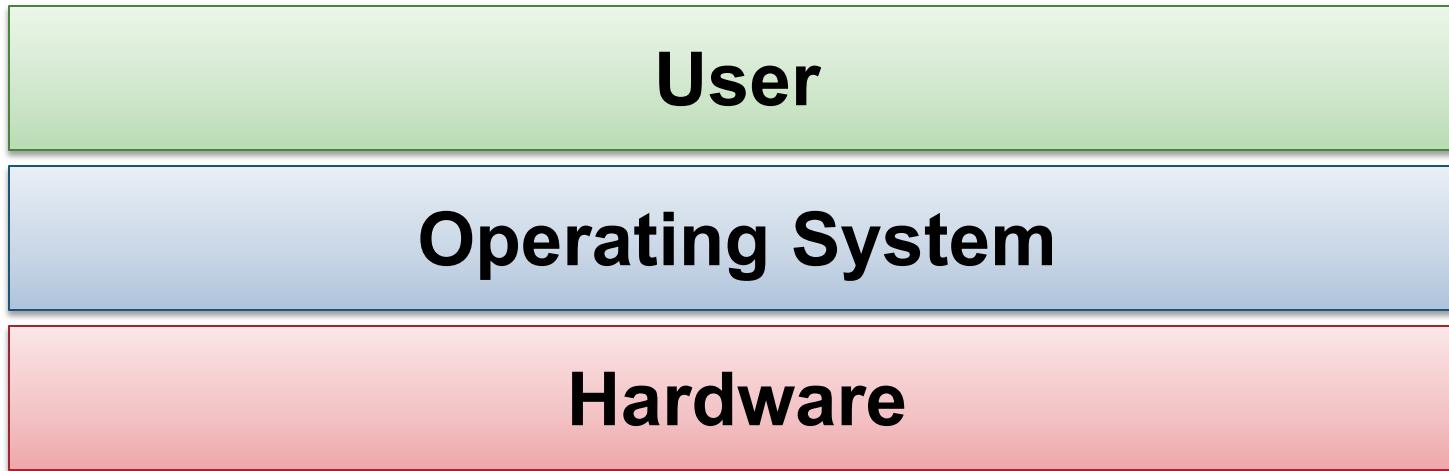
- OS components
- Types of kernel

■ Virtual Machines

What is an OS?

- Incorrect/Incomplete definition:
 - It is the desktop when you boot up your PC
 - The "thing" that stores your games
 - Windows! (or Mac!) (or Linux!)
- One simple definition:
 - A **program** that acts as an **intermediary** between a **computer user** and the **computer hardware**
 - **Wikipedia:** An operating system (OS) is system software that manages computer hardware, software resources, and provides common services for computer programs.

Illustration: What is an OS?



- A simplified view:
 - Will be refined as we move along
- The most general version:
 - Hardware (not only computer!)
 - User (can be application programs or actual person!)

Example of Common OS

- On Computer:
 - Windows 11/10/8/XP
 - Mac OS X
 - Linux distros: Ubuntu, Redhat, Debian, Fedora, CentOS
 - Solaris, FreeBSD
- On Smartphone:
 - iOS, Android
- Other hardware with OS:
 - Game console: PS5, Xbox, Nintendo Switch, ...
 - Home appliance: Smart TV, Smart Watch, ...

To invent the future, you must understand the past

BRIEF HISTORY OF OS

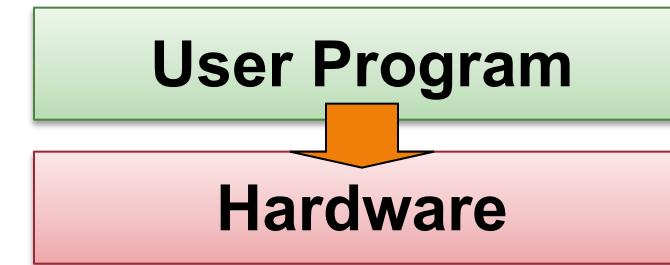
Brief History of OS

- Essentially, OS evolves with:
 - Computer hardware
 - User application and usage pattern
- The "first" computers:
 - Electronic Numerical Integrator And Computer (**ENIAC**)
 - 1945
 - Program controlled by cables and switches
 - Harvard Mark I:
 - 1944
 - Program controlled by punched paper tape

OS for the first computers

- OS Type:

- NO OS



- Programs directly interact with hardware

- Reprogram by changing **physical** configuration of hardware

- **Advantage:**

- Minimal overhead

- **Disadvantage:**

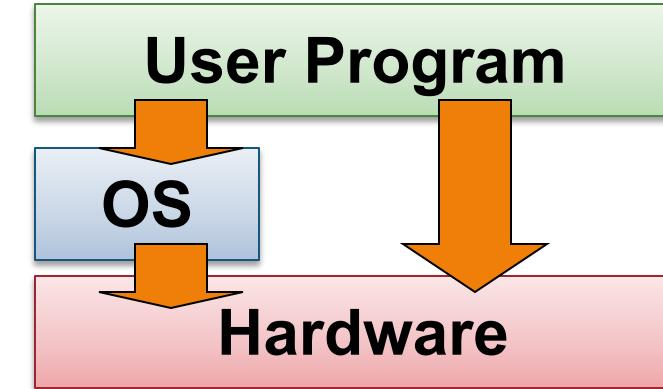
- Not portable
 - Inefficient use of computer!

Mainframes: The "Big Iron"

- Commonly used by large corporations in 60s, 70s
- Common features:
 - No interactive interface
 - Accept programs in the form of:
 - Paper tape, magnetic tape, punch card
 - Support batch processing only
 - Very costly
 - Usually "rented" instead of owned
- Example:
 - IBM 360
 - Cost 5 billion US dollars in 1964 to develop
 - Cost 130k US dollar in 1965 to buy

OS for Mainframes

- OS Type:
 - Batch OS



- Batch OS:
 - Execute user program (a.k.a *job*) one at a time
 - Load job from media, execute, collect result
- User Job:
 - Still interact with hardware directly
 - With additional information for the OS
 - Resource required
 - Job specification

OS for Mainframes: Improvements

- Simple batch processing is inefficient:
 - CPU idle when perform I/O
- One possible Improvements:
 - Multiprogramming:
 - Loads multiple jobs and runs other jobs when I/O needs to be done
 - Overlaps computation with I/O
- Another development of OS during this period (70s):
 - Time-Sharing OS

Time-Sharing OS

■ Features:

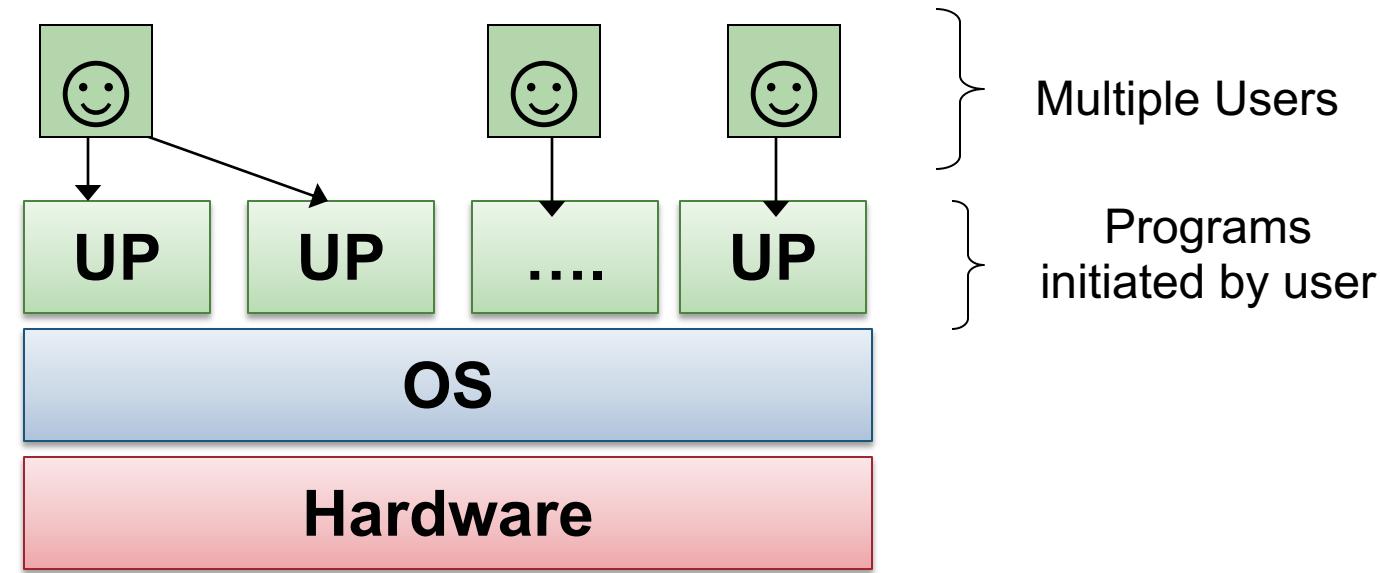
- Allow multiple users to interact with machine using terminals (*teletypes*)
- User job scheduling
 - Illusion of Concurrency
- Memory management

■ Famous Examples:

- CTSS developed at MIT 1960s
- Multics (1970s)
 - Considered as the *parent of Unix*
- Pushed the state of art in virtual memory, security

■ Similar to Unix servers today but more primitive

Time-sharing OS: Illustration



- OS manages the sharing of:
 - CPU time, memory, and storage
- ***Virtualization*** of hardware:
 - Each program executes **as if** it has all the resources to itself

Personal Computer

- Apple II PC (1977):
 - First successfully produced mass home computer
 - Designed by Steve Wozniak (alone!)
- IBM PC (1981):
 - The first generic PC
 - PC becoming a collection of **commodity** hardware components
 - Leads to dominance of Microsoft OSes on PCs: MSDOS (1981) then Windows (1985)

OS on Personal Computer

- Machine (can be) dedicated to user, not timeshared between multiple users
 - Give rise to **personal OS**
- Several Models:
 - Windows model:
 - Single user at a time but possibly more than 1 user can access
 - Dedicated machine
 - Unix model:
 - One user at the workstation but other users can access remotely
 - General time sharing model

Why do we need OS?

MOTIVATIONS OF OS

Motivation for OS: Abstraction

- Large variation in hardware configurations
- Example (Hard disk):
 - Different capacity (500MB, 320GB, 1.5TB, etc.)
 - Different capabilities:
 - Rotation per minutes (RPM)
 - Access (read/write) speed
 - ...
- However, hardware in the same category has well defined and common functionality
 - Example (Hard disk): store and retrieve information

Motivation for OS: Abstraction

- Operating System serves as an **abstraction**:
 - Hide the different low-level details
 - Present the common high-level functionality to user
- The user can then perform essential tasks **through** operating system
 - No need to concern with low-level details
- Provides:
 - Efficiency, programmability, and portability

Motivation for OS: Resource Allocator

- Program execution requires multiple resources:
 - CPU, memory, I/O devices, etc.
- For better utilization of resources, multiple programs should be allowed to execute simultaneously
- OS is a **resource allocator**
 - Manages all resources
 - CPU, Memory, I/O devices
 - Arbitrate potentially conflicting requests
 - for efficient and fair resource use

Motivation for OS: Control Program

- Program can misuse the computer:
 - Accidentally: due to coding bugs
 - Maliciously: virus, malware, etc.
- Multiple users can share the computer:
 - Tricky to ensure isolation among users
- OS is a **control program**
 - Controls execution of programs
 - Prevent errors and improper use of the computer
 - Provides security, isolation, and protection

Motivation for OS: Summary

- Manage resources and coordination
 - Process synchronization, resource sharing
- Simplify programming
 - Abstraction of hardware, convenient services
- Enforce usage policies
- Security and protection
- User program portability:
 - Across different hardware
- Efficiency
 - Sophisticated implementations
 - Optimized for particular usage and hardware

The families of modern OS

OVERVIEW OF MODERN OS

Modern OS: Overview

Desktop



Mobile

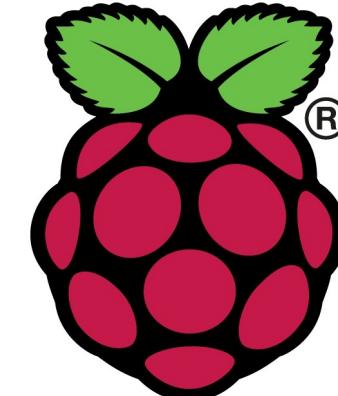


ANDROID

Real-Time

freeRTOS

Embedded



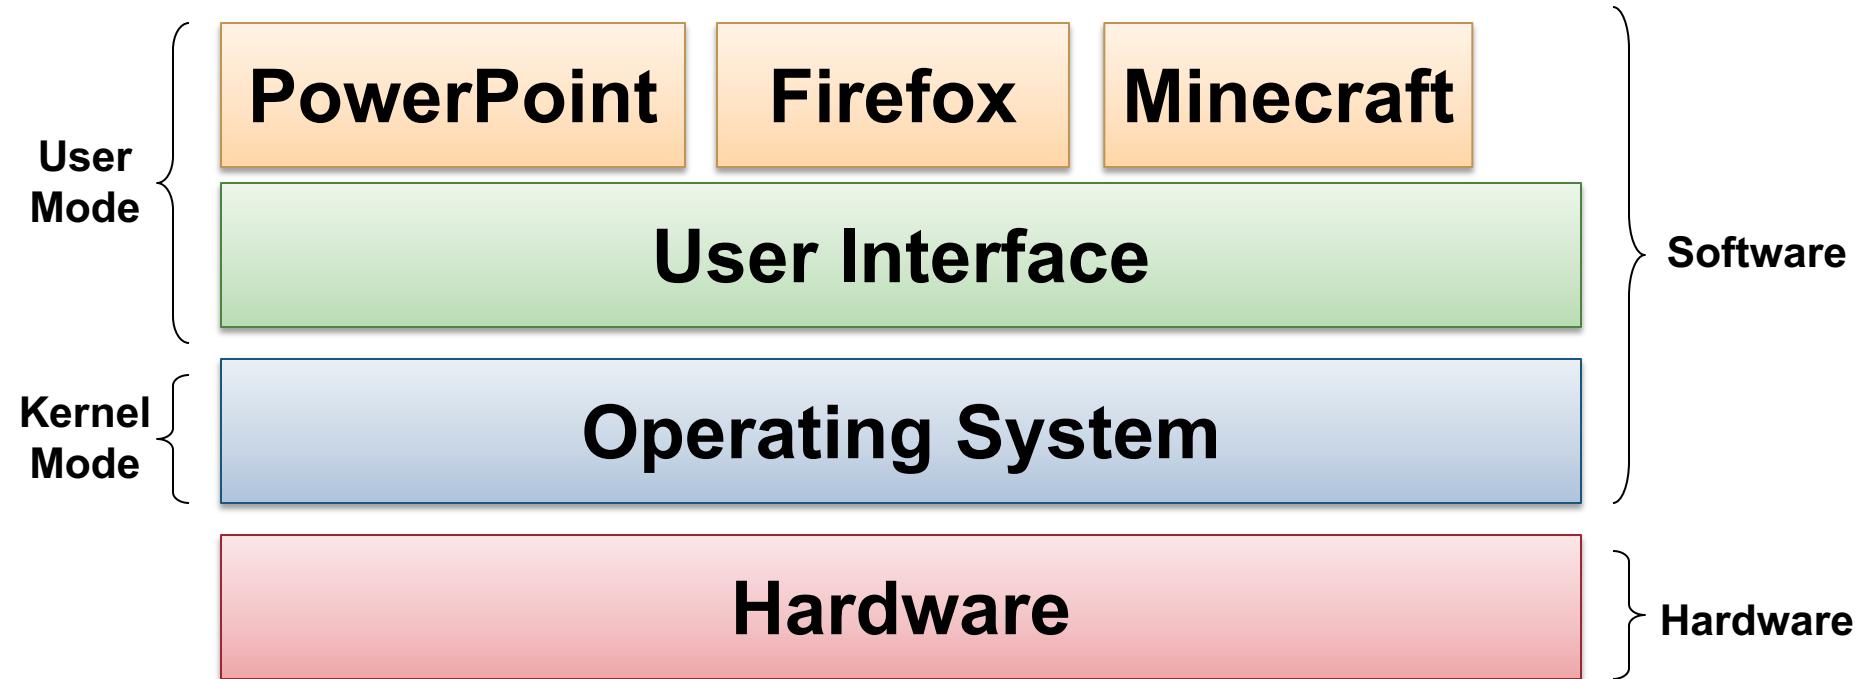
Common Architecture for OS

OS STRUCTURE

Operating System Structures

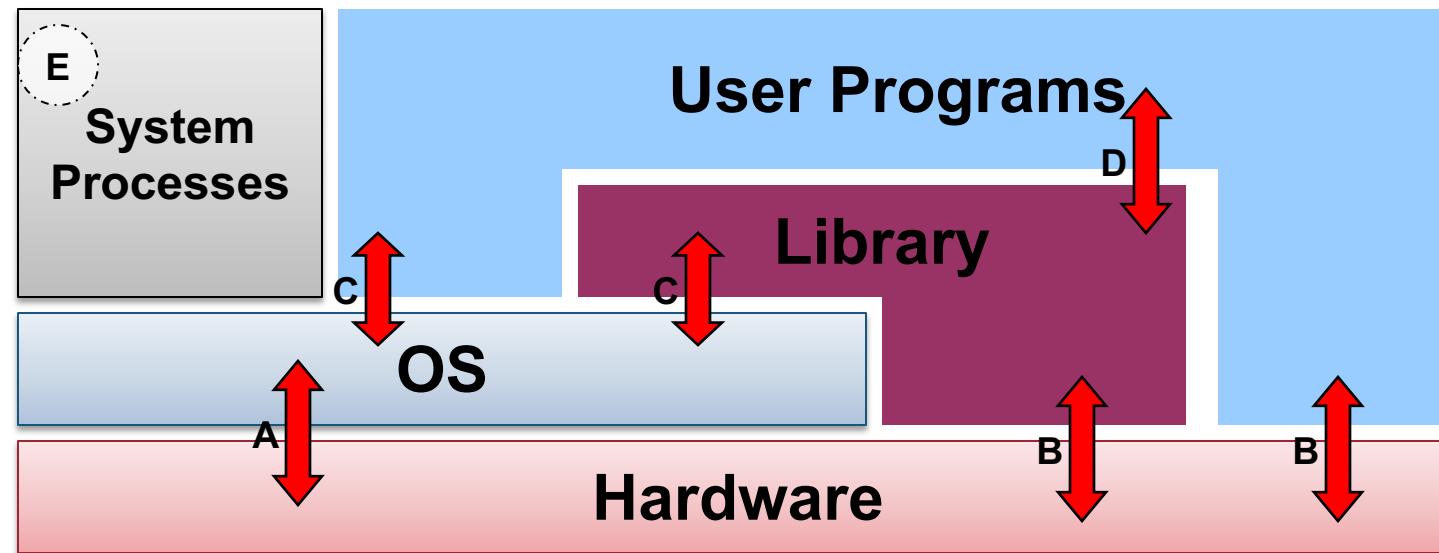
- We have identified the major capabilities/functionalities of an OS
 - i.e., the *specification* of the OS
- Let us now consider:
 - The best way to provide these capabilities
 - i.e., the *implementations* of the OS
- Operating system structure:
 - *Organization* of the various components
 - Important factors:
 - Flexibility
 - Robustness
 - Maintainability
 - Performance

Illustration: High-level view of OS



- Operating System is essentially a software
 - Runs in ***kernel mode***: direct access to all hardware resources
- Other software executes in ***user mode***
 - With limited (or controlled) access to hardware resources

Illustration: Generic OS Components



- **A:** OS executing machine instructions
- **B:** Normal machine instructions executed (program/library code)
- **C:** Calling OS using ***system call interface***
- **D:** User program calls library code
- **E:** System processes
 - Provide high-level services, usually part of OS

OS as a Program

- A critical part of an OS is the **kernel**
 - Just another program with some special features
 - Deals with hardware issues
 - Provides system call interface
 - Special code for interrupt handlers, device drivers
- Kernel code has to be different than normal programs:
 - Can't use system calls in kernel code
 - Can't use normal libraries
 - No “normal” I/O
- Consider this:
 - Normal programs use OS: what does OS use? ☺

Implementing Operating System

- **Programming Language:**
 - Historically in assembly/machine code
 - Now in **HLLs**:
 - Especially C/C++
 - Heavily hardware architecture dependent
- **Common code organization:**
 1. Machine independent HLL
 2. Machine dependent HLL
 3. Machine dependent assembly code
- **Challenges:**
 - “No one else” to rely on for nice services
 - Debugging is hard
 - Complexity
 - Enormous Codebase

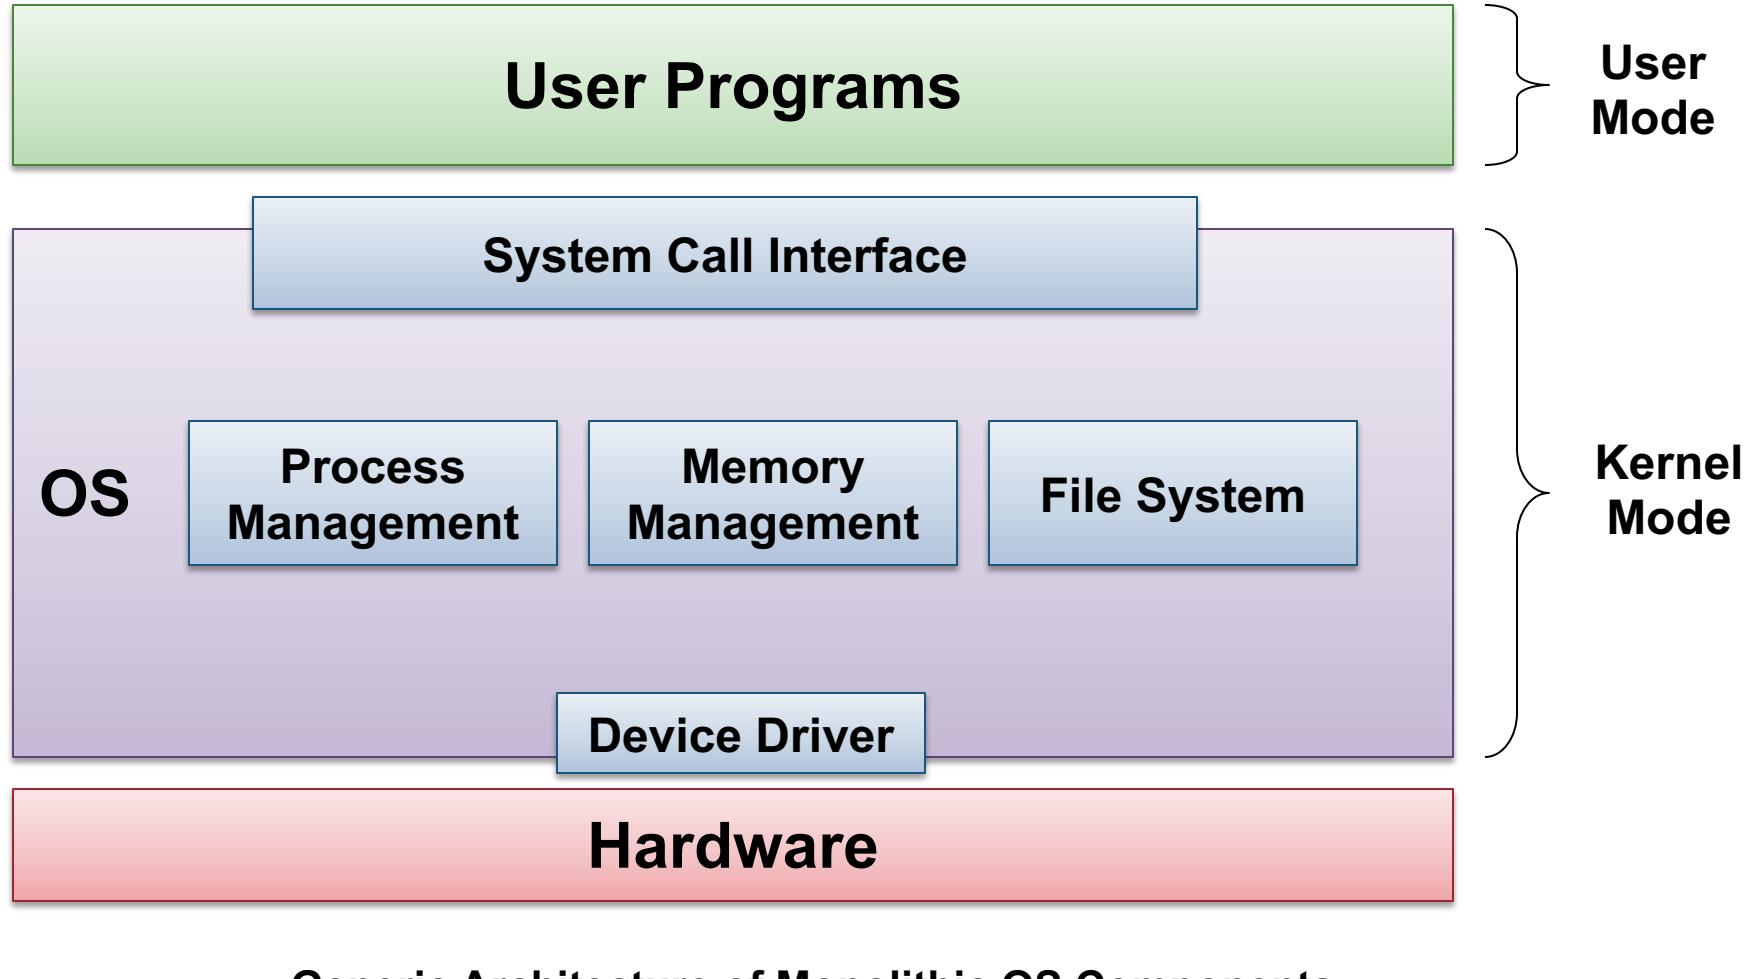
OS Structures

- Several ways to structure an OS:
 - **Monolithic**
 - **Microkernel**
 - Layered
 - Client-Server
 - Exokernel
 - Hybrid
 - ...
- We will cover the first two in details:
 - They represent the whole range of possibilities
 - Most other approaches are variant or improvement

Monolithic OS

- Kernel is:
 - One **BIG** special program
 - Various services and components are integral part
 - Good software engineering principles are still possible with:
 - Modularization
 - Separation of interfaces and implementation
- This is the traditional approach taken by:
 - Most Unix variants, DOS, Windows 9x
- **Advantages:**
 - Well understood
 - Good performance
- **Disadvantages:**
 - Highly coupled components
 - Usually very complicated internal structure

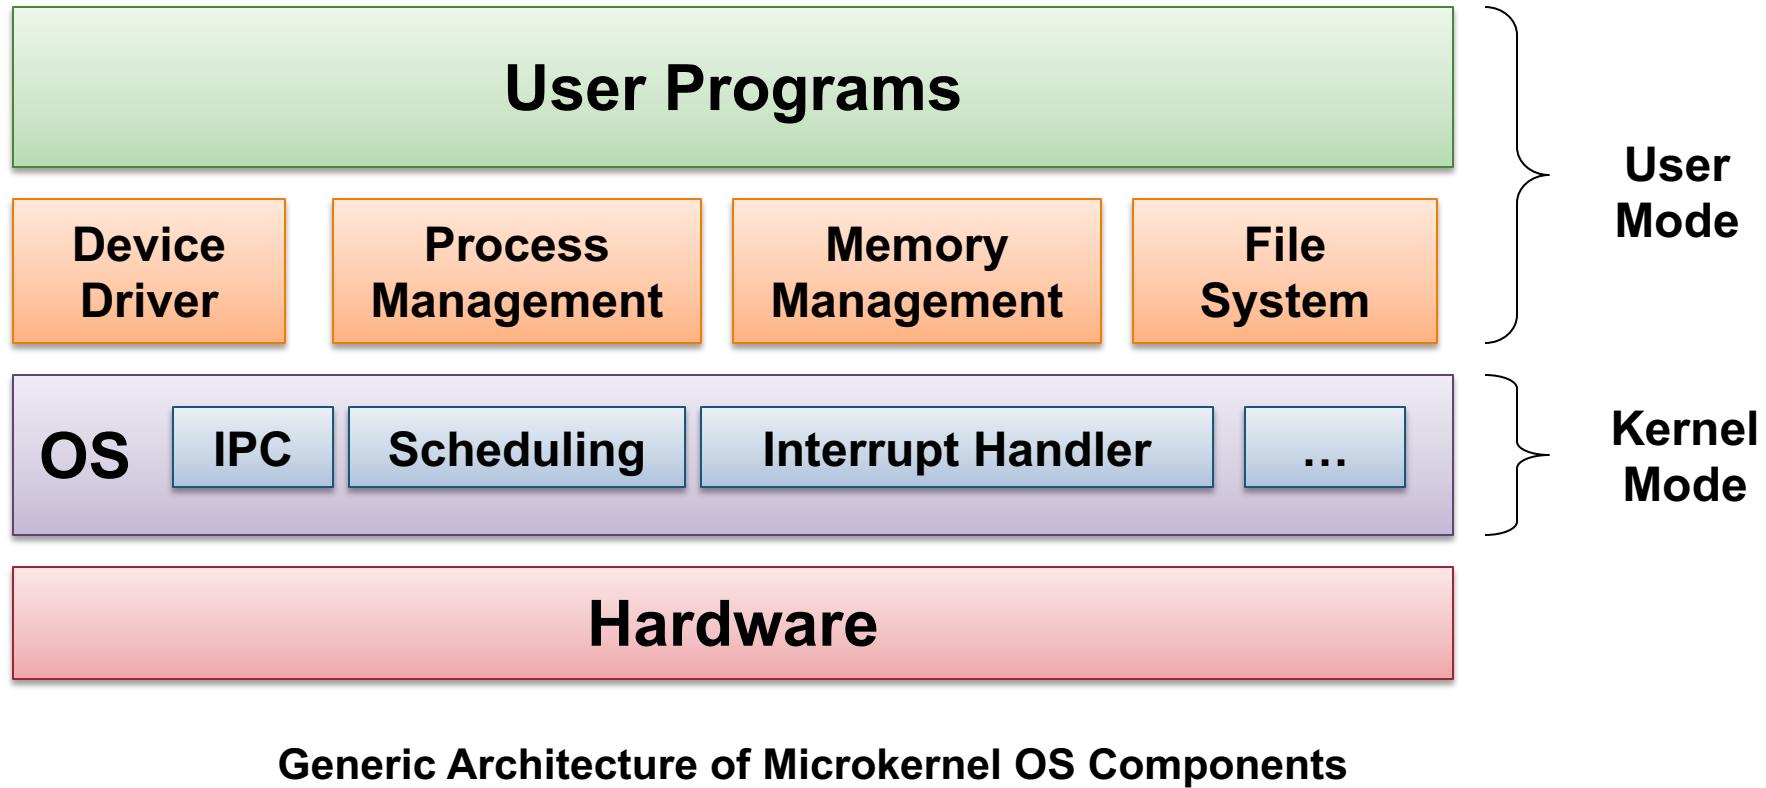
Monolithic Kernel Illustration



Microkernel OS

- Kernel is:
 - Very small and clean
 - Only provides basic and essential facilities:
 - Inter-Process Communication (IPC)
 - Address space management
 - Thread management
 - ...
- Higher-level OS services:
 - Built on top of the basic facilities
 - Run as server process **outside** of the kernel
 - Use IPC to communicate
- Advantages:
 - Kernel is generally more robust and more extensible
 - Better isolation and protection between kernel and high-level services
- Disadvantages:
 - Lower Performance

Microkernel Components



Other Operating System Structure

- Layered Systems:
 - Generalization of monolithic system
 - Organize the components into hierarchy of layers
 - Upper layers make use of the lower layers
 - Lowest layer is the hardware
 - Highest layer is the user interface
- Client-Server Model
 - Variation of microkernel
 - Two classes of processes:
 - Client process request service from server process
 - Server process built on top of the microkernel
 - Client and server process can be on separate machine!

Ways of running OSes

VIRTUAL MACHINES

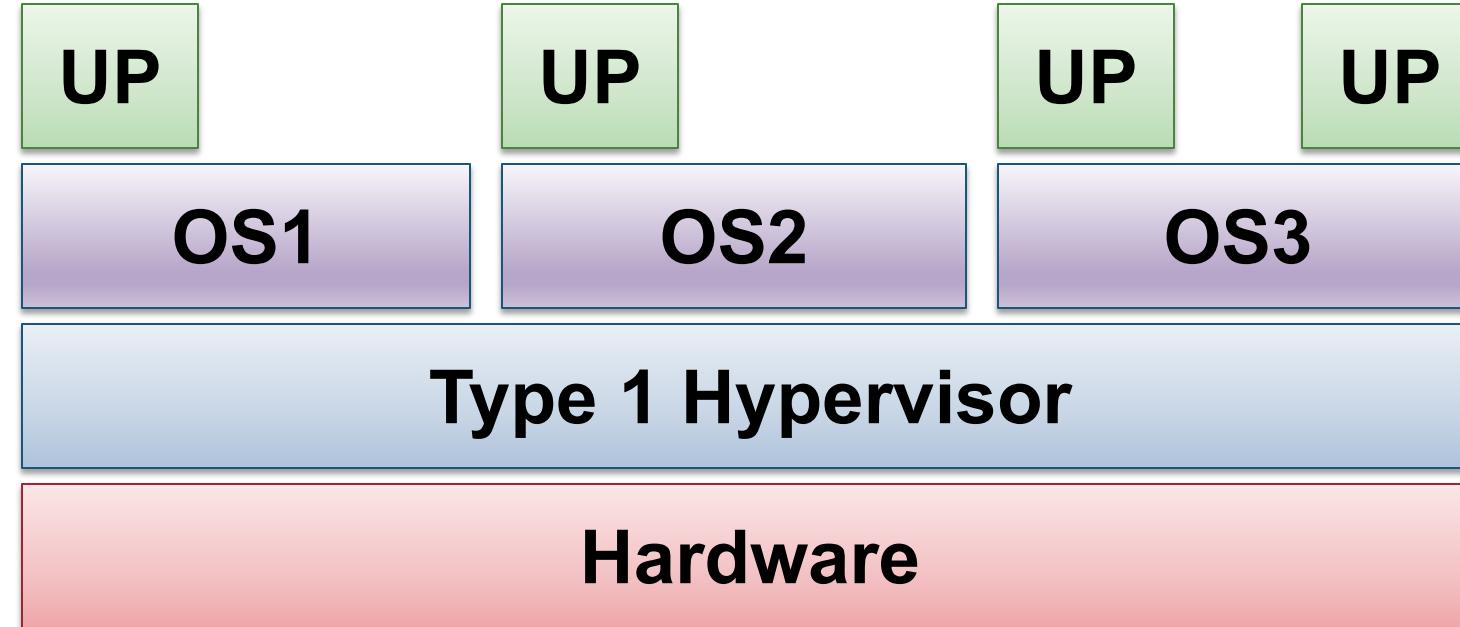
Motivation: Why Virtual Machines

- OS assumes total control of the hardware:
 - What if we want to run several OSes on the same hardware at the same time?
 - Cloud computing (Infrastructure as a Service)
- OS is hard to debug / monitor:
 - How do we observe the working of the OS?
 - How do we test a potentially destructive implementation?

Definition: Virtual Machine

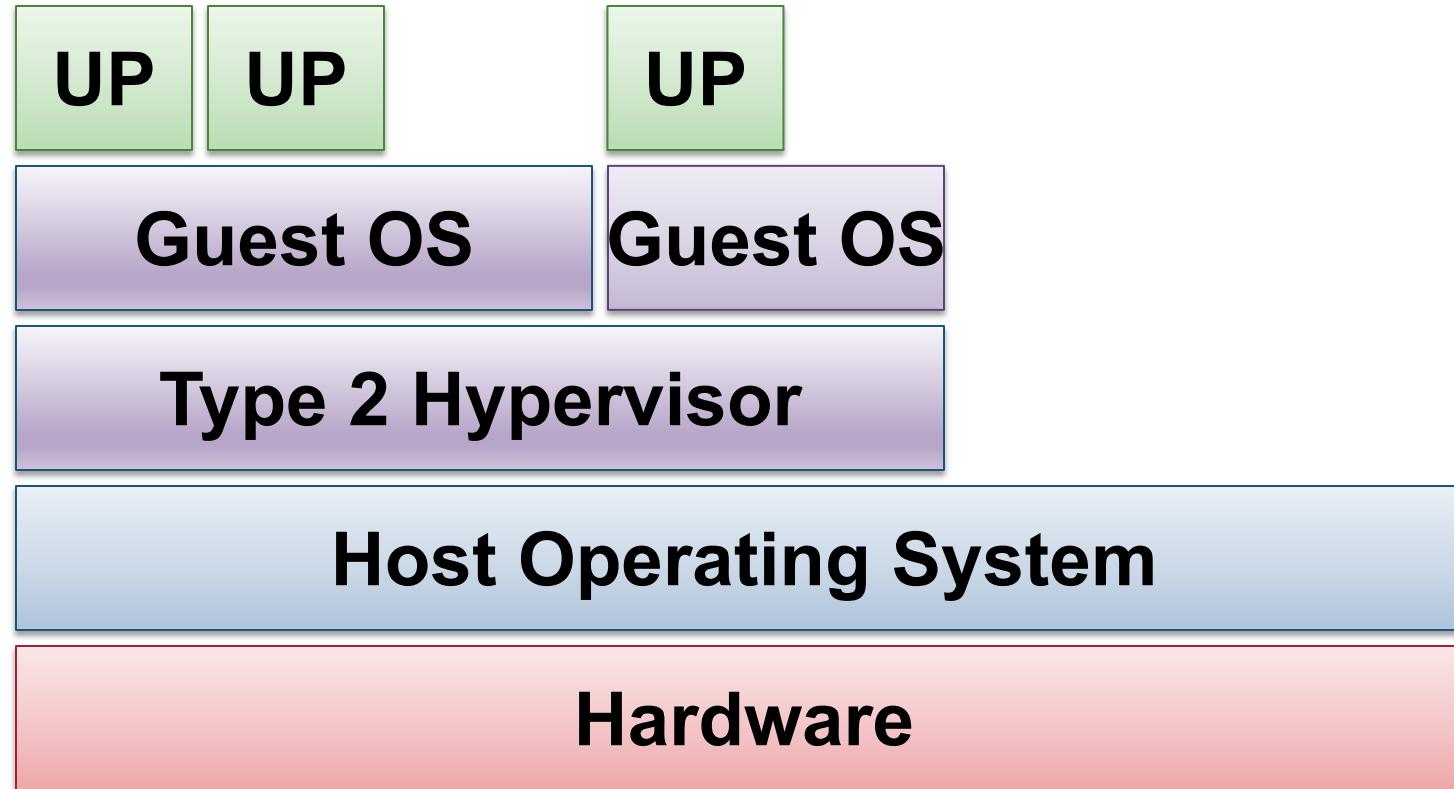
- **Virtual Machine:**
 - A software emulation of hardware
 - **Virtualization** of underlying hardware
 - Illusion of complete hardware to level above: memory, CPU, hard disk, ...
 - Normal (primitive) operating system can then run on top of the virtual machine
- Created and managed by **Hypervisor**
 - Also known as **Virtual Machine Monitor (VMM)**
 - Two classes of implementations shown next

Type 1 Hypervisor



- Type 1 hypervisor:
 - AKA bare-metal hypervisor
 - Provides individual *virtual machines* to guest OSes
 - E.g., IBM VM/370, VMware ESXi

Type 2 Hypervisor



- Type 2 hypervisor
 - Runs in host OS
 - Guest OS runs inside Virtual Machine
 - e.g., VMware Workstation, VirtualBox, QEMU

Summary

- Definition of Operating System
- Roles of Operating System
- Common Operating System families
- Operating System structure

Reference

- Modern Operating System (4th Edition)
 - By Andrew S.Tanenbaum
 - Published by Pearson

- Operating System Concepts (8th Edition)
 - By Abraham Silberschatz, Peter Baer Galvin & Greg Gagne
 - Published by McGraw Hill