

Process Management

Process Abstraction

Lecture 2

Overview

- Program execution:
 - Hardware context
 - Memory context
 - Code & data
 - Function call
 - Dynamically allocated memory
- Introduction to process management
 - OS context
 - Process state
 - Process Control Block and Process Table
- OS interaction with process

Recap: C Sample Program and Assembly Code


```
int i = 0;  
  
i = i + 20;
```

C Code Fragment


```
lw    $1, 4096          // Assume address of i = 4096  
addi $1, $0, 0          // register $1 = 0  
sw    $1, 4096          // i = 0  
  
lw    $2, 4096          // read i  
addi $3, $2, 20          // $3 = $2 + 20  
sw    $3, 4096          // i = i + 20
```

Corresponding MIPS-like Assembly Code

Recap: Program Execution (Memory)



Recap: Generic Computer Organization



Recap: Component Description

- Memory:
 - Storage for instruction and data
- Cache:
 - Duplicate part of the memory for faster access
 - Usually split into instruction cache and data cache
- Fetch unit:
 - Loads instruction from memory
 - Location indicated by a special register: **Program Counter (PC)**

Recap: Component Description (cont)

- Functional units:
 - Carry out the instruction execution
 - Dedicated to different instruction type
- Registers:
 - Internal storage for the fastest access speed
 - General Purpose Register (GPR):
 - Accessible by user program
 - Special Register:
 - Program Counter (PC)
 - Stack Pointer (SP)
 - Frame Pointer (FP)
 - ...

Recap: Basic Instruction Execution

- Instruction X is fetched
 - Memory location indicated by **Program Counter**
- Instruction X dispatched to the corresponding Functional Unit
 - Read operands if applicable
 - Usually from memory or GPR
 - Result computed
 - Write value if applicable
 - Usually to memory or GPR
- Instruction X is completed
 - PC updated for the next instruction

Recap: What you should know 😊

- An executable (binary) consists of two major components:
 - Instructions and Data
- When a program is **under execution**, there are **more information**:
 - Memory context:
 - **Text and Data, ...**
 - Hardware context:
 - **General purpose registers, Program Counter, ...**
- Actually, there are **other types of memory usage** during program execution
 - Coming up next

Memory Context

Function Call

What if `f()` calls `u()` calls `n()`?

Function Call : Challenges

```
int i = 0;  
i = i + 20;
```

C Code Fragment

VS

```
int g(int i, int j)  
{  
    int a;  
  
    a = i + j  
    return a;  
}
```


C Code with Function

Consider:

- How do we allocate memory space for variables **i**, **j** and **a**?
 - Can we just make use of the "data" memory space?
- What are the key issues?

Function Call : Control Flow and Data

- **f() calls g()**
 - f() is the **caller**
 - g() is the **callee**
- Important Steps:
 1. Setup the parameters
 2. Transfer control to callee
 3. Setup local variable
 4. Store result if applicable
 5. Return to caller



Function Call : Control Flow and Data

■ Control Flow Issues:

- Need to jump to the function body
- Need to resume when the function call is done
- ➔ Minimally, need to store the PC of the caller

■ Data Storage Issues:

- Need to pass parameters to the function
- Need to capture the return result
- May have local variables declaration

➔ Need a **new region of memory** that dynamically used by function invocations

Introducing Stack Memory

■ **Stack Memory Region:**

- The new memory region to store information for function invocation

■ Information of a function invocation is described by a **stack frame**

■ Stack frame contains:

- Return address of the caller
- Arguments (parameters) for the function
- Storage for local variables
- Other information.... (more later)

Stack Pointer

- The top of stack region (first unused location) is logically indicated by a **Stack Pointer**:
 - Most CPU has a specialized register for this purpose
 - Stack frame is added on top when a function is invoked
 - Stack "grows"
 - Stack frame is removed from top when a function call ends
 - Stack "shrinks"

Illustration: Stack Memory




Illustration: Stack Memory Usage (1 / 5)

```
void f()
{
    ...
    g(); ← At this
    ...
}

void g()
{
    h();
    ...
}

void h()
{
    ...
}
```




Illustration: Stack Memory Usage (2 / 5)

```
void f()
{
    ...
    g();
    ...
}

void g()
{
    h();           ← At this
    ...
}

void h()
{
    ...
}
```




Illustration: Stack Memory Usage (3 / 5)

```
void f()
{
    ...
    g();
    ...
}

void g()
{
    h();
    ...
}

void h()
{
    ...
}
```




Illustration: Stack Memory Usage (4 / 5)

```
void f()
{
    ...
    g();
    ...
}

void g()
{
    h();
    ...
}

void h()
{
    ...
}
```

At this point




Illustration: Stack Memory Usage (5 / 5)

```
void f()
{
    ...
    g();
    ...
}

void g()
{
    h();
    ...
}

void h()
{
    ...
}
```

At this point





Illustration: Stack Frame v1.0



Function Call Convention

- Different ways to setup stack frame:
 - Known as **function call convention**
 - Main differences:
 - What information is stored in stack frame vs. in registers?
 - Which portion of stack frame is prepared by caller / callee?
 - Which portion of stack fame is cleared by caller / callee?
 - Who between caller / callee to adjust the stack pointer?
- No universal way
 - Hardware and programming language dependent
- An example scheme is described next

Stack Frame Setup



■ Prepare to make a function call:


- **Caller:** Pass parameters using registers and/or stack
 - **Caller:** Save Return PC on stack
- **Transfer Control from Caller to Callee**

- **Callee:** Save the old Stack Pointer (SP)
- **Callee:** Allocate space for local variables of callee on stack
- **Callee:** Adjust SP to point to new stack top

Illustration: Calling function ***g()***

```
void f(int a, int b)
{
    int c;
    a = 123;
    b = 456;
    c = g(a, b);
    ...
}

int g(int i, int j)
{
    int a;
    a = i + j
    return a * 2;
}
```



Stack Frame Teardown



- On returning from function call:
 - **Callee**: Place return result in register (if applicable)
 - **Callee**: Restore saved Stack Pointer
- Transfer control back to caller using saved PC
- **Caller**: Utilize return result (if applicable)
- **Caller**: Continues execution in caller

Illustration: Function *g()* finishes



```
void f(int a, int b)
{
    int c;

    a = 123;
    b = 456;
    c = g(a, b);
    . . .
}

int g(int i, int j)
{
    int a;

    a = i + j
    return a * 2;
}
```

Execution resumes here



Other Information in Stack Frame

- We have described the basic idea of:
 - Stack frame
 - Calling Convention: Setup and Teardown
- A few common additional information in the stack frame:
 - Frame Pointer
 - Saved Registers


Frame Pointer

- To facilitate the access of various stack frame items:
 - Stack Pointer is hard to use as it can change
 - Some processors provide a dedicated register **Frame Pointer**
- The Frame Pointer points to a fixed location in a stack frame
 - Other items are accessed as a displacement from the Frame Pointer
- The usage of FP is platform dependent

Saved Registers

- The number of general purpose register (GPR) on most processors are very limited:
 - E.g., MIPS has 32 GPRs, x86 has 16 GPRs
- When GPRs are exhausted:
 - Use memory to temporary hold the GPR value
 - GPR can then be reused for other purposes
 - GPR value can be restored afterwards
 - Known as **register spilling**
- Similarly, a function can spill the registers it intend to use before the function starts (callee-saved)
 - Restore those registers at the end of function

Illustration: Stack Frame v2.0



Stack Frame Setup / Teardown [Updated]

- On executing function call:

- **Caller**: Pass arguments with registers and/or stack
 - **Caller**: Save Return PC on stack
- **Transfer control from caller to callee**

- **Callee**: Save registers used by callee. Save old FP, SP
 - **Callee**: Allocate space for local variables of callee on stack
 - **Callee**: Adjust SP to point to new stack top; adjust FP

- On returning from function call:

- **Callee**: Restore saved registers, FP, SP
 - Transfer control from callee to caller using saved PC
 - **Caller**: Continues execution in caller

- Remember, just an example!

Function Call Summary

- In this part, we learned:
 - Another portion of memory space is used as a **Stack Memory**
 - Stack Memory stores the executing function using **Stack Frame**
 - Typical information stored on a stack frame
 - Typical scheme of setting up and tearing down a stack frame
 - The usage of Stack Pointer and Frame Pointer

Memory Context

Dynamically Allocated Memory

Hmm... I need more memory

Dynamically Allocated Memory

- Most programming languages allow dynamically allocated memory:
 - i.e., acquire memory space during **execution time**
- Examples:
 - In C, the ***malloc()*** function call
 - In C++, the ***new*** keyword
 - In Java, the ***new*** keyword
- Question:
 - Can we use the existing "Data" or "Stack" memory regions?

Dynamically Allocated Memory


■ Observations:

- These memory blocks have different behaviors:
 1. Allocated only at runtime, i.e., size is not known during compilation time → Cannot place in **Data** region
 2. No definite deallocation timing, e.g., can be explicitly freed by programmer in C/C++, can be implicitly freed by garbage collector in Java → Cannot place in **Stack** region

■ Solution:

- Setup a separate **heap memory region**

Illustration for Heap Memory



Managing Heap Memory

- Heap memory is a lot trickier to manage due to its nature:
 - Variable size
 - Variable allocation / deallocation timing
- You can easily construct a scenario where heap memory are allocated /deallocated in such a way to create "holes" in the memory
 - Free memory block squeezed in between of occupied memory block
- We will learn more in the memory management (much) later in the course

Checkpoint: Contexts updated

■ Information describing a **program execution**:

- Memory context:
 - Text, Data, **Stack** and **Heap**
- Hardware context:
 - General purpose registers, Program Counter, **Stack pointer**, **Stack frame pointer**,

Overview

- Program execution:
 - Hardware Context
 - Memory Context
 - Code & Data
 - Function call
 - Dynamically allocated memory
- **Introduction to Process Management**
 - OS Context
 - Process State
 - Process Control Block and Process Table
- OS interaction with Process

Recap: Efficient Hardware Utilization

- OS should provide efficient use of hardware resources:
 - by managing the programs executing on the hardware
- Observation:
 - If there is only **one program executing at any point in time**, how can we utilize hardware resources effectively?
 - Batch processing?
- Solution:
 - Allow **multiple programs to share the hardware**
 - e.g., multiprogramming, time-sharing

Introduction to Process Management

- As the OS, to switch from running program **A** to program **B** requires:
 1. Information regarding the execution of program **A** needs to be stored
 2. Program **A**'s information is replaced with the information required to run program **B**
- Hence, we need:
 - An **abstraction** to describe a running program
 - aka **process**

Key Topics

Process Abstraction

- Information describing an executing program

Process Scheduling

- Deciding which process get to execute

Inter-Process Communication & Synchronization




- Passing information between processes

Alternative to Process

- Light-weight process aka Thread

Process Abstraction

- **(Process / Task / Job)** is a dynamic abstraction for executing program
 - information required to describe a **running program**



OS Context

Process Id & Process State

Your ID? Give me a status report!


Process Identification

- To distinguish processes from each other
 - Common approach is to use process ID (**PID**)
 - Just a number
 - Unique among processes
- There are a couple of OS dependent issues:
 - Are PIDs reused?
 - Does it limit the maximum no. of processes?
 - Are there reserved PIDs?

Introducing Process State


- In the multitasking scenario:
 - A process can be:
 - Running OR
 - Not-running, e.g., another process running
- A process can be **ready to run**
 - But not actually executing
 - E.g., waiting for its turn to use the CPU
- Hence, each process should have a **process state**:
 - As an indication of the execution status

(Simple) Process Model State Diagram



- The set of states and transitions are known as **process model**
 - Describes the behaviors of a process

Generic 5-State Process Model



Notes: generic process states, details vary in actual OS

Process States for 5-Stage Model

■ **New:**

- ❑ New process created
- ❑ May still be under initialization → not yet ready

■ **Ready:**

- ❑ Process is waiting to run

■ **Running:**

- ❑ Process being executed on CPU

■ **Blocked:**

- ❑ Process waiting (sleeping) for event
- ❑ Cannot execute until event is available

■ **Terminated:**

- ❑ Process has finished execution, may require OS cleanup

Process State Transitions in 5-Stage Model

- **Create** (`nil` → `New`):

- New process is created

- **Admit** (`New` → `Ready`):

- Process ready to be scheduled for running

- **Switch** (`Ready` → `Running`):

- Process selected to run

- **Switch** (`Running` → `Ready`):

- Process gives up CPU voluntarily or *preempted* by scheduler

Process State Transitions

■ Event wait (Running → Blocked):

- Process requests event/resource/service which is not available/in progress
- Example events:
 - Acquiring lock, waiting for I/O, (*more later*)


■ Event occurs (Blocked → Ready):

- Event occurs → process can continue

Global View of Process States

- Given n processes:
 - With 1 CPU (core):
 - ≤ 1 process in running state
 - conceptually 1 transition at a time
 - With m CPUs (cores):
 - $\leq m$ process in running state
 - possibly parallel transitions
- Different processes may be in different states
 - each process may be in different part of its state diagram
- **Assumption in CS2106: Our CPU has 1 core!**

Queuing Model of 5 state transition



Notes:

- More than 1 process can be in ready + blocked queues
- May have separate event queues
- Queuing model gives global view of the processes, i.e., how the OS views them

Checkpoint: Contexts updated

- When a program is **under execution**, there are **more information**:
 - Memory context:
 - Text and Data, Stack, and Heap
 - Hardware context:
 - General Purpose Registers, Program Counter, Stack Pointer, Frame Pointer,
...
 - OS context:
 - **Process ID, Process State, ...**

Process Table & Process Control Block

Putting it together


Process Control Block & Table

- The entire execution context for a process
 - Traditionally called **Process Control Block (PCB)** or **Process Table Entry**
- Kernel maintains PCB for all processes
 - Conceptually stored as one table representing all processes

Interesting Issues:

- Scalability
 - How many concurrent processes can you have?
- Efficiency
 - Should provide efficient access with minimum space wastage

Illustration of a Process Table



Process interaction with OS

System Calls

Can you please do this for me?

System Calls

- Application Program Interface (API) to OS
 - Provides way of calling facilities/services in kernel
 - **NOT** the same as normal function call
 - have to change from user mode to kernel mode
- Different OS have different APIs:
 - Unix Variants:
 - Most follows **POSIX** standards
 - Small number of calls: ~100
 - Windows Family:
 - Uses **Win API** across different Windows versions
 - New version of windows usually adds more calls
 - Huge number of calls:~1000

Unix System Calls in C/C++ program

- In C/C++ program, system call can be invoked *almost directly*
 - Majority of the system calls have a library version with the **same name** and the same parameters
 - The library version act as a **function wrapper**
 - Other than that, a few library functions present a more user friendly version to the programmer
 - E.g., lesser number of parameters, more flexible parameter values, etc.
 - The library version acts as a **function adapter**

Example

```
#include <unistd.h>
#include <stdio.h>

int main()
{
    int pid;

    /* get Process ID */
    pid = getpid();

    printf("process id = %d\n", pid);

    return 0;
}
```

Library call that has the same name as a system call

Library call that make a system call

- System Calls invoked in this example:
 - **getpid()**
 - **write()** – made by **printf()** library call


General System Call Mechanism

1. User program invokes the library call
 - Using the normal function call mechanism as discussed
2. Library call (usually in assembly code) places the **system call number** in a designated location
 - E.g., a register
3. Library call executes a special instruction to switch from user mode to kernel mode
 - That instruction is commonly known as **TRAP**
 - Saves CPU state

General System Call Mechanism (cont)

4. Now in kernel mode, the appropriate system call handler is determined:
 - Using the system call number as index
 - This step is usually handled by a **dispatcher**
5. System call handler is executed:
 - Carry out the actual request
6. System call handler ended:
 - Restore CPU state, and return to the library call
 - Switch from kernel mode to user mode
7. Library call return to the user program:
 - via normal function return mechanism

Illustration: System Call Mechanism



Process interaction with OS

Exception and Interrupt

Ops!


Exception

- Executing a **machine level instruction** can cause exception
- For example:
 - Arithmetic Errors
 - Overflow, Division by Zero
 - Memory Accessing Errors
 - Illegal memory address, Misaligned memory access
 - ...
- Exception is **Synchronous**
 - occur due to program execution
- Effect of exception:
 - Have to execute an **exception handler**
 - Similar to a **forced function call**

Interrupt

- External events can interrupt the execution of a program
- Usually hardware related, e.g.:
 - Timer, mouse movement, keyboard pressed, etc.
- Interrupt is **asynchronous**
 - Events that occurs **independent** of program execution
- Effect of interrupt:
 - Program execution is suspended
 - Have to execute an **interrupt handler**

Exception/Interrupt Handler: Illustration



1. Exception/Interrupt occurs:

- Control transfer to a handler routine **automatically**

2. Return from handler routine:

- Program execution resume
- **May** behave as if nothing happened

Summary

- Using process as an abstraction of running program:
 - Necessary information (environment) of execution
 - Memory, Hardware and OS contexts
- Process from OS perspective:
 - PCB and process table
- OS $\leftarrow \rightarrow$ Process interactions
 - System calls
 - Exception / Interrupt

References

- Modern Operating System (3rd Edition)
 - Section 2.1
- Operating System Concepts (8th Edition)
 - Section 3.1