

Process Management

Process Alternative - Threads

Lecture 5

Overview

■ **Threads**

- Motivation
- Basic Idea

■ **Threads models**

- Kernel vs User Thread
- Hybrid model

■ **Thread in Unix**

- POSIX thread:
 - Create, Exit, Synchronization
 - Exploration

Motivation for Thread

- Process is expensive:
 - Process creation under the `fork()` model:
 - Duplicate memory space
 - Duplicate most of the process context, etc.
 - Context switch:
 - Requires saving/restoration of process Information
- It is hard for independent processes to communicate with each other:
 - Independent memory space → No easy way to pass information
 - Requires Inter-Process Communication (IPC)

Motivation for Thread (cont)

- Thread was invented to overcome the problems with process model
 - Started out as a "quick hack" and eventually matured to be very popular mechanism
- **Basic Idea:**
 - A traditional process has a single thread of control:
 - Only one instruction of the whole program is executing at any time
 - We "simply" add more threads of control to the same process:
 - Multiple parts of the programs is executing at the same time conceptually

Threads of control: Illustration

- Suppose we are preparing lunch, which consists of the following tasks:
 - Steam rice
 - Fry fish
 - Cook soup
- A pseudo-C program:

```
int main()
{
    steamRice( twoBowls );
    fryFish( bigFish );
    cookSoup( cornSoup );


    printf( "Lunch READY!!\n" );
    return 0;
}
```

Threads of control: Illustration (cont)

- Process with a single thread will go through the functions sequentially

```
int main()
{
    steamRice( twoBowls );
    fryFish( bigFish );
    cookSoup( cornSoup );

    printf( "Lunch Ready\n" );
    return 0;
}
```



- Suppose the tasks are independent:
 - Let's try to have multiple threads of control

Multiple "threads of control" with `fork()`

```
int result;
result = fork();
if (result != 0){
    result = fork();
    if (result != 0){
        steamRice( twoBowls );
    } else {
        fryFish( bigFish );
    }
} else {
    cookSoup( cornSoup );
}
...  ...
```

PC 1

```
int result;
result = fork();
if (result != 0){
    result = fork();
    if (result != 0){
        steamRice( twoBowls );
    } else {
        fryFish( bigFish );
    }
} else {
    cookSoup( cornSoup );
}
...  ...
```

PC 2

We effectively
have two
"threads" of
control after
`fork()`

New approach: Multi-Threading


```
int main()
{
    Thread 1 do { steamRice( twoBowls ); }

    Thread 2 do { fryFish( bigFish ); }

    Thread 3 do { cookSoup( cornSoup ); }

    Wait for all threads to finish;

    printf( "Lunch Ready\n" );
    return 0;
}
```




■ Note:

- Pseudocode ☺
- The three threads are concurrent

Process and Thread

- A single process can have multiple threads:
 - known as **multithreaded process**
- Threads in the same process shares:
 - **Memory Context:** Text, Data, Heap
 - **OS Context:** Process id, other resources like files, etc.
- Unique information needed for each thread:
 - Identification (usually **thread id**)
 - Registers (General purpose and special)
 - “Stack” (more about this later)

Process and thread: Illustration



- Taken from Operating System Concepts (7th Edition) by Silberschatz, Galvin & Gagne, published by Wiley

Process Context Switch VS Thread Switch

- Process context switch involves:
 - OS Context
 - Hardware Context
 - Memory Context
- Thread switch within the same process involves:
 - Hardware context:
 - Registers
 - "Stack" (actually just changing FP and SP registers)
- Thread is much “lighter” than process:
 - a.k.a. ***lightweight process***

Threads: Benefits

Economy:

- Multiple threads in the same process requires much less resources to manage compared to multiple processes

Resource sharing:

- Threads share most of the resources of a process
- No need for additional mechanism for passing information around

Responsiveness:

- Multithreaded programs can appear much more responsive

Scalability:

- Multithreaded program can take advantage of multiple CPUs

Thread: Problems

■ System Call Concurrency

- Parallel execution of multiple threads → parallel system call possible
 - Have to guarantee correctness and determine the correct behavior

■ Process Behavior:

- Impact on process operations
- Example:
 - `fork()` duplicate process, how about threads?
 - If a single thread executes `exit()`, how about the whole process?
 - If a single thread calls `exec()`, how about other threads?

Thread Models

What are the ways to support threads?

Two Major Thread Implementations


■ User Thread:

- Thread is implemented as a **user library**
 - A runtime system (in the process) will handle thread related operation
- Kernel is **not aware** of the threads in the process

■ Kernel Thread:

- Thread is implemented in the OS
 - Thread operation is handled as system calls
- Thread-level scheduling is possible:
 - Kernel schedule by threads, instead of by process
- Kernel may make use of threads for its own execution

User Thread: Illustration



- Adapted from Modern Operating System Concepts (3rd Edition) by Andrew Tanenbaum, published by Pearson

User Thread: Pros and Cons


■ Advantages:

- Can have multithreaded program on ANY OS
- Thread operations are just library calls
- Generally more configurable and flexible
 - e.g., customized thread scheduling policy

■ Disadvantages:

- OS is not aware of threads, scheduling is performed at process level
 - One thread blocked → Process blocked → All threads blocked
 - Cannot exploit multiple CPUs!

Kernel Thread: Illustration



- Adapted from Modern Operating System Concepts (3rd Edition) by Andrew Tanenbaum, published by Pearson

Kernel Threads: Pros and Cons

■ Advantages:

- Kernel can schedule on thread levels:
 - More than 1 thread in the same process can run simultaneously on multiple CPUs


■ Disadvantages:

- Thread operations is now a system call!
 - Slower and more resource intensive
- Generally less flexible:
 - Used by all multithreaded programs
 - If implemented with many features:
 - Expensive, overkill for simple program
 - If implemented with few features:
 - Not flexible enough for some programs


Hybrid Thread Model

- Have both Kernel and User threads
 - OS schedule on kernel threads only
 - User thread can bind to a kernel thread

- Offer great flexibility
 - Can limit the concurrency of any process / user



Hybrid Model Example: Solaris



Threads on Modern Processor

- Threads started off as a software mechanism
 - User space library → OS aware mechanism
- There are hardware support on modern processors:
 - Essentially supply multiple sets of registers (GPRs, and special registers) to allow threads to run natively and in parallel on the same core
 - known as **simultaneous multi-threading (SMT)**
 - Example: **Hyperthreading** on Intel processor

POSIX Threads

A popular thread API

POSIX Threads: ***pthread***

- Standard defined by the IEEE
 - Supported by most Unix variants
- Defines the API as well as the behavior
 - However, **implementation** is not specified
 - So, ***pthread*** can be implemented as user / kernel thread
- Will show a few example to highlight the differences between process and thread only

Basics of pthread

- Header File:
 - `#include <pthread.h>`
- Compilation (flag is system dependent):
 - `gcc xxxx.c -lpthread`
- Useful datatypes:
 - `pthread_t` : Data type to represent a thread id (TID)
 - `pthread_attr`: Data type to represents attributes of a thread

pthread Creation Syntax

```
int pthread_create(  
    pthread_t* tidCreated,  
    const pthread_attr_t* threadAttributes,  
    void* (*startRoutine) (void*),  
    void* argForStartRoutine );
```

- Returns (0 = success; !0 = errors)
- Parameters:
 - **tidCreated**: Thread Id for the created thread
 - **threadAttributes**: Control the behavior of the new thread
 - **startRoutine**: Function pointer to the function to be executed by thread
 - **argForStartRoutine**: Arguments for the `startRoutine` function

pthread Termination Syntax

```
int pthread_exit( void* exitValue );
```

- Parameters:
 - **exitValue**: Value to be returned to whoever synchronize with this thread (more later)
- If ***pthread_exit()*** is not used, a *pthread* will terminate automatically at the end of the **startRoutine**
 - If a “return XYZ;” statement is used, then “XYZ” is captured as the **exitValue**
 - Otherwise, the **exitValue** is not well defined

pthread Creation & Termination: Example

```
//header files not shown

void* sayHello( void* arg )
{
    printf("Just to say hello!\n");
    pthread_exit( NULL );
}

int main()
{
    pthread_t tid;

    pthread_create( &tid, NULL, sayHello, NULL );
    printf("Thread created with tid %i\n", tid);

    return 0;
}
```

Function to be executed by a pthread

Pthread Termination

Pthread Creation

pthread: Sharing of memory space

```
//header files not shown
int globalVar; Variable shared between pthreads

void* doSum( void* arg )
{
    int i;
    for (i = 0; i < 1000; i++)
        globalVar++;
}

int main()
{
    pthread_t tid[5];      //5 threads id
    int i;
    for (i = 0; i < 5; i++)
        pthread_create( &tid[i], NULL, doSum, NULL );
    printf("Global variable is %i\n", globalVar);
    return 0;
}
```

Using a shared variables

- What is the sum that we get (or should get)?

pthread Simple Synchronization - Join

```
int pthread_join( pthread_t threadID,  
                    void **status );
```

- To wait for the termination of another **pthread**:
- Returns(0 = success; !0 = errors)
- Parameters:
 - **threadID**: TID of the **pthread** to wait for
 - **status**: Exit value returned by the target **pthread**

Pthread: Sharing of memory space V2.0

```
//header files not shown
int globalVar;

void* doSum( void* arg )
{ //same as before }

int main()
    pthread_t tid[5];      //5 threads id
    int i;
    for (i = 0; i < 5; i++)
        pthread_create( &tid[i], NULL, doSum, NULL );

    //Wait for all threads to finish
    for (i = 0; i < 5; i++)
        pthread_join( tid[i], NULL );
    printf("Global variable is %i\n", globalVar);
    return 0;
}
```

Pthread
Synchronization

Pthread: A lot more!

- There are more interesting stuff about ***pthread***:
 - Yielding (giving up CPU voluntarily)
 - Advanced synchronization
 - Scheduling policies
 - Binding to kernel threads
 - Etc.
- As we cover new topics, you can explore the ***pthread*** library to see the application!

Summary

- Thread:
 - Motivation
 - Difference between thread and process
- Thread Models:
 - User vs Kernel thread
- Simple introduction to *pthread* library

Reference

- Modern Operating System (3rd Edition)
 - Chapter 2.2
- Operating System Concepts (7th Edition)
 - Chapter 4
- Linux Pthread Man Pages
 - “man pthread.....” for more