

CS2106 Operating Systems
Tutorial 10 Suggested Solutions
File Abstraction

1. [Working Set Model] Working set model allows OS to make better decision for frame allocation among processes. Let us consider several aspects of the idea in this question.
 - a. [Basic Idea] Given the following memory reference string (from lecture 9 – page replacement section):

Time	1	2	3	4	5	6	7	8	9	10	11	12
Page	2	3	2	1	5	2	4	5	3	2	5	2

Given that $W(T, \Delta)$ gives the working set at time T with a window size of length Δ , give the following:

$$W(9, 3), W(11, 3), W(9, 4), W(11, 4)$$

- b. [Application] Suppose we know the $W(T, \Delta)$ value for all processes, how can OS use this information? Hint: think about virtual memory + demand paging. Demand paging refers to the idea that OS only loads a page into RAM upon page fault instead of trying to predict and pre-load pages for processes.

(Exploration Question) Using the same idea, is there any easy way to dynamically adjust the Δ value?

- c. [Implementation] Consider the following "mysterious" algorithm:

1. Every Page Table Entry has an additional K-bit mysterious value, $M_0, M_1, \dots M_{K-1}$. All K-bit are initialized to '0'.
2. Whenever a page P is referenced, its corresponding Mysterious value are updated as follows:
 - a. $M_{K-1} \leftarrow M_{K-2}$
 - b. ...
 - c. $M_2 \leftarrow M_1$
 - d. $M_1 \leftarrow M_0$
 - e. $M_0 \leftarrow 1$
3. All other Non-P pages' Mysterious value are updated as follows:
 - a. $M_{K-1} \leftarrow M_{K-2}$
 - b. ...
 - c. $M_2 \leftarrow M_1$
 - d. $M_1 \leftarrow M_0$
 - e. $M_0 \leftarrow 0$

What does the mysterious value represents? If the above algorithm is handled by the hardware, what should be done every K memory accesses?

ANS:

a.

$$W(9, 3) = \{ 3, 4, 5 \}$$

$$W(11, 3) = \{ 2, 3, 5 \}$$

$$W(9, 4) = \{ 2, 3, 4, 5 \}$$

$$W(11, 4) = \{ 2, 3, 5 \}$$

b.

$W()$ gives the number of frames required by a process. So, we can utilize the $W()$ value to:

- When a process become active (e.g. from blocked \square running), load all working set pages into frames for that process. Similarly, when the process become inactive (e.g. from running \square blocked), we start to migrate those pages from physical frame into secondary storage.
- We can also use the total $W()$ values among all "active" processes, if the total exceeds the physical frame number, we may want to stop allowing more process to become runnable as that may induce thrashing. It is common to refer to the two sets of processes the "active set" and the "inactive set" in this context.

(Optional) The second point above suggests that if the CPU is well utilized and page fault activity is low \square **Δ value is well chosen** as the pages in memory are good representation of working set pages for all processes. If the page fault activity is very high (coupled with low CPU utilization) \square **Δ may be too small**. Similarly, if both page fault activity and CPU utilization are low \square **Δ may be too large** (which prevented more processes to be runnable to utilize the CPU).

c.

The mysterious value indicates whether a page has been used in K memory accesses window. Every K-accesses, the hardware should trigger an interrupt to bring in the OS so that the OS can take note of the W() (essentially all pages with non-zero K-bit).

2. [Wrapping File Operation] File operations are very expensive in terms of time. There are several reasons: a) As we learned in the lecture, each file operation is a system call, which requires an execution mode change (user \square kernel); b) Secondary storage has high latency access time.

This leads to a strange phenomenon: it is generally true that the total time to perform 100 file operations for 1 item each is **much longer** than performing a single file operation for 100 items instead. e.g. write one byte 100 times takes longer than writing 100 bytes in one go.

So, most high level programming languages provide **buffered file operations** that wraps around the primitive file operations. The buffered version essentially maintains an internal intermediate storage in memory (i.e. buffer) to store user read/write values from/to the file. For example, a **buffered write operation** will wait until the internal memory buffer is full before doing a large one time file write operation to **flush** the buffer content into file.

- a. [Generalization] Give one or two examples of buffered file operations found in your favorite programming language(s). Other than the "chunky" read/write benefit, are there any other additional features provided by those high level buffered file operations?
- b. [Application] Take a look at the given "**weird.c**" source code. Compiles and performs the following experiments: Change the trigger value from 100, 200, ... until you see values printed on screen **before the program crashes**. Can you explain both the behavior and the significance of the "trigger" value? If you add a new line character "\n" to the the **printf()** statement, how does the output pattern changes? How can this information be useful?
- c. [Design] Give **a high level pseudo code** to provide buffered file read operation. Use the following "function header" as a starting point:

```
BufferedFileRead( File, outputArray, arraySize )
//Read "arraySize" items from "File" and place in
// the "outputArray"
```

ANS:

a.

C: printf, scanf, fprintf, fscanf

Java: FileInputStream, FileOutputStream

C++: "<<", ">>" stream operators

Common additional features: error checking, packing/unpacking of datatypes (e.g. low level file operations are usually operating in bytes, it is useful to be able to directly operate on other

datatype (int, float, etc) without worrying about how to translate the value into/from human readable string).

b.

- printf buffers the user output until the internal buffer is full before actually output to the screen.
- The trigger indicates the probable size of the internal buffer.
- If a newline character is added, the output is performed "immediately".
- If printf() or similar is used as a debugging mechanism, the buffered output sequence may confuse the coder. e.g. in the original "weird.c", the program can crash without showing any printout, which can easily lead to the wrong conclusion ("the while loop is not executed!").

c.

Only key points are mentioned.

```
BufferedFileRead( File, outputArray, requestSize )    //pseudo code
// Read arraySize bytes from File, place the file content in outputArray

//Assume we have an internal buffer "Buffer" with "size" and "availableItems" attributes
// The buffer can be implemented as a circular array


If Buffer.availableItems < requestSize          //not enough!
    read(File, Buffer, Buffer.size - Buffer.availableItems) //low level file reading
    Buffer.availableItems = Buffer.size //Buffer is full now

    outputValues □ Buffer( requestSize )    //copy the user requested items over to output
    buffer.availableItems -= requestSize      // assume requestSize <= buffer.size

//Scenarios not considered for the above simple key points:
//1. When requestSize is >> Buffer.size
//2. When the file "File" is exhausted (i.e. cannot refill buffer anymore)

//These scenarios can be easily supported by building on the key points above.
```

3. [Open File Table] Below is the illustration taken from lecture 10:



Discuss how this organization helps OS to handle the following scenarios. Your answer should refer to the relevant structure(s) if possible.

- Process A tries to open a file that is currently being written by Process B.
- Process A tries to use a bogus file descriptor in a file-related system call.
- Process A can never "accidentally" access files opened by Process B.
- Process A and Process B reads from the same file. However, their reading should not affect each other.
- Redirect Process A's standard input / output, e.g. "`a.out < test.in > test.out`". (Hint: *nix processes has 3 default file descriptors: 0 = stdin, 1 = stdout, 2 = stderr)

ANS:

- OS uses the Open File Table to check for existing opened file. Since the file is already opened for reading, it can reject the file open system call from process A.
- Since Process A passed the file descriptor (fd for short) to OS as parameter, OS can check whether that particular entry is valid (or even exists) in the PCB of A. If the fd is out of range, non-existent etc, OS can reject the file-related system calls made by Process A.
- Since the fd index into process specific PCB, there is no way Process A can access Process B's file descriptor table.
- As discussed in lecture, Process A and Process B can have their own **fds**, which refers to two distinct locations in the open file table. Each entry of the open file table keep track of the current location separately. This enables Process A and Process B to read from the same file independently.
- Key points: There are 3 standard file descriptors for every program in Unix (0 = stdin, 1 = stdout, 2 = stderr). These are "opened" automatically and linked to the


corresponding files. Note that screen (terminal), keyboard are represented as *special files* in unix.

So, for all file redirections, it is a simple question of:

- a. Opening and possibly creating the file.
- b. Replace the corresponding file descriptor to point to the entry from (1) in the open file table.

The following illustrate the last case:

a. `out < test.in > test.out`



4. [Cover if time permits - Understanding directory permission] In *nix system, a directory has the same set of permission settings as a file. For example:

```
sooyj@sunfire [13:22:52] ~/tmp/Parent $ ls -l
total 8
drwxr-x--x 2 sooyj compsc 4096 Nov 8 13:22 Directory
sooyj@sunfire [13:22:53] ~/tmp/Parent $
```

You can see that directory **Directory** has the read, write, execute permission for owner, but only execution permission for group and others. It is easy to understand the permission bits for a regular file (read = can only access, write = can modify, execute = can execute this file). However, the same cannot be said for the directory permission bits.

Let's perform a few experiments to understand the permission bits for a directory.

Setup:

- Unzip **DirExp.zip** on any *nix platform (Solaris, Mac OS X included).
- Change directory to the **DirExp/** directory, there are 4 subdirectories with the same set of files. Let's set their permission as follows:

<code>chmod 700 NormDir</code>	NormDir is a normal directory with read, write and execute permissions.
<code>chmod 500 ReadExeDir</code>	ReadExeDir has read and execute permission.
<code>chmod 300 WriteExeDir</code>	WriteExeDir has write and execute permission.
<code>chmod 100 ExeOnlyDir</code>	ExeOnlyDir has only execute permission.

Perform the following operations on each of the directory and note down the result. Make sure you are at the **DirExp/** directory at the beginning. DDDD is one of the subdirectories.

- a. Perform "**ls -l DDDD**".
- b. Change into the directory using "**cd DDDD**".
- c. Perform "**ls -l**".
- d. Perform "**cat file.txt**" to read the file content.
- e. Perform "**touch file.txt**" to modify the file.
- f. Perform "**touch newfile.txt**" to create a new file.

Can you deduce the meaning of the permission bits for directory after the above? Can you use the "directory entry" idea to explain the behavior?

ANS:

	ReadExeDir	WriteExeDir	ExeOnlyDir
a	ok	nope	nope
b	ok	ok	ok
c	ok	nope	nope
d	ok	ok	ok
e	ok	ok	ok

f	nope	ok	nope
----------	------	----	------

The responses may seem arbitrary unless you apply the understanding that "Directory == the list of directory entries (file/subdir)".

Then, the permission can be understood as:

Read = Can you read this list? (impact: ls, <tab> auto-completion)

Write = Can you change this list? (impact: create, rename, delete file/subdir). Note the interaction with the execute permission bit.

Execute = Can you use this directory as your working directory? (impact: cd).

Since modifying the directory entries (i.e. with write permission) requires us to set the current working directory, execute bit is needed for the write-related operations under the target directory.

Some interesting scenarios:

- a. Directory permission is independent from the file permission. So, you still can modify a file under a "read only" directory if the file allows write.
 - b. If you want to allow outsider to access a particular deeply nested file, e.g. A/B/C/D/file, you **only need** execute bit on A, B, C, D directory (i.e. read permission is not needed). This is a great way to hide the content of the directory and only allow access to specific file given the full pathname.
-

Question for your own exploration

1. Explain the following concepts in your own words clearly; the shorter, the better! You explanation should be easily understandable for non-CS2106 students.
 - a. What is a file?

A file is the smallest amount of information that can be written to secondary memory. It is a named collection of data, used for organizing secondary memory.
 - b. Name and describe the two basic classifications of files.

Text files: Files that contain text. Each byte is an ANSI character or each 2 bytes is a Unicode character.
Binary files: The bytes in a binary file do not necessarily contain characters. These files require a special interpretation.
 - c. Distinguish between a file type and a file extension.

A file type is a description of the information contained in the file. A file extension is a part of the file name that follows a dot and identifies the file type.
 - d. What does it mean to open and close a file?

Operating systems keep a table of currently open files. The open operation enters the file into this table and places the file pointer at the beginning of the file. The close operation removes the file from the table of open files.
 - e. What does it mean to truncate a file?

Truncating a file means that all the information on the file is erased but the administrative entries remain in the file tables. Occasionally, the truncate operation removes the information from the file pointer to the end.
2. [Adapted from (SGG)] Why do many operating systems have a system call to “open” a file, rather than just passing a path to the read or write system calls each time?

ANS:

 - Either the read/write system calls return a handle to a file descriptor (like open normally does), or they do not and we pass a path to read/write each time.
 - In the former case, it would complicate the interface of the system calls, because they would then need to also accept a handle instead of a path somehow.
 - In the latter case, we would incur the cost of permission checking and path resolution each and every time the system call is made, which are non-trivial and expensive. Also, we lose the ability to have the OS track our offset within the file.
 - In either case, such an interface would not generalise well to file descriptors that do not come from paths on the filesystem, like pipes and anonymous sockets.
 - Some history: open, read, write and close were present in the original Unix OS from the start. Although pipes and sockets, etc., didn’t come around until much later, the design generalised well to those. See <https://unix.stackexchange.com/a/265627> , <http://www.read.seas.harvard.edu/~kohler/class-aosref/ritchie84evolution.pdf>