

CS2106 Operating Systems
25/26 Semester 1
Tutorial 2
Process Abstraction in Unix

1. [Process Creation Recap - Taken from AY18/19 S1 Midterms] Each of the following cases insert zero or more lines at the Point α and β . Evaluate whether the described behaviour is correct or incorrect. (Note that `wait()` does not block when a process has no children.)

	C code:
00	<code>int main() {</code>
01	<code>//This is process P</code>
02	<code>if (fork() == 0) {</code>
03	<code>//This is process Q</code>
04	<code>if (fork() == 0) {</code>
05	<code>//This is process R</code>
06	<code>.....</code>
07	<code>return 0;</code>
08	<code>}</code>
09	<code><Point α></code>
10	<code>}</code>
11	<code><Point β></code>
12	<code>return 0;</code>
13	<code>}</code>
14	

Point α	Point β	Behavior
<code>Nothing</code>	<code>wait(NULL);</code>	Process Q <i>always</i> terminate before P. Process R can terminate at any time w.r.t. P and Q.
<code>wait(NULL);</code>	<code>nothing</code>	Process Q <i>always</i> terminate before P. Process R can terminate at any time w.r.t. P and Q.
<code>execl(valid executable....);</code>	<code>wait(NULL);</code>	Process Q <i>always</i> terminate before P. Process R can terminate at any time w.r.t. P and Q.
<code>wait(NULL);</code>	<code>wait(NULL);</code>	Process P never terminates .

2. [Behavior of **fork**] The C program below attempts to highlight the behavior of the **fork()** system call:

C code:

```
int dataX = 100;

int main( ) {
    pid_t childPID;

    int dataY = 200;
    int* dataZptr = (int*) malloc(sizeof(int));

    *dataZptr = 300;

    //First Phase
    printf("PID[%d] | X = %d | Y = %d | Z = %d |\n",
getpid(), dataX, dataY, *dataZptr);

    //Second Phase
    childPID = fork();
    printf("*PID[%d] | X = %d | Y = %d | Z = %d |\n",
getpid(), dataX, dataY, *dataZptr);

    dataX += 1;
    dataY += 2;
    (*dataZptr) += 3;
    printf("#PID[%d] | X = %d | Y = %d | Z = %d |\n",
getpid(), dataX, dataY, *dataZptr);

    //Insertion Point

    //Third Phase
    childPID = fork();
    printf("**PID[%d] | X = %d | Y = %d | Z = %d |\n",
getpid(), dataX, dataY, *dataZptr);

    dataX += 1;
    dataY += 2;
    (*dataZptr) += 3;
    printf("##PID[%d] | X = %d | Y = %d | Z = %d |\n",
getpid(), dataX, dataY, *dataZptr);

    return 0;
}
```

The code above can also be found in the given program "**ForkTest.c**". Please run it on your system before answering the questions below.

- a. What is the difference between the 3 variables: **dataX**, **dataY** and the memory location pointed by **dataZptr**?
- b. Focusing on the messages generated by second phase (they are prefixed with either "*" and "#"), what can you say about the behavior of the **fork()** system call?
- c. Using the messages seen on your system, draw a **process tree** to represent the processes generated. Use the process tree to explain the values printed by the child processes.
- d. Do you think it is possible to get different ordering between the output messages, why?
- e. Can you point how which pair(s) of messages can never swap places? i.e. their relative order is always the same?
- f. If we insert the following code at the insertion point:

Sleep Code
<pre>if (childPID == 0) { sleep(5); //sleep for 5 seconds }</pre>

- How does this change the ordering of the output messages? State your assumption, if any.
- g. Instead of the code in (f), we insert the following code at the insertion point:

Wait Code
<pre>if (childPID != 0) { wait(NULL); // NULL means we don't care // about the return result }</pre>

How does this change the ordering of the output messages? State your assumption, if any.

3. [Parallel Computation] Even with the crude synchronization mechanism, we can solve programming problems in new (and exciting) ways. We will attempt to utilize multiple processes to work on a problem simultaneously in this question.

You are given two C source code "**Parallel.c**" and "**PrimeFactors.c**". The "**PrimeFactors.c**" is a simple prime factorization program. "**Parallel.c**" use the "**fork()**" and "**exec()**" combination to spawn off a new process to run the prime factorization.

Let's setup the programs as follows:

1. Compiles "**PrimeFactors.c**" to get a executable with name "**PF**": **gcc PrimeFactors.c -o PF**
2. Compiles "**Parallel.c**": **gcc Parallel.c**

Run the **a.out** generated from step (2). Below is a sample session:

```
$ ./a.out
1024
1024 has 10 prime factors // note: not unique prime factors
```

If you try large prime numbers, e.g. 111113111, the program may take a while.

Modify only Parallel.c such that we can now initiate prime factorization on [1-9] user inputs simultaneously. More importantly, we want to report result as soon as they are ready regardless of the user input order. Sample session below:

```
$ ./a.out
5 // 5 user inputs
44721359
99999989
9
111113111
118689518
9 has 2 prime factors // Results
118689518 has 3 prime factors
44721359 has 1 prime factors
99999989 has 1 prime factors
111113111 has 1 prime factors
```

Note the order of the result may differ on your system. Most of time, they should follow roughly the computation time needed (composite number < prime number and small number < large number). Two simple test cases are given **test1.in** and **test2.in** to aid your testing. If you are using a rather powerful machine (e.g. the SoC Compute Cluster), you can use the **test3.in** to provide a bit more grind.

Most of what you need is already demonstrated in the original **Parallel.c** (so that this is more of a mechanism question rather than a coding question). You only need "**fork()**", "**exec1()**" and "**wait()**" for your solution.

After you have solved the problem, find a way to change your **wait()** to **waitpid()**, **what do you think is the effect of this change?**

For your own exploration:

4. (Process Creation) Given the following full program, give and explain the execution output. The source code **FF.c** is also given for your own test.

C code:

```
int factorial(int n) {
    if (n == 0) {
        fork();           // NOTE the change
        return 1;
    }
    return factorial(n-1) * n;
}

int main() {
    printf("fac(2) = %d\n", factorial(2));

    return 0;
}
```

5. (Process Creation) Consider the following sequence of instructions in a C program:

C code:

```
int x = 10;
int y = 123;

y = fork();
if (y == 0)
    x--;

y = fork();
if (y == 0)
    x--;

printf("[PID %d]: x=%d, y=%d\n", getpid(), x, y);
```

You can assume that the first process has process number **100** (and so **getpid()** returns the value **100** for this process), and that the processes created (in order) are **101,102** and so on.

Give:

- a) A possible final set of printed messages.
- b) An impossible final set of printed messages.

6. (Parent-Child Synchronization) Consider the following sequence of instructions in a C program:

C code:

```
int i;
pid_t cPidArray[3]; // an array of 3 child pids

for (i = 0; i < 3; i++) {
    cPidArray[i] = fork();
    if (cPidArray[i] == 0) {
        // do something
        printf("Child [%d] is done!\n", getpid());
        return 0; // exit
    }
}

// Code insert point here
printf("Parent [%d] is done!\n", getpid());
```

Similar to Q1, let's assume the first process has pid of 100, and that the processes created (in order) are **101,102** and so on.

Suppose we insert the following code fragments at the end of the program above, describe the effects on the synchronization / timing property of the program. If it helps, you can give a sample output to aid your explanation.

a) wait()

```
for (i = 0; i < 3; i++) {  
    wait(NULL);  
    printf("Parent: one child exited\n");  
}
```

b) waitpid()

```
for (i = 0; i < 3; i++) {  
    waitpid(cPidArray[i], NULL, 0);  
    printf("Parent: one child exited\n");  
}
```