

**CS2109S: Introduction to AI and Machine Learning**

# Lecture 8: Unsupervised Learning

18 March 2025

# Last time

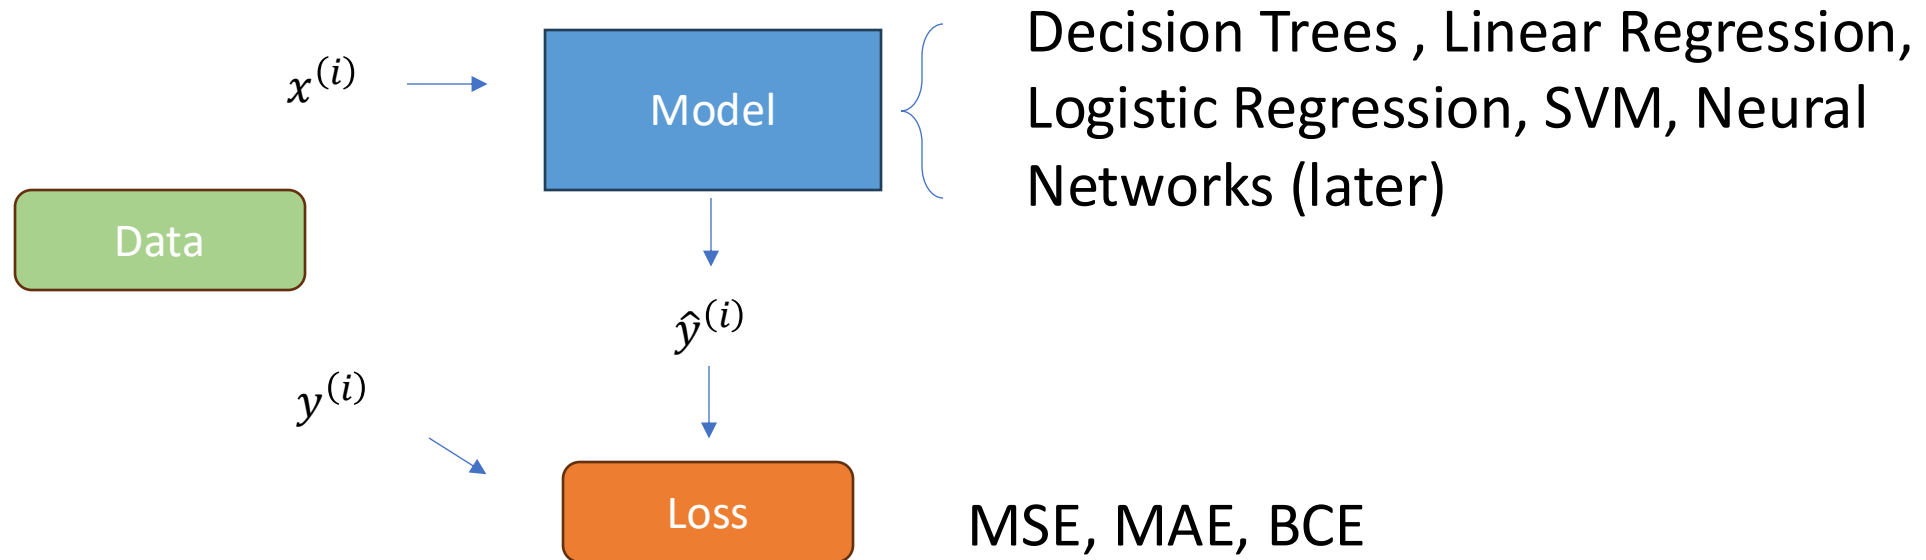
- Regularization
- Kernel method
- Support Vector Machine

# Outline

- Unsupervised Learning
- K-means clustering
  - Algorithm
  - Measuring the goodness of clusters
  - Picking the number of clusters
  - Variants
- Dimensionality Reduction
  - Singular Value Decomposition (SVD)
  - Principal Component Analysis (PCA)

# Supervised Learning

Given a set of  $N$  input-output pairs (training samples)  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(N)}, y^{(N)})\}$ , learn a model that makes predictions.



# Unsupervised Learning - Motivation

- We often have data where we don't have the ground truth outputs.
  - In some cases, obtaining the ground truth can be very expensive.
  - Examples:
    - Topic modelling: documents/images/videos, where we do not know the category like health, finance, ...
    - Preventative medicine: health data, where we do not know disease/no disease.
    - ...
- Human: much of our own learning happens by experiencing data with no ground truths.
- Examples:
  - Clustering, dimensionality reduction, image segmentation, GPT (contains unsupervised learning), ...

# Unsupervised Learning

Given a set of  $N$  data points  $\{x^{(1)}, \dots, x^{(N)}\}$ , learn patterns in the data.

## Types of unsupervised learning:

- **Clustering**: identify clusters in the data
- **Dimensionality reduction**: find a lower-dimensional representation of the data
- ...

# Clustering

Clustering is a type of unsupervised machine learning technique used to **group similar data points into clusters or groups**.

The goal is to organize a set of objects in such a way that objects within the same group (cluster) are more similar to each other than to those in other groups.

In clustering, **the number of clusters or groups is not predefined by the data**.

## **Common applications:**

- Data segmentation
- Anomaly detection

# Dimensionality Reduction

Dimensionality reduction is a technique to **reduce the number of features (variables or dimensions)** in a dataset while **retaining as much of the relevant information as possible**.

## Common applications:

- Visualizations
- Feature extractions / transformations



# Outline

- Unsupervised Learning
- **K-means clustering**
  - Algorithm
  - Measuring the goodness of clusters
  - Picking the number of clusters
  - Variants
- Dimensionality Reduction
  - Singular Value Decomposition (SVD)
  - Principal Component Analysis (PCA)

# Background: Distance

Suppose that we have two  $d$ -dimensional vectors,  $u$  and  $v$ .

The straight-line **Euclidean distance** is defined as  $\|u - v\| = \sqrt{\sum_{i=1}^d (u_i - v_i)^2}$

Machine learning models that rely only on distances between data points are called **distance-based models**.

(Do you recall similarity functions from kernel methods?)


# Example: Clothing Company

Suppose that

- You have data about your customers' **width** and **height**
- You want to mass-produce clothes of **K distinct sizes** such that the **fit for everyone is pretty good** (you only have K machines!)
- What would you do?

# Example: Clothing Company

Goal: Mass-produce clothes of **K distinct sizes** such that the **fit for everyone is pretty good**.



Group people together!


How many groups?

Up to us

This is called **clustering**.

# Example: Clothing Company

Goal: Mass-produce clothes of **K distinct sizes** such that the **fit for everyone is pretty good**.




# Centroid

A **centroid** is the average of a set of points  $x^{(i)}$  in a cluster.

For cluster  $j$  with  $N_j$  data points:


$$\mu_j = \frac{1}{N_j} \sum_{i=1}^{N_j} x^{(i)}$$

A centroid defines a cluster.



# How to <sup>Procedurally</sup> Group Data?

Find K=3 Clusters



Let's start with the **centroids**  
(which define the clusters)

Where should we put them (initially)?  
We don't know the "correct" grouping.

➤ Randomly put the centroids.

What are the groups?  
How should we assign the points to the group?

➤ Assign points to the nearest centroids.

Are the centroids "centroids"?  
i.e., the center of the groups

➤ No! Need to recompute.

Are all points correctly grouped?  
i.e., assigned to the nearest centroids

➤ No! need to reassign.

# Assignments

Given  $N$  data points and  $K$  clusters. Define the variables

$$c^{(1)}, c^{(2)}, \dots, c^{(N)} \in \{1, \dots, K\},$$

which represent the assignments to a cluster for each data point. These variables are updated in the K-Means algorithm together with the centroids.




# K-Means

- Randomly initialize  $K$  centroids:  
 $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  
 $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$   
assigned to cluster  $k$


# K-Means

- Randomly initialize  $K$  centroids:  
 $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  
 $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$   
assigned to cluster  $k$




# K-Means

- Randomly initialize  $K$  centroids:  
 $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$  assigned to cluster  $k$




# K-Means

- Randomly initialize  $K$  centroids:  $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$  assigned to cluster  $k$




# K-Means

- Randomly initialize  $K$  centroids:  
 $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  
 $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$   
assigned to cluster  $k$




# K-Means

- Randomly initialize  $K$  centroids:  
 $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  
 $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$   
assigned to cluster  $k$




# K-Means

- Randomly initialize  $K$  centroids:  $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$  assigned to cluster  $k$



# K-Means

- Randomly initialize  $K$  centroids:  
 $\mu_1, \dots, \mu_K$
- Repeat until convergence:
  - For  $i = 1, \dots, N$ :
    - $c^{(i)} \leftarrow$  index of cluster centroid  
 $(\mu_1, \dots, \mu_K)$  closest to  $x^{(i)}$
  - For  $k = 1, \dots, K$ :
    - $\mu_k \leftarrow$  centroid of data points  $x^{(i)}$   
assigned to cluster  $k$




No more change: **converged!**




# K-Means: Measuring Quality of Clusters

We can measure the **distortion** of the clusters: the average square distance of each sample to its assigned centroid.

$$J(c^{(1)}, c^{(2)}, \dots, c^{(N)}, \mu_1, \mu_2, \dots, \mu_K) = \frac{1}{N} \sum_{i=1}^N \|x^{(i)} - \mu_{c^{(i)}}\|^2$$




Converged




Converged, but lower distortion

# K-Means: Local Optima



Ideal outcome



Possible outcome

Outcome is a stable configuration (the centroids will not move)

# K-Means: Convergence

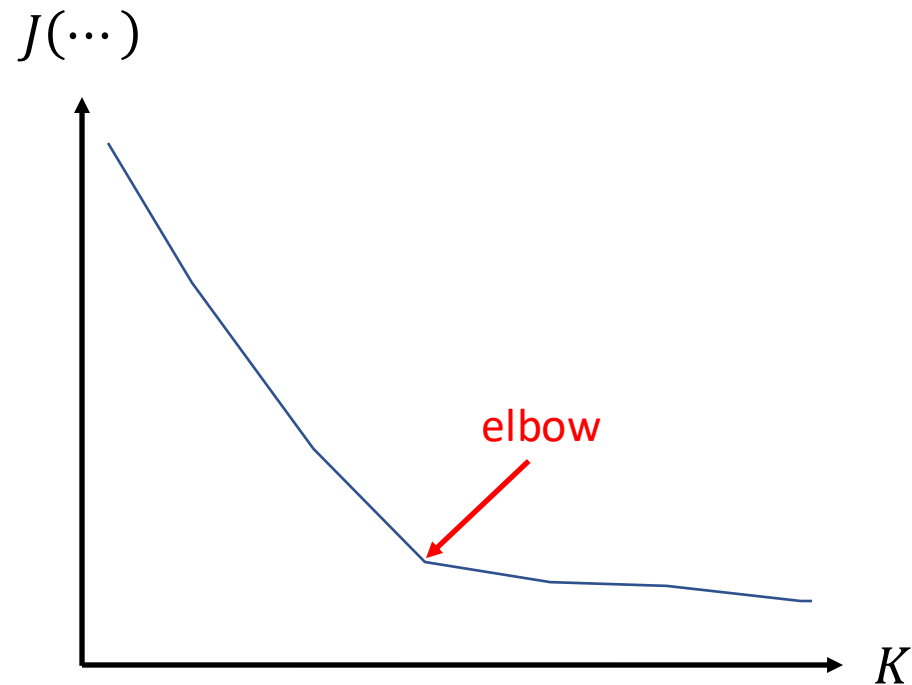
## Theorem (Informal):

Each step in the K-Means algorithm never increases **distortion**.

More on convergence will be discussed in **Tutorial 7**.



# K-Means: Picking the Number of Clusters

Application dependent




## Elbow Method

Disclaimer: some data may not have elbow, or have multiple elbows.



# K-Means: Variants

- Pick  $K$  initial centroids randomly from the points in the data
- K-Medoids: pick the data points that are closest to the centroids, and use them as the centroids.
  - “Snap” the centroids to the nearest data points



# Clustering vs Classification

Aspect	Classification	Clustering
Type of Feedback	Supervised Learning	Unsupervised Learning
Input	Input-output pairs $\{(x^{(i)}, y^{(i)})\}$	Input only $\{x^{(i)}\}$
Output	A model: $h(x) \rightarrow \hat{y}$	Clustering/grouping Example: $x^{(1)}$ is assigned to group 5
Methods	Hyperplane-based (e.g., SVM), probability-based (e.g., logistic regression)	Distance-based
Number of • Classes (for classification) • Clusters (for clustering)	Defined by the dataset	Up to us*

\*) the ones covered in this course

# Poll Everywhere

What is an **incorrect** statement about K-means?

- A: The K parameter is chosen by the user.
- B: The centroids are always located at the given data points.
- C: K-means minimizes the distortion.
- D: Multiple restarts of K-means give different solutions.

# Poll Everywhere

What is an **incorrect** statement about K-means?

A: The K parameter is chosen by the user.

**B: The centroids are always located at the given data points.**

C: K-means minimizes the distortion.

D: Multiple restarts of K-means give different solutions.



# Outline

- Unsupervised Learning
- K-means clustering
  - Algorithm
  - Measuring the goodness of clusters
  - Picking the number of clusters
  - Variants
- Dimensionality Reduction
  - Singular Value Decomposition (SVD)
  - Principal Component Analysis (PCA)

# High-dimensional Features

Many machine learning problems have data with **high-dimensional features**.

For example:

- HD images have  $1280 \times 720 = 921,600$  features

Is there any problem with this?

# Curse of Dimensionality

## Theorem (Sample Complexity, Informal):

Number of samples  $N$  needed to learn a hypothesis class increases exponentially with the number of features  $d$ :

$$N = O(2^d)$$

This is often called **the curse of dimensionality**.

# High-dimensional Features

Many machine learning problems have data with **high-dimensional features**.


For example:

- HD images have  $1280 \times 720 = 921,600$  features


Machine learning model that takes in this high-dimensional input may suffer from the curse of dimensionality.

How to reduce the number of features?

# Key Idea: Remove Redundant Features




Feature  $x_1$  is important  
Feature  $x_2$  is redundant




One of the features can be  
removed if we change the basis

# Key Idea: Remove Redundant Features



Feature  $x_1$  is important  
Feature  $x_2$  is redundant



Feature  $x'_1$  is important  
Feature  $x'_2$  is redundant

# Recall: Data Matrix

$$X = \begin{matrix} & \text{Features} & & \\ & & & \\ \begin{matrix} \text{Data points} \\ \vdots \end{matrix} & \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_d^{(1)} \\ 1 & x_1^{(2)} & & x_d^{(2)} \\ 1 & \vdots & & \vdots \\ 1 & x_1^{(N)} & & x_d^{(N)} \end{bmatrix} & & \end{matrix}$$

This matrix is everything we have in the unsupervised setting.

# Recall: Data Matrix (Transposed)

$$X^T = \begin{bmatrix} 1 & 1 & \dots & 1 \\ x_1^{(1)} & x_1^{(2)} & \dots & x_1^{(N)} \\ \vdots & \vdots & \dots & \vdots \\ x_d^{(1)} & x_d^{(2)} & \dots & x_d^{(N)} \end{bmatrix}$$

Data points

Features

We work with this matrix in the following slides.



# Identifying Important Features in Data Matrix

How do we identify important features in the following data matrix?

$$X^T = \begin{bmatrix} 1 & 1 & \dots & 1 \\ x_1^{(1)} & x_1^{(2)} & \dots & x_1^{(N)} \\ \vdots & \vdots & \dots & \vdots \\ x_d^{(1)} & x_d^{(2)} & \dots & x_d^{(N)} \end{bmatrix}$$

Use tools from [linear algebra](#) to find **better** and **more compact basis**.

# Singular Value Decomposition (SVD)

Given a matrix  $A$ , we can decompose it into 3 matrices:  $U$ ,  $\Sigma$ , and  $V^T$ .


$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

# Background: Orthonormal Basis

**Orthonormality:** Two vectors  $a$  and  $b$  are called orthonormal if they are:

- Normalized:  $\|a\|^2 = a^T a = 1$  and  $\|b\|^2 = 1$ , and
- Orthogonal:  $a^T b = 0$ .

**Orthonormal basis:** A set of vectors such that every vector can be expressed as a unique linear combination of the vectors in the set.



# SVD: Existence

**Theorem:** Without loss of generality, let  $d > N$ . For any  $d \times N$  real-valued matrix  $\mathbf{A}$ , there exists a factorization  $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$  called SVD, such that:

- $\mathbf{U}$  is  $d \times d$  and has  $d$  orthonormal columns left singular vectors
- $\mathbf{\Sigma}$  is  $d \times N$  and is a diagonal matrix with  $\sigma_j \geq 0$  singular values
  - Ordering:  $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_N \geq 0$
- $\mathbf{V}$  is  $N \times N$  and has  $N$  orthonormal columns and rows right singular vectors

Why “singular”? Possibly from integral theory, however it’s unclear:  
On the Early History of the Singular Value Decomposition (G. W. Stewart)

# SVD: Schematic for $X^T$

$$A = U\Sigma V^T$$

$$X^T = \begin{bmatrix} | & | & \dots & | \\ x^{(1)} & x^{(2)} & \dots & x^{(N)} \\ | & | & \dots & | \end{bmatrix}$$

$$= U\Sigma V^T$$

$$= \begin{bmatrix} | & | & \dots & | \\ u^{(1)} & u^{(2)} & \dots & u^{(d)} \\ | & | & \dots & | \end{bmatrix} \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sigma_N \\ 0 & 0 & \dots & 0 \end{bmatrix} \begin{bmatrix} | & | & \dots & | \\ v^{(1)} & v^{(2)} & \dots & v^{(N)} \\ | & | & \dots & | \end{bmatrix}^T$$

$$d \times d$$

$$d \times N$$

$$N \times N$$

# SVD: Interpretation

$$\mathbf{X}^T = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

$$\mathbf{X}^T = \begin{bmatrix} | & | & \dots & | \\ x^{(1)} & x^{(2)} & \dots & x^{(N)} \\ | & | & \dots & | \end{bmatrix}$$

$$= \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

$$= \begin{bmatrix} | & | & \dots & | \\ u^{(1)} & u^{(2)} & \dots & u^{(d)} \\ | & | & \dots & | \end{bmatrix} \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sigma_N \\ 0 & 0 & \dots & 0 \end{bmatrix} \begin{bmatrix} | & | & \dots & | \\ v^{(1)} & v^{(2)} & \dots & v^{(N)} \\ | & | & \dots & | \end{bmatrix}^T$$

New (Orthonormal) Basis

Importance

Linear combination coefficients

# SVD: Example

$$X^T = U \Sigma V^T$$

$$X^T = \begin{bmatrix} x_1^{(1)} & x_1^{(2)} \\ x_2^{(1)} & x_2^{(2)} \end{bmatrix} = \begin{bmatrix} | & | \\ u^{(1)} & u^{(2)} \\ | & | \end{bmatrix} \begin{bmatrix} \sigma_1 & 0 \\ 0 & \sigma_2 \end{bmatrix} \begin{bmatrix} v_1^{(1)} & v_1^{(2)} \\ v_2^{(1)} & v_2^{(2)} \end{bmatrix}^T$$

$$= \begin{bmatrix} | & | \\ \sigma_1 u^{(1)} & \sigma_2 u^{(2)} \\ | & | \end{bmatrix} \begin{bmatrix} v_1^{(1)} & v_2^{(1)} \\ v_1^{(2)} & v_2^{(2)} \end{bmatrix}$$

$$= \begin{bmatrix} v_1^{(1)} \sigma_1 u^{(1)} + v_1^{(2)} \sigma_2 u^{(2)} & v_2^{(1)} \sigma_1 u^{(1)} + v_2^{(2)} \sigma_2 u^{(2)} \\ | & | \end{bmatrix}$$

Data points  $x^{(i)}$  are linear combinations of basis vectors  $u^{(1)}$  and  $u^{(2)}$

# SVD: Another Interpretation

$$X^T = U\Sigma V^T$$

$$X^T = \begin{bmatrix} \text{Face 1} & \text{Face 2} & \dots & \text{Face N} \end{bmatrix}$$

$$= U\Sigma V^T$$

$$= \begin{bmatrix} \text{Template 1} & \text{Template 2} & \dots & \text{Template M} \end{bmatrix} \begin{bmatrix} 45 & 0 & \dots & 0 \\ 0 & 10 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & 0.01 \\ 0 & 0 & \dots & 0 \end{bmatrix} \begin{bmatrix} 0.9 & -0.5 & \dots & \dots \\ -0.2 & 0.8 & \dots & \dots \\ \dots & \dots & \dots & \dots \\ -0.1 & 0.4 & \dots & \dots \end{bmatrix}$$

Templates
Template Importance
Combination coefficients

Face 1




Face 2





# Dimensionality Reduction via SVD

$$\mathbf{X}^T = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$$


**Key idea:** The  $N$  singular values tell us about the importance of the new basis vectors. **Remove** less important basis vectors.

**How?** Recall that singular values are ordered:  $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_N > 0$

We can set all singular values except the first  $r$  to 0.

$$\begin{bmatrix} \sigma_1 & 0 & \vdots & \dots & 0 \\ 0 & \sigma_2 & \vdots & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & \dots & \sigma_N & \\ 0 & 0 & \dots & 0 & \end{bmatrix}$$

$$r = 2$$

# SVD: Effect of Reduction

$$X^T = U \Sigma V^T$$

$$X^T = \begin{bmatrix} | & | & \dots & | \\ x^{(1)} & x^{(2)} & \dots & x^{(N)} \\ | & | & \dots & | \end{bmatrix}$$

$$= U \Sigma V^T$$

$$= \begin{bmatrix} | & | & \dots & | \\ u^{(1)} & u^{(2)} & \dots & u^{(d)} \\ | & | & \dots & | \end{bmatrix} \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sigma_N \\ 0 & 0 & \dots & 0 \end{bmatrix} \begin{bmatrix} | & | & \dots & | \\ v^{(1)} & v^{(2)} & \dots & v^{(N)} \\ | & | & \dots & | \end{bmatrix}^T$$

New (Orthonormal) Basis

Importance

Linear combination coefficients

# SVD: Effect of Reduction

$$\mathbf{X}^T = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

Recall: Data points  $x^{(i)}$  are linear combinations of basis vectors  $u^{(j)}, j \in \{1, \dots, d\}$

$$\mathbf{X}^T = \left[ \begin{array}{c|c|c} \begin{array}{c} | \\ v_1^{(1)}\sigma_1 u^{(1)} + v_1^{(2)}\sigma_2 u^{(2)} + \dots \\ | \end{array} & \begin{array}{c} | \\ v_2^{(1)}\sigma_1 u^{(1)} + v_2^{(2)}\sigma_2 u^{(2)} + \dots \\ | \end{array} & \begin{array}{c} | \\ \vdots \\ | \end{array} \\ \hline \begin{array}{c} | \\ \vdots \\ | \end{array} & \begin{array}{c} | \\ \vdots \\ | \end{array} & \begin{array}{c} | \\ \vdots \\ | \end{array} \end{array} \right]$$

After **dimensionality reduction** with  $r = 2$ , data points  $x^{(i)}$  are linear combinations of  $u^{(j)}, j \in \{1, 2\}$

$$\tilde{\mathbf{X}}^T = \left[ \begin{array}{c|c|c} \begin{array}{c} | \\ v_1^{(1)}\sigma_1 u^{(1)} + v_1^{(2)}\sigma_2 u^{(2)} \\ | \end{array} & \begin{array}{c} | \\ v_2^{(1)}\sigma_1 u^{(1)} + v_2^{(2)}\sigma_2 u^{(2)} \\ | \end{array} & \begin{array}{c} | \\ \vdots \\ | \end{array} \\ \hline \begin{array}{c} | \\ \vdots \\ | \end{array} & \begin{array}{c} | \\ \vdots \\ | \end{array} & \begin{array}{c} | \\ \vdots \\ | \end{array} \end{array} \right]$$

# SVD: Dimensionality Reduction

$$r < N < d$$

Recap: Dimensionality reduction starts at the singular value matrix:

$$\Sigma = \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sigma_N \\ 0 & 0 & \dots & 0 \end{bmatrix} \quad \tilde{\Sigma} := \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sigma_N \\ 0 & 0 & \dots & 0 \end{bmatrix}$$

The truncated matrix  $\tilde{\Sigma}$  implies a smaller basis matrix:

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(d)} \\ | & | & & | \end{bmatrix} \quad \tilde{U} := \begin{bmatrix} | & | \\ u^{(1)} & u^{(2)} \\ | & | \end{bmatrix} \in d \times r$$


Can we use  $\tilde{U}$  to compress vectors?

# Background: Dot Product and Projection

Let  $u, v$  be two vectors.

$$u^T v = \sum_{j=1}^d u_j v_j = z \|v\|$$

The dot product gives the **length** of the projection of  $u$  onto  $v$  times the norm of  $v$ .



$$z = \begin{bmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{bmatrix}^T \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \frac{1}{\sqrt{2}}$$


# Projection using New Basis

The matrix  $\tilde{U}$  defines 2-dimensional subspace of the feature space. (In general, r-dimensional)

$$\tilde{U} = \begin{bmatrix} | & | \\ \textcolor{red}{u}^{(1)} & \textcolor{green}{u}^{(2)} \\ | & | \end{bmatrix}$$

$\nearrow$   
 $d$ -dimensional vector


Other dimensions ( $d - 2$ )



For any feature vector  $x$ , multiplying with  $\tilde{U}$  obtains the linear combination coefficients in the 2-dimensional subspace:

$$\tilde{U}^T = \begin{bmatrix} - & \textcolor{red}{u}^{(1)T} & - \\ - & \textcolor{green}{u}^{(2)T} & - \end{bmatrix} \quad x = \begin{bmatrix} 1 \\ \dots \\ x_d \end{bmatrix}$$

$$\tilde{U}^T x = \begin{bmatrix} \textcolor{red}{z}_1 \\ \textcolor{green}{z}_2 \end{bmatrix}$$



# SVD: Compression

$$r < N < d$$

We can use the matrix  $\tilde{U}^T$  to compress the data matrix  $X^T$ :

$$Z := \tilde{U}^T X^T$$

Here,  $Z$  is  $r \times N$  matrix.

**Intuition:** we project the data points  $x^{(i)}$  to the new (lower-dimensional) basis  $\tilde{U}$ .  
 $Z$  contains all the projection coefficients.

Can we de-compress (reconstruct) the original data?

# SVD: Reconstruction

$$r < N < d$$

**Theorem (informal):** For a fixed  $r$ , it holds that  $\tilde{U}\tilde{U}^T \approx \mathbb{I}$ , with approximation error dependent on  $r$ .

Thus, we can reconstruct the data in the original feature dimension  $d$ :

$$\tilde{X}^T := \tilde{U}Z$$

$$\tilde{X}^T = \tilde{U}\tilde{U}^T X^T \approx X^T$$

How much information is retained?



# Background: Expected Value

Given a discrete random variable  $X \in \{1, 2, \dots, M\}$  and a probability  $p(X)$  for the outcomes of the random variable.

The **expected value** is defined as  $\mathbb{E}_p[X] := \sum_{x=1}^M p(x)x$

Example: Uniform 6-sided dice.

$$\begin{aligned}\mathbb{E}_p[X] &= \sum_{x=1}^6 p(x)x \\ &= \frac{1}{6}(1 + 2 + 3 + 4 + 5 + 6) = \frac{21}{6} = 3.5\end{aligned}$$



# Background: Variance

Given a random variable  $X$  and a probability  $p(X)$  for the outcomes of the random variable.

The **variance** is defined as  $\mathbb{V}_p[X] := \mathbb{E}_p \left[ (X - \mathbb{E}_p[X])^2 \right]$ .

Example: Uniform 6-sided dice.

$$\begin{aligned}\mathbb{V}_p[X] &= \sum_{x=1}^6 p(x)(x - 3.5)^2 \\ &= \frac{1}{6}(2.5^2 + 1.5^2 + 0.5^2 + 0.5^2 + 1.5^2 + 2.5^2) = \frac{35}{12} \approx 2.92\end{aligned}$$



# Mean-centred Data

Assume we work with mean-centred data.

1. Compute mean feature vector over samples:  $\bar{x} = \frac{1}{N} \sum_{i=1}^N x^{(i)}$
2. Compute mean-centred data points:  $\hat{x}^{(i)} = x^{(i)} - \bar{x}$  for all  $i$ .
3. Use mean-centred data matrix  $\hat{X}^T \leftarrow [\hat{x}^{(1)}, \dots, \hat{x}^{(N)}]$ .

# SVD and Variances

$$\hat{X}^T = \begin{matrix} d \times d \\ \left[ \begin{array}{c|c|c|c} | & | & \dots & | \\ u^{(1)} & u^{(2)} & \dots & u^{(d)} \\ | & | & & | \end{array} \right] \end{matrix} \begin{matrix} d \times N \\ \left[ \begin{array}{c|c|c|c} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sigma_N \\ 0 & 0 & \dots & 0 \end{array} \right] \end{matrix} \begin{matrix} N \times N \\ \left[ \begin{array}{c|c|c|c} | & | & \dots & | \\ v^{(1)} & v^{(2)} & \dots & v^{(N)} \\ | & | & & | \end{array} \right]^T \end{matrix}$$

**Theorem (informal):** Given a mean-centred data matrix  $\hat{X}^T$ , the values  $\frac{\sigma_j^2}{N-1}$  are variances of the data in the basis defined by the vectors  $u^{(j)}$ .

# Compressing Data while Retaining Variance

Task: “Retain at least 99% of variance in the data”

Solution: Choose minimum  $r$ , such that  $\frac{\sum_{i=1}^r \sigma_i^2}{\sum_{i=1}^N \sigma_i^2} \geq 0.99$ .

**Theorem:** Using this  $r$ , the original data points  $\hat{x}^{(i)}$  (mean-centred) and reconstructed data points  $\tilde{x}^{(i)}$  are close as follows:

$$\frac{\sum_{i=1}^N \|\hat{x}^{(i)} - \tilde{x}^{(i)}\|^2}{\sum_{i=1}^N \|\hat{x}^{(i)}\|^2} \leq 0.01$$

# Principal Component Analysis (PCA)

**Statistics application** of SVD. Capture components that maximize the *statistical variations* of the data. Same idea, but uses sample covariance matrix as an input (will be discussed in tutorial).




Image credit: numxl

# Singular Value Decomposition (SVD)

## Algorithm:

- Outside of the scope of this course
- If you are curious:
  - <https://web.stanford.edu/class/cme335/spr11/lecture6.pdf>

## Implementations:

- Numpy: `numpy.linalg.svd`
  - <https://numpy.org/doc/stable/reference/generated/numpy.linalg.svd.html>
- Matlab: `svd`
  - <https://www.mathworks.com/help/matlab/ref/double.svd.html>

# Summary

- Unsupervised Learning: learn patterns of the data without labels
- K-means clustering
  - Algorithm: find **centroids** based on the data points, **assign each point to the closest centroid**
  - Measuring the quality of clusters: distance of each point to their centroid
  - Picking the number of clusters: **elbow method**, business needs
  - Variants: **K-medoids**, etc
- Dimensionality Reduction: finding **new basis** that best captures the data
  - Singular Value Decomposition (SVD)
  - Principal Component Analysis (PCA): **statistical application** of SVD



# Coming Up Next Week

- Neural networks
- Dr Conghui will take over starting from next week!

# To Do

- **Lecture Training 8**
  - +250 Free EXP
  - +100 Early bird bonus
- **Problem Set 4**
- **Mini Project**
  - ~3 weeks left!

