

National University of Singapore
School of Computing
CS2109S: Introduction to AI and Machine Learning
Semester II, 2024/2025

Tutorial 9
Backpropagation

Summary of Key Concepts

In this tutorial, we will discuss and explore the following learning points from Lecture:

1. Backpropagation
2. Matrix calculus for backpropagation
3. Potential issues with backpropagation
4. Dying ReLU problem

A Backpropagation (Warm Up)

Grace has a wine dataset which comprises of 1100 samples. Each wine sample has a label indicating which plant variety it comes from, and two features: colour intensity and alcohol level.

Now, she wants to build a classifier that can predict which plant variety a wine sample comes from using the two features. She decided to train a neural network with the architecture shown in the figure below.




Figure 1: neural network architecture for the i -th data point (out of 1100 samples). $X_{0i} = 1$ represents its bias term, and X_{1i} , X_{2i} represent its 1st and 2nd features, i.e., the color intensity and alcohol level, respectively.

Mathematically, this is given by:

$$\begin{aligned} f^{[1]} &= W^{[1]T} X \\ \hat{Y} &= g^{[1]}(f^{[1]}) \end{aligned}$$

where

$$\text{the weight matrix } W^{[1]} = \begin{bmatrix} W_{00}^{[1]} \\ W_{10}^{[1]} \\ W_{20}^{[1]} \end{bmatrix} \in \mathbb{R}^{3 \times 1}, \quad \text{input data matrix } X = \begin{bmatrix} 1 & 1 & \cdots & 1 \\ X_{10} & X_{11} & \cdots & X_{1,(n-1)} \\ X_{20} & X_{21} & \cdots & X_{2,(n-1)} \end{bmatrix} \in \mathbb{R}^{3 \times n},$$

$$\text{weighted sum matrix } f^{[1]} = \begin{bmatrix} f_1^{[1]} & f_2^{[1]} & \cdots & f_n^{[1]} \end{bmatrix} \in \mathbb{R}^{1 \times n},$$

$$\text{predicted value matrix } \hat{Y} = \begin{bmatrix} \hat{Y}_1 & \hat{Y}_2 & \cdots & \hat{Y}_n \end{bmatrix} \in \mathbb{R}^{1 \times n}.$$

activation function $g^{[1]}(s) = \sigma(s) = \frac{1}{1 + e^{-s}}$, which is applied pointwise to $f^{[1]}$ to produce \hat{Y} .

NOTE: Each data point corresponds to a column in input data matrix X , as well as an entry in $f^{[1]}$ and \hat{Y} . To emphasize, you should treat $f^{[1]}$, $W^{[1]}$, \hat{Y} , Y and X as matrices, and refer to their scalar entries using subscripts, e.g., $f_0^{[1]}$, $W_{10}^{[1]}$, \hat{Y}_2 .

In this classifier, she decided to use the following loss function:¹

$$\mathcal{E} = -\frac{1}{n} \sum_{i=0}^{n-1} \left\{ [Y_i \cdot \log(\hat{Y}_i)] + [(1 - Y_i) \log(1 - \hat{Y}_i)] \right\}$$

where $\hat{Y} \in (0, 1)^{1 \times n}$ after applying sigmoid function, $Y \in \{0, 1\}^{1 \times n}$ such that $Y_i = 1$ if the i -th wine sample is from plant variety A and $Y_i = 0$ if it is from plant variety B, and n is the number of wine samples (i.e., $n = 1100$ in this case). Furthermore, we take $\log(x)$ to be the natural logarithm in this tutorial.

To illustrate, we calculate the loss function for the following sample of two labels:

$$\begin{aligned} Y &= [0 \ 1] \\ \hat{Y} &= [0.2 \ 0.9] \end{aligned}$$

The loss function will be calculated as follows:

$$\begin{aligned} \mathcal{E} &= -\frac{1}{2} \left[(1 - 0) \cdot \log(1 - 0.2) + 1 \cdot \log(0.9) \right] \\ &= -\frac{1}{2} \left[\log(0.8) + \log(0.9) \right] \approx 0.16425 \end{aligned}$$

Note that $\log(1) = 0$ and $\log(0) \rightarrow -\infty$. We assume that $\log(x) = \ln(x)$ in this tutorial.

¹This loss function is usually known as *log loss* or *binary cross-entropy*.

For subquestions (1) and (2), to keep things simple, let us consider the case where $n = 1$.

1. **Show that**

$$(a) \quad \frac{\partial \mathcal{E}}{\partial \hat{Y}} = \left[-\frac{Y_0}{\hat{Y}_0} + \frac{1-Y_0}{1-\hat{Y}_0} \right]$$

We provide the answer for this part as a guiding example:

Note that for n (all) data points, $\frac{\partial \mathcal{E}}{\partial \hat{Y}} = \left[\frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \frac{\partial \mathcal{E}}{\partial \hat{Y}_1} \cdots \frac{\partial \mathcal{E}}{\partial \hat{Y}_n} \right]$.

If we assume $n = 1$, it becomes $\frac{\partial \mathcal{E}}{\partial \hat{Y}} = \left[\frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \right]$. Then, the only entry in matrix is

$$\frac{\partial \mathcal{E}}{\partial \hat{Y}_0} = -\frac{Y_0}{\hat{Y}_0} + \frac{1-Y_0}{1-\hat{Y}_0}$$

Therefore, the partial derivative matrix is

$$\frac{\partial \mathcal{E}}{\partial \hat{Y}} = \left[-\frac{Y_0}{\hat{Y}_0} + \frac{1-Y_0}{1-\hat{Y}_0} \right] \text{ when } n = 1$$

Try to plug in the case where (1) the true label $Y_0 = 0$ and (2) $Y_0 = 1$. This equation tells us how the change in predicted value \hat{Y} will be able to influence the change of error/loss value.

$$(b) \quad \frac{\partial \mathcal{E}}{\partial f^{[1]}} = \hat{Y} - Y$$

Hint: Since $n = 1$, $\frac{\partial \mathcal{E}}{\partial f^{[1]}} = \left[\frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \right]$. By chain rule, $\frac{\partial \mathcal{E}}{\partial f_0^{[1]}} = \frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \frac{\partial \hat{Y}_0}{\partial f_0^{[1]}} \dots$

$$(c) \quad \frac{\partial \mathcal{E}}{\partial W_{20}^{[1]}} = \left(\frac{\partial \mathcal{E}}{\partial f^{[1]}} \right)_0 X_{20}$$

Hint: Since $n = 1$, $f^{[1]} = [f_0^{[1]}]$. Thus, $\frac{\partial f^{[1]}}{\partial W_{20}^{[1]}} = \left[\frac{\partial f_0^{[1]}}{\partial W_{20}^{[1]}} \right] \dots$

NOTE: $\left(\frac{\partial \mathcal{E}}{\partial f^{[1]}} \right)_0$ refers to the first entry (at the 0-th index) of the matrix $\frac{\partial \mathcal{E}}{\partial f^{[1]}}$.

Solution:

(b) Since $n = 1$, $\frac{\partial \mathcal{E}}{\partial f^{[1]}} = \left[\frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \right]$. By chain rule, $\frac{\partial \mathcal{E}}{\partial f_0^{[1]}} = \frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \frac{\partial \hat{Y}_0}{\partial f_0^{[1]}}$.

Now, let us compute $\frac{\partial \hat{Y}_0}{\partial f_0^{[1]}}$.

$$\frac{\partial \hat{Y}_0}{\partial f_0^{[1]}} = \sigma(f_0^{[1]}) (1 - \sigma(f_0^{[1]}))$$

It is helpful for us to simplify this equation. Notice that

$$\hat{Y}_0 = \sigma(f_0^{[1]})$$

Thus,

$$\begin{aligned} \frac{\partial \hat{Y}_0}{\partial f_0^{[1]}} &= \hat{Y}_0 (1 - \hat{Y}_0) \\ \frac{\partial \hat{Y}}{\partial f^{[1]}} &= \left[\frac{\partial \hat{Y}_0}{\partial f_0^{[1]}} \right] = [\hat{Y}_0 (1 - \hat{Y}_0)] \end{aligned}$$

Then,

$$\begin{aligned}
 \frac{\partial \mathcal{E}}{\partial f^{[1]}} &= \left[\frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \right] \\
 &= \left[\frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \frac{\partial \hat{Y}_0}{\partial f_0^{[1]}} \right] \\
 &= \left[\left(-\frac{Y_0}{\hat{Y}_0} + \frac{1-Y_0}{1-\hat{Y}_0} \right) \times \hat{Y}_0(1-\hat{Y}_0) \right] \\
 &= \left[-Y_0(1-\hat{Y}_0) + (1-Y_0)\hat{Y}_0 \right] \\
 &= \left[\hat{Y}_0 - Y_0 \right] \\
 &= \hat{Y} - Y.
 \end{aligned}$$

(c) Since $n = 1$, $f^{[1]} = [f_0^{[1]}]$. Thus, $\frac{\partial f^{[1]}}{\partial W_{20}^{[1]}} = \left[\frac{\partial f_0^{[1]}}{\partial W_{20}^{[1]}} \right]$. Moreover,

$$f_0^{[1]} = \sum_{i=0}^2 (W^{[1]T})_{0i} X_{i0} = \sum_{i=0}^2 W_{i0}^{[1]} X_{i0} \in R$$

As such,

$$\frac{\partial f_0^{[1]}}{\partial W_{20}^{[1]}} = X_{20}$$

Then, by chain rule, we get

$$\begin{aligned}
 \frac{\partial \mathcal{E}}{\partial W_{20}^{[1]}} &= \frac{\partial \mathcal{E}}{\partial f_0^{[1]}} X_{20} \\
 &= \left(\frac{\partial \mathcal{E}}{\partial f^{[1]}} \right)_0 X_{20}
 \end{aligned}$$

To summarise, $\frac{\partial \mathcal{E}}{\partial \hat{Y}} = \left[-\frac{Y_0}{\hat{Y}_0} + \frac{1-Y_0}{1-\hat{Y}_0} \right]$, $\frac{\partial \mathcal{E}}{\partial f^{[1]}} = \hat{Y} - Y$, and $\frac{\partial \mathcal{E}}{\partial W_{20}^{[1]}} = \left(\frac{\partial \mathcal{E}}{\partial f^{[1]}} \right)_0 X_{20}$.

NOTE: $\frac{\partial \mathcal{E}}{\partial \hat{Y}}$, and $\frac{\partial \mathcal{E}}{\partial f^{[1]}}$ are vectors since \mathcal{E} is a scalar, but \hat{Y} and $f^{[1]}$ are vectors. However, $\frac{\partial \mathcal{E}}{\partial W_{20}^{[1]}}$ is a scalar since $W_{20}^{[1]}$ is a scalar.

2. Using your answer in (1), derive an expression for $\frac{\partial \mathcal{E}}{\partial W^{[1]}}$. What can you observe from this expression?

Solution:

Observe that our solution in (1) also applies to other $\frac{\partial \mathcal{E}}{\partial W_{i0}^{[1]}}$, where $0 \leq i \leq 2$.

Furthermore,

$$\begin{aligned}
 \frac{\partial \mathcal{E}}{\partial W^{[1]}} &= \left[\frac{\partial \mathcal{E}}{\partial W_{00}^{[1]}} \frac{\partial \mathcal{E}}{\partial W_{10}^{[1]}} \frac{\partial \mathcal{E}}{\partial W_{20}^{[1]}} \right]^T \\
 &= \left[\frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \frac{\partial f_0^{[1]}}{\partial W_{00}^{[1]}} \quad \frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \frac{\partial f_0^{[1]}}{\partial W_{10}^{[1]}} \quad \frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \frac{\partial f_0^{[1]}}{\partial W_{20}^{[1]}} \right]^T \\
 &= \frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \times \left[\frac{\partial f_0^{[1]}}{\partial W_{00}^{[1]}} \quad \frac{\partial f_0^{[1]}}{\partial W_{10}^{[1]}} \quad \frac{\partial f_0^{[1]}}{\partial W_{20}^{[1]}} \right]^T \\
 &= \frac{\partial \mathcal{E}}{\partial f_0^{[1]}} \times \begin{bmatrix} X_{00} & X_{10} & X_{20} \end{bmatrix}^T
 \end{aligned}$$

Hence, we can conclude that

$$\frac{\partial \mathcal{E}}{\partial W^{[1]}} = \left(\frac{\partial \mathcal{E}}{\partial f^{[1]}} \right)_0 X$$

Note that $\frac{\partial \mathcal{E}}{\partial W^{[1]}} \in \mathbb{R}^{3 \times 1}$, $\left(\frac{\partial \mathcal{E}}{\partial f^{[1]}} \right)_{00}$ is a scalar, and that $X \in \mathbb{R}^{3 \times 1}$.

We can observe that the gradient $\frac{\partial \mathcal{E}}{\partial W^{[1]}}$ is directly proportional to the input X .

3. Let us consider a general case where $n \in \mathbb{N}$. Using your answer to (1), find $\frac{\partial \mathcal{E}}{\partial f^{[1]}}$ when $n \in \mathbb{N}$. Check your answer using the Python notebook.

Hint: Read the Python notebook.

Solution:

$$\begin{aligned}
 \frac{\partial \mathcal{E}}{\partial f^{[1]}} &= \frac{1}{n} \left[(\hat{Y}_0 - Y_0) \quad (\hat{Y}_1 - Y_1) \quad \dots \quad (\hat{Y}_{(n-1)} - Y_{(n-1)}) \right] \\
 &= \frac{1}{n} (\hat{Y} - Y)
 \end{aligned}$$

Note that the $\frac{1}{n}$ comes from the loss function.

A more rigorous way of deriving the above is as follows:

$$\begin{aligned}
 \frac{\partial \mathcal{E}}{\partial f^{[1]}} &= \frac{\partial \mathcal{E}}{\partial \hat{Y}} \frac{\partial \hat{Y}}{\partial f^{[1]}} \\
 &= \frac{1}{n} \left[\frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \frac{\partial \mathcal{E}}{\partial \hat{Y}_1} \frac{\partial \mathcal{E}}{\partial \hat{Y}_2} \dots \frac{\partial \mathcal{E}}{\partial \hat{Y}_{(n-1)}} \right] \left[\frac{\partial \hat{Y}_i}{\partial f_j^{[1]}} \right]_{\substack{i=0,1,\dots,(n-1) \\ j=0,1,\dots,(n-1)}}
 \end{aligned}$$

where $\left[\frac{\partial \hat{Y}_{0i}}{\partial f_j^{[1]}} \right]_{\substack{i=0,1,\dots,(n-1) \\ j=0,1,\dots,(n-1)}} \in \mathbb{R}^{n \times n}$ and $\frac{\partial \hat{Y}_{0i}}{\partial f_j^{[1]}}$ indicates the (i, j) -th cell of the matrix.

We then have that:

$$\frac{\partial \mathcal{E}}{\partial \hat{Y}_i} = -\frac{Y_i}{\hat{Y}_i} + \frac{1 - Y_i}{1 - \hat{Y}_i}$$

$$\frac{\partial \hat{Y}_i}{\partial f_j^{[1]}} = \begin{cases} 0, & \text{if } i \neq j \\ \sigma(f_i^{[1]})(1 - \sigma(f_i^{[1]})) = \hat{Y}_i(1 - \hat{Y}_i) & \text{otherwise} \end{cases}$$

Hence we can conclude that:

$$\frac{\partial \mathcal{E}}{\partial f^{[1]}} = \frac{1}{n}(\hat{Y} - Y)$$

4. Let's say that Grace has 100 samples from plant variety A and 1000 samples from plant variety B that make up the 1100 total samples. To deal with the imbalanced data set, she decided to introduce two hyper-parameters α and β in the loss function, as shown below:

$$\mathcal{E} = -\frac{1}{n} \sum_{i=0}^{n-1} \left\{ \alpha [Y_i \cdot \log(\hat{Y}_i)] + \beta [(1 - Y_i) \cdot \log(1 - \hat{Y}_i)] \right\}$$

For example, using the same sample of two labels from part (1), the loss function can be computed as follows:

$$\begin{aligned} \mathcal{E} &= -\frac{1}{2} \left\{ \beta [(1 - 0) \cdot \log(1 - 0.2)] + \alpha [1 \cdot \log(0.9)] \right\} \\ &= -\frac{1}{2} \left\{ \beta \cdot \log(0.8) + \alpha \cdot \log(0.9) \right\} \end{aligned}$$

Why do you think that she introduced the hyper-parameters α and β ? How should she set their values?

Solution:

The hyper-parameters act as weights that determine how much each class contributes to the loss function. By setting $\alpha > \beta$, we can mitigate the problem of having a highly unbalanced dataset.

There are multiple possible answers for α and β , but generally, their relationship should be $\alpha \approx 10\beta$ (since the dataset has 10 times more samples from plant variety B than those from plant variety A).

In other words, we punish the machine more heavily if it misclassifies A. That way, the machine won't be biased towards predicting all samples as B.

B Backpropagation for a Deep(er) Network

After training the neural network described in question 1, Grace observed that the training error is high. She thus decided to introduce a hidden layer, resulting in the architecture shown in the figure below.




Figure 2: Figure for Question B

Mathematically, the network is given by

$$f^{[1]} = W^{[1]T} X$$

$$a^{[1]} = g^{[1]}(f^{[1]})$$

$$f^{[2]} = W^{[2]T} a^{[1]}$$

$$\hat{Y} = g^{[2]}(f^{[2]})$$

where $g^{[1]}(s) = \text{ReLU}(s)$, $g^{[2]}(s) = \sigma(s) = \frac{1}{1+e^{-s}}$, $W^{[1]} \in \mathbb{R}^{3 \times 2}$ and $W^{[2]} \in \mathbb{R}^{2 \times 1}$.

The *ReLU* function is defined as follows:

$$\text{ReLU}(s) = \begin{cases} s & \text{if } s > 0 \\ 0 & \text{otherwise} \end{cases}$$

Note: In Question 1, we only had **one neuron** in $f^{[1]}$, making $f^{[1]}$ a matrix with a single row.

However, in this question, we have two layers of activations, denoted as $f^{[1]}$ and $f^{[2]}$.

- First layer $f^{[1]} \in \mathbb{R}^{2 \times n}$:

Since this layer has **two neurons**, each producing an output for every data point, $f^{[1]}$ is a matrix with **two rows**—one for each neuron.

- Second layer $f^{[2]} \in \mathbb{R}^{1 \times n}$:

This layer has **only one neuron**, so $f^{[2]}$ is a matrix with **single row**.

Mathematically, we represent them as follows:

$$\text{first layer } f^{[1]} = \begin{bmatrix} f_{00}^{[1]} & f_{01}^{[1]} & \cdots & f_{0n}^{[1]} \\ f_{10}^{[1]} & f_{11}^{[1]} & \cdots & f_{1n}^{[1]} \end{bmatrix} = \begin{bmatrix} \text{outputs of 1st neuron} \\ \text{outputs of 2nd neuron} \end{bmatrix}$$

$$\text{second layer } f^{[2]} = \begin{bmatrix} f_0^{[2]} & f_1^{[2]} & \cdots & f_n^{[2]} \end{bmatrix}$$

Similar to question 1, let us keep things simple and consider what happens when $n = 1$. In addition, assume that the loss function used remains the same.

1. Compute $\frac{\partial \mathcal{E}}{\partial W_{11}^{[1]}}$.

Your answer should not contain any partial derivatives on the right hand side (i.e., they should all be evaluated). What can you observe from this expression?

Hint: The results found/observations made in question 1 are likely to be useful here.

Solution:

Based on our previous observations in question 1, and the fact that the only entry in $f^{[1]}$ that is dependent on W_{11} is its (1, 0) entry, we get

$$\begin{aligned}\frac{\partial \mathcal{E}}{\partial W_{11}^{[1]}} &= \frac{\partial \mathcal{E}}{\partial \hat{Y}_0} \frac{\partial \hat{Y}_0}{\partial f_0^{[2]}} \frac{\partial f_0^{[2]}}{\partial a_{10}^{[1]}} \frac{\partial a_{10}^{[1]}}{\partial f_{10}^{[1]}} \frac{\partial f_{10}^{[1]}}{\partial W_{11}^{[1]}} \\ \frac{\partial \mathcal{E}}{\partial \hat{Y}_0} &= -\frac{\alpha Y_0}{\hat{Y}_0} + \frac{\beta(1 - Y_0)}{1 - \hat{Y}_0} \\ \frac{\partial \hat{Y}_0}{\partial f_0^{[2]}} &= \sigma(f_0^{[2]}) (1 - \sigma(f_0^{[2]})) \\ \frac{\partial f_0^{[2]}}{\partial a_{10}^{[1]}} &= W_{10}^{[2]} \\ \frac{\partial a_{10}^{[1]}}{\partial f_{10}^{[1]}} &= \begin{cases} 0, & \text{if } f_{10}^{[1]} \leq 0 \\ 1, & \text{otherwise} \end{cases} \\ &= 1_{f_{10}^{[1]} > 0}\end{aligned}$$

where $1_{f_{10}^{[1]} > 0}$ is an indicator function which makes notation simpler, especially for the final expression which we are supposed to derive.

$$\frac{\partial f_{10}^{[1]}}{\partial W_{11}^{[1]}} = X_{10}$$

Therefore,

$$\frac{\partial \mathcal{E}}{\partial W_{11}^{[1]}} = \left[-\frac{\alpha Y_0}{\hat{Y}_0} + \frac{\beta(1 - Y_0)}{1 - \hat{Y}_0} \right] \sigma(f_0^{[2]}) (1 - \sigma(f_0^{[2]})) W_{10}^{[2]} 1_{f_{10}^{[1]} > 0} X_{10}$$

Notice that in the first question with no hidden layer, the partial derivative computed is directly proportional to X . However, notice that in this case, the partial derivative in the earlier layer is influenced by the weights in the later layers.

C Potential Issues with Training Deep Neural Networks

Suppose we use σ (the sigmoid function) as our activation function in a neural network with 50 hidden layers, as per the code in the accompanying Python notebook.

1. Play around with the code. Notice that when performing backpropagation, the gradient magnitudes of the first few layers are extremely small. What do you think

causes this problem?

Solution:

This problem is known as the *vanishing gradient* problem.

Referring to the previous question, observe that to compute $\frac{\partial \mathcal{E}}{\partial W^{[l]}}$, we need to take a product of many derivatives. Moreover, recall that the derivative of σ returns a value that is certainly between 0 and 0.25. Therefore, if we multiply many of such values together, we will end up with a very small number. Consequently, our change in weights for each iteration may be very small, causing convergence to be slow, if not impossible.

- Based on what we have learnt thus far, how can we **mitigate** this problem? Test out your solution by modifying the code and checking the gradient magnitudes.

Hint: You don't need to change much.

Solution:

A possible mitigation is to use the ReLU function (or its variants). As long as the input into the ReLU function is non-negative, its derivative will be 1, thereby precluding the aforementioned issue.

D Dying ReLU Problem

This problem occurs when majority of the activations are 0 (meaning the underlying pre-activations are mostly non-positive), resulting in the network dying midway. The gradients passed back are also 0 which leads to poor gradient descent performance and hence poor learning. Refer to the figure below on the ReLU and Leaky ReLU activation functions.




Figure 3: The Rectified Linear Unit (ReLU) (left) vs The Leaky Rectified Linear Unit (Leaky ReLU) with a as the slope when the values are negative. (right)

- How does Leaky ReLU fix this? What happens if we set $a = 1$ in Leaky ReLU?

Solution:

If a ReLU activation is dead, it will always output 0 for any input. When a ReLU unit ends up in this state, it is unlikely to recover because the function

gradient at any $x \leq 0$ is 0. Leaky ReLU has a small positive gradient a for negative inputs. Having a small positive gradient for the negative inputs gives the network a chance to recover. However, this depends on the dataset used and the value of a set prior to training. When $a = 1$, our activation function simply becomes a linear function and therefore our neural network layers degenerates back to a single matrix multiplication as shown in Tutorial 8.

E Appendix

To learn more about matrix differentiation—particularly how dimensions change throughout the calculations—refer to the following article: <https://medium.com/analytics-vidhya/matrix-calculus-for-machine-learning-e0262f0eaa8e>